



19. Oktober 2017

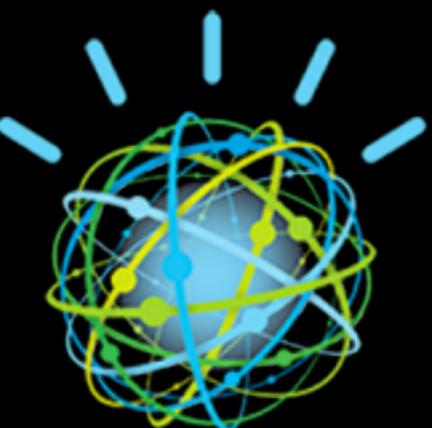
Markthalle Basel

Reserviere jetzt Dein Ticket

DeepLearning on the JVM - ApacheSystemML and DeepLearning4J on ApacheSpark

DeepLearning frameworks are popping up at very high frequency but only a few of them are suitable to run on clusters, use GPUs and supporting topologies beyond Feed-Forward at the same time. DeepLearning4J features all this without forcing you to learn new exotic programming languages and in addition also scales-out on well established infrastructures like ApacheSpark and Hadoop/YARN.

In this talk we will introduce DeepLearning4J on top of ApacheSpark with an example to create an anomaly detector for IoT sensor data with a LSTM auto encoder neural network.



IBM Watson



IBM Disclaimer

- THE INFORMATION CONTAINED IN THIS PRESENTATION IS PROVIDED FOR INFORMATIONAL PURPOSES ONLY.
- WHILE EFFORTS WERE MADE TO VERIFY THE COMPLETENESS AND ACCURACY OF THE INFORMATION CONTAINED IN THIS PRESENTATION, IT IS PROVIDED "AS IS", WITHOUT WARRANTY OF ANY KIND, EXPRESS OR IMPLIED.
- IBM SHALL NOT BE RESPONSIBLE FOR ANY DAMAGES ARISING OUT OF THE USE OF, OR OTHERWISE RELATED TO, THIS PRESENTATION OR ANY OTHER DOCUMENTATION.
- NOTHING CONTAINED IN THIS PRESENTATION IS INTENDED TO, OR SHALL HAVE THE EFFECT OF:
 - CREATING ANY WARRANTY OR REPRESENTATION FROM IBM (OR ITS AFFILIATES OR ITS OR THEIR SUPPLIERS AND/OR LICENSORS); OR
 - ALTERING THE TERMS AND CONDITIONS OF THE APPLICABLE LICENSE AGREEMENT GOVERNING THE USE OF IBM SOFTWARE.

IBM's statements regarding its plans, directions, and intent are subject to change or withdrawal without notice at IBM's sole discretion. Information regarding potential future products is intended to outline our general product direction and it should not be relied on in making a purchasing decision. The information mentioned regarding potential future products is not a commitment, promise, or legal obligation to deliver any material, code or functionality. Information about potential future products may not be incorporated into any contract. The development, release, and timing of any future features or functionality described for our products remains at our sole discretion.

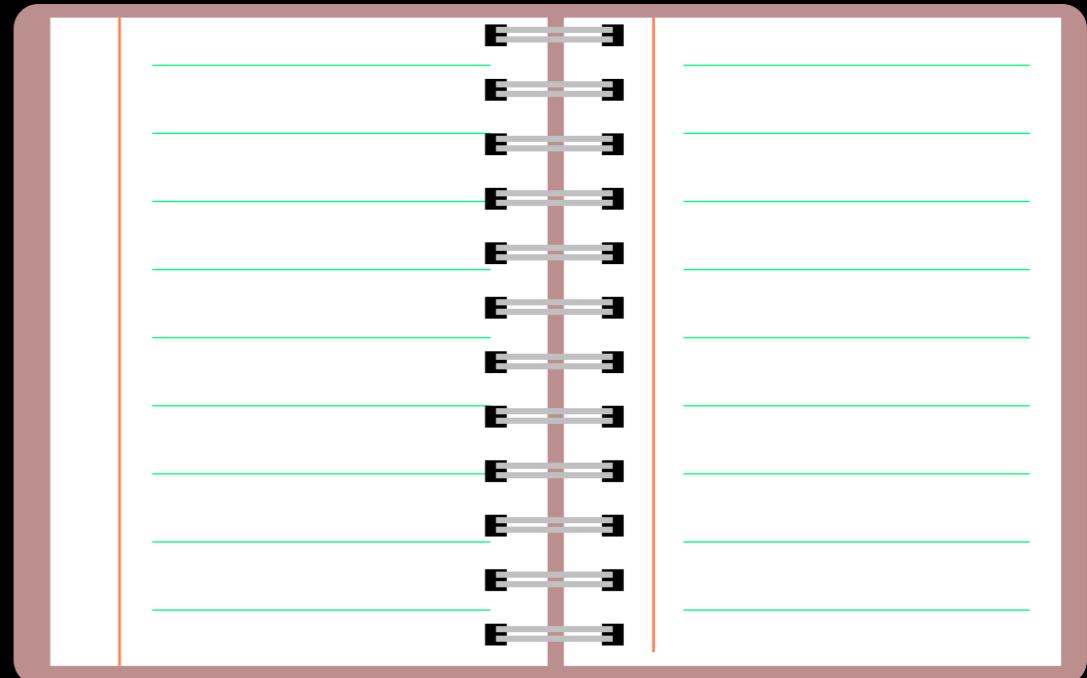
We are looking forward to hearing your suggestions and feedback so that we can improve our solutions and products.

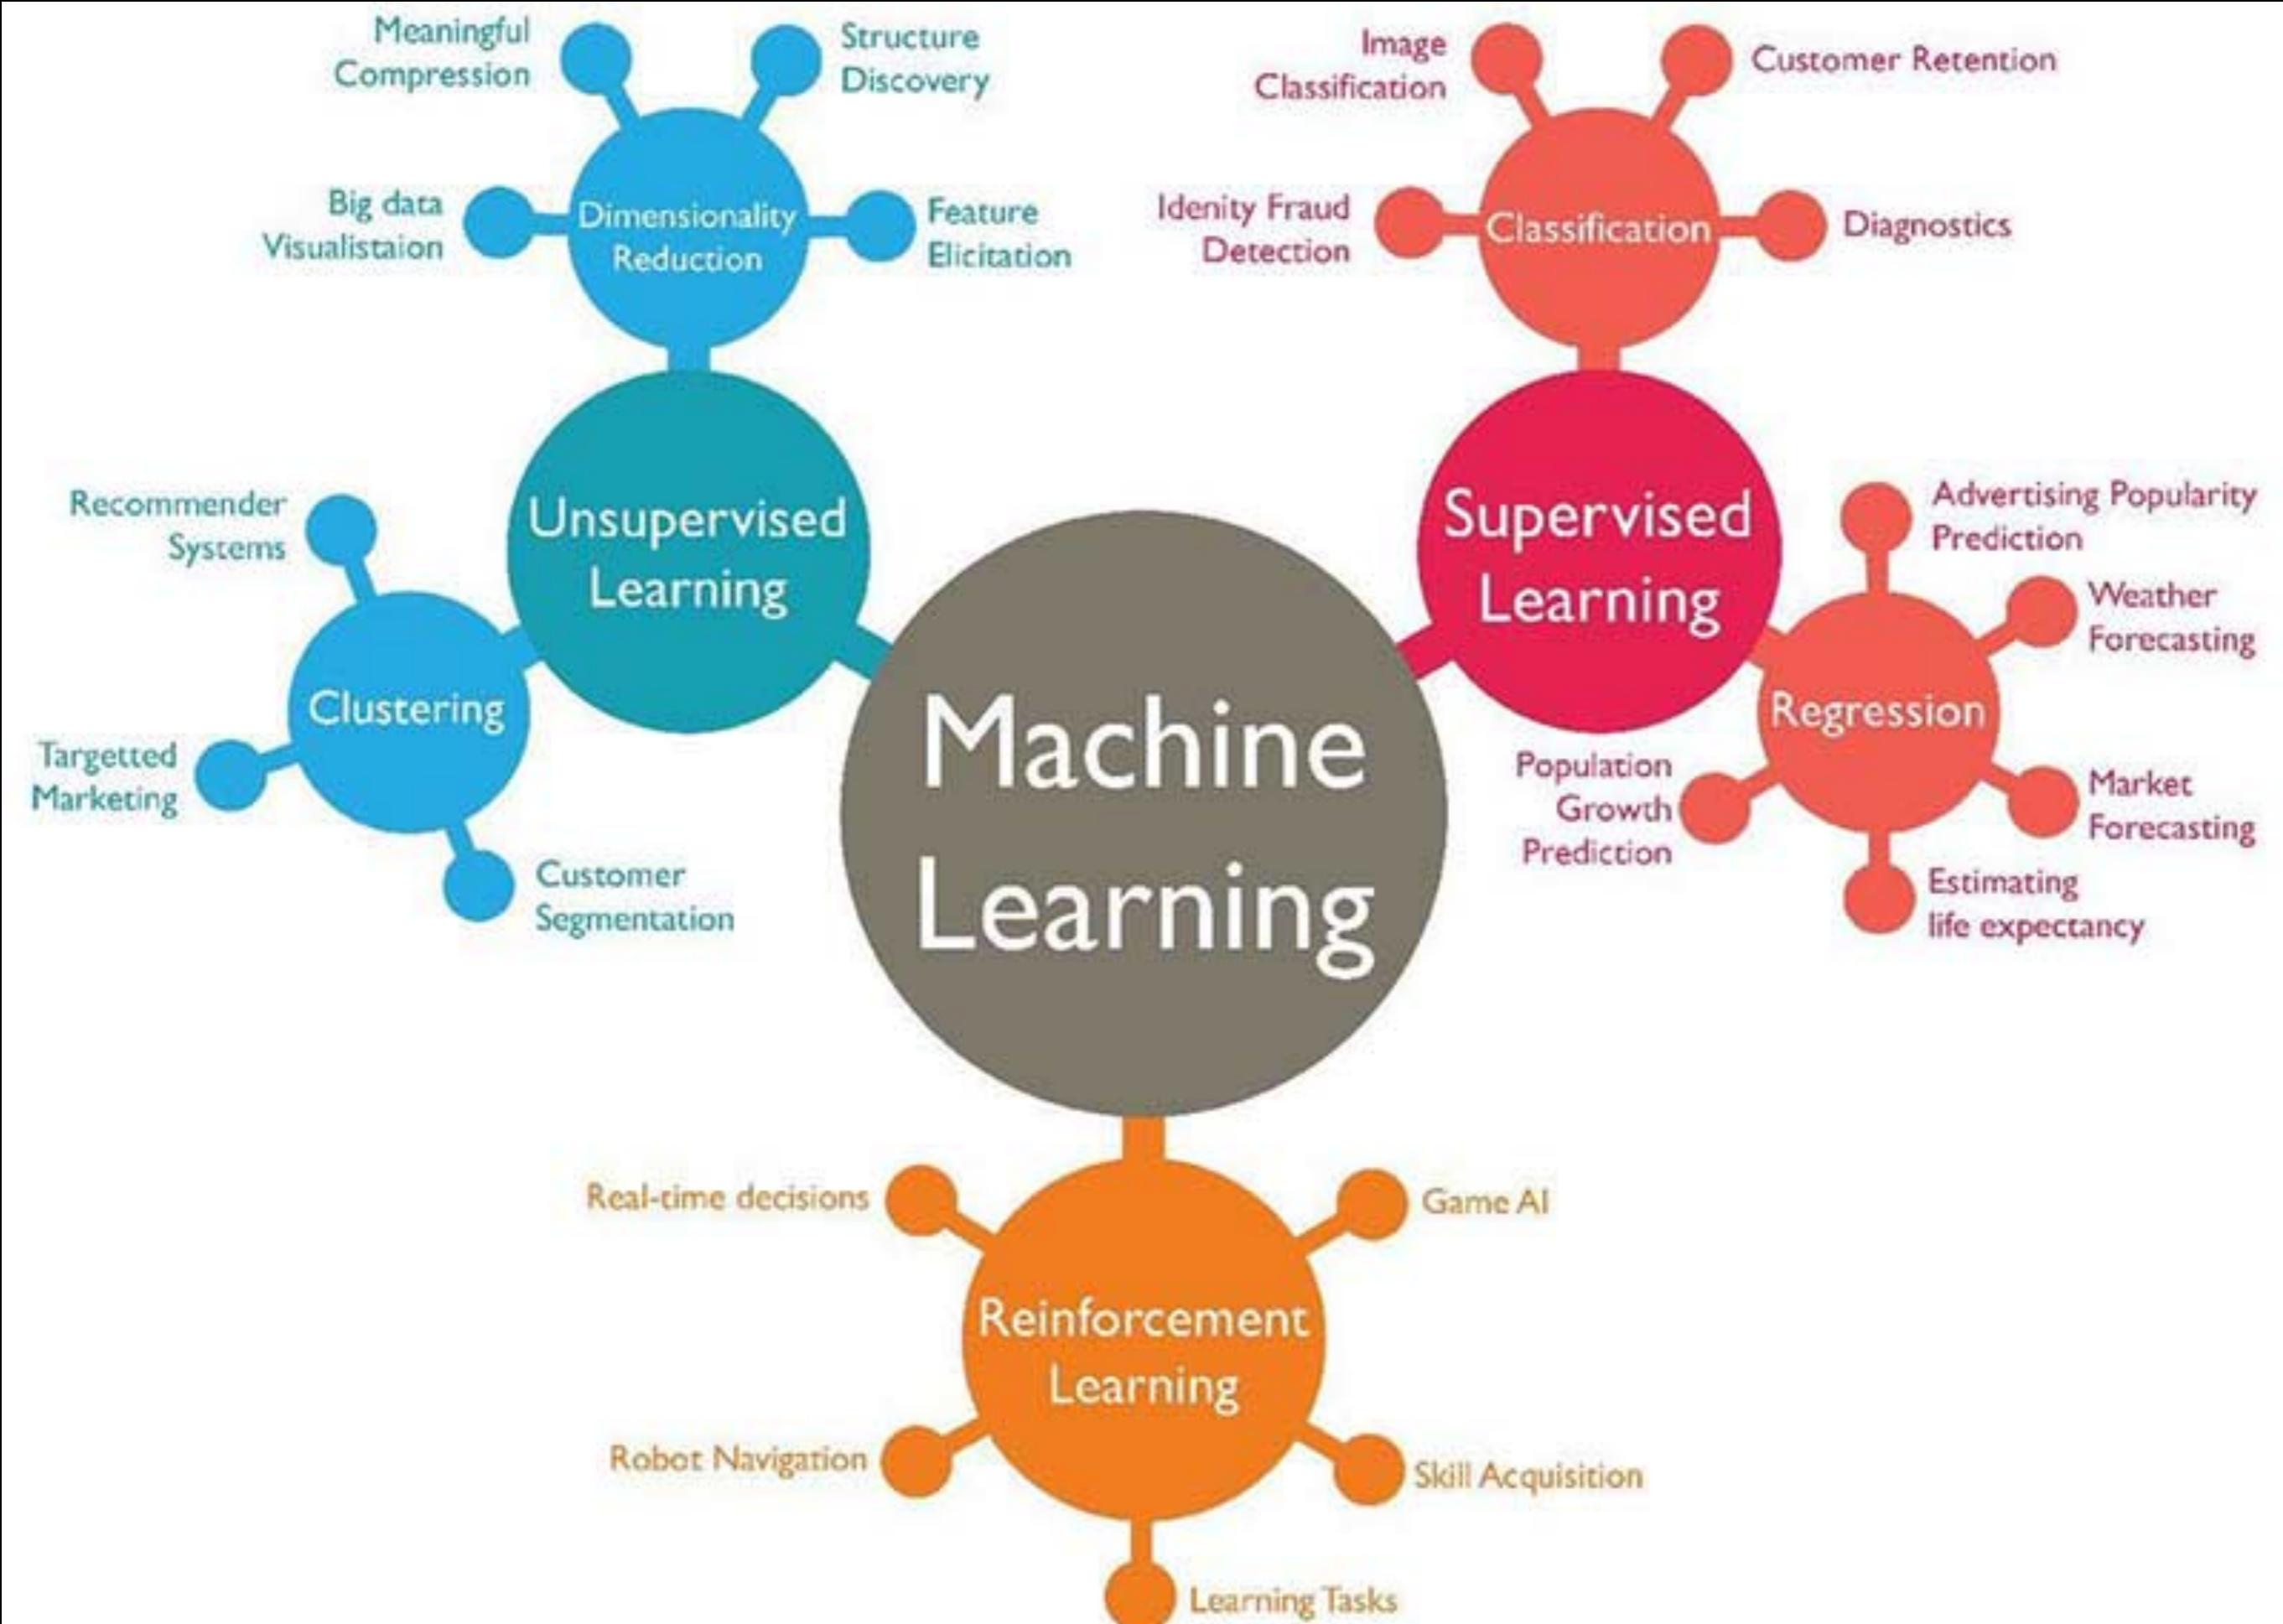
By joining this web conference, or clicking the "Join" link to join the ITM vNext Open Beta Community, you agree to the following terms, in addition to those set forth in the developerWorks terms of use:

IBM shall be free to use for any purpose and without restriction any oral or written suggestions or feedback that you provide to IBM. By providing IBM with any information or material, you grant IBM an unrestricted, irrevocable license to copy, reproduce, publish, upload, post, transmit, distribute, publicly display, perform, modify, create derivative works from, and otherwise freely use, those materials or information. You also agree that IBM is free to use any ideas, concepts, know-how, or techniques that you provide us for any purpose.

Agenda

- **Introduction to DeepLearning** ✓
- **Introduction to ApacheSpark**
- **Introduction to DL4J**
- **Introduction to SystemML**
- **Demos**





Neural Networks

Part No.	Max Temp. 1	Min Temp 1	Max Vibration 1	Outcome
100	35	35	12	1
101	46	35	21	1
130	56	46	3412	0

Neural Networks

Part No.	Max Temp. 1	Min Temp 1	Max Vibration 1	Outcome
100	35	35	12	1
101	46	35	21	1
130	56	46	3412	0

```
def predict(datapoint):
    if datapoint.MaxVibration1>100:
        return 0
    else:
        return 1
```

Neural Networks

Part No.	Max Temp. 1	Min Temp 1	Max Vibration 1	Outcome
100	35	35	12	1
101	46	35	21	1
130	56	46	3412	0

```
def predict(dp):
    return a+b*dp.MaxTemp1+c*dp.MinTemp1+d*dp.MaxVibration1
```

Neural Networks

Part No.	Max Temp. 1	Min Temp 1	Max Vibration 1	Outcome
100	35	35	12	1
101	46	35	21	1
130	56	46	3412	0

```
import math

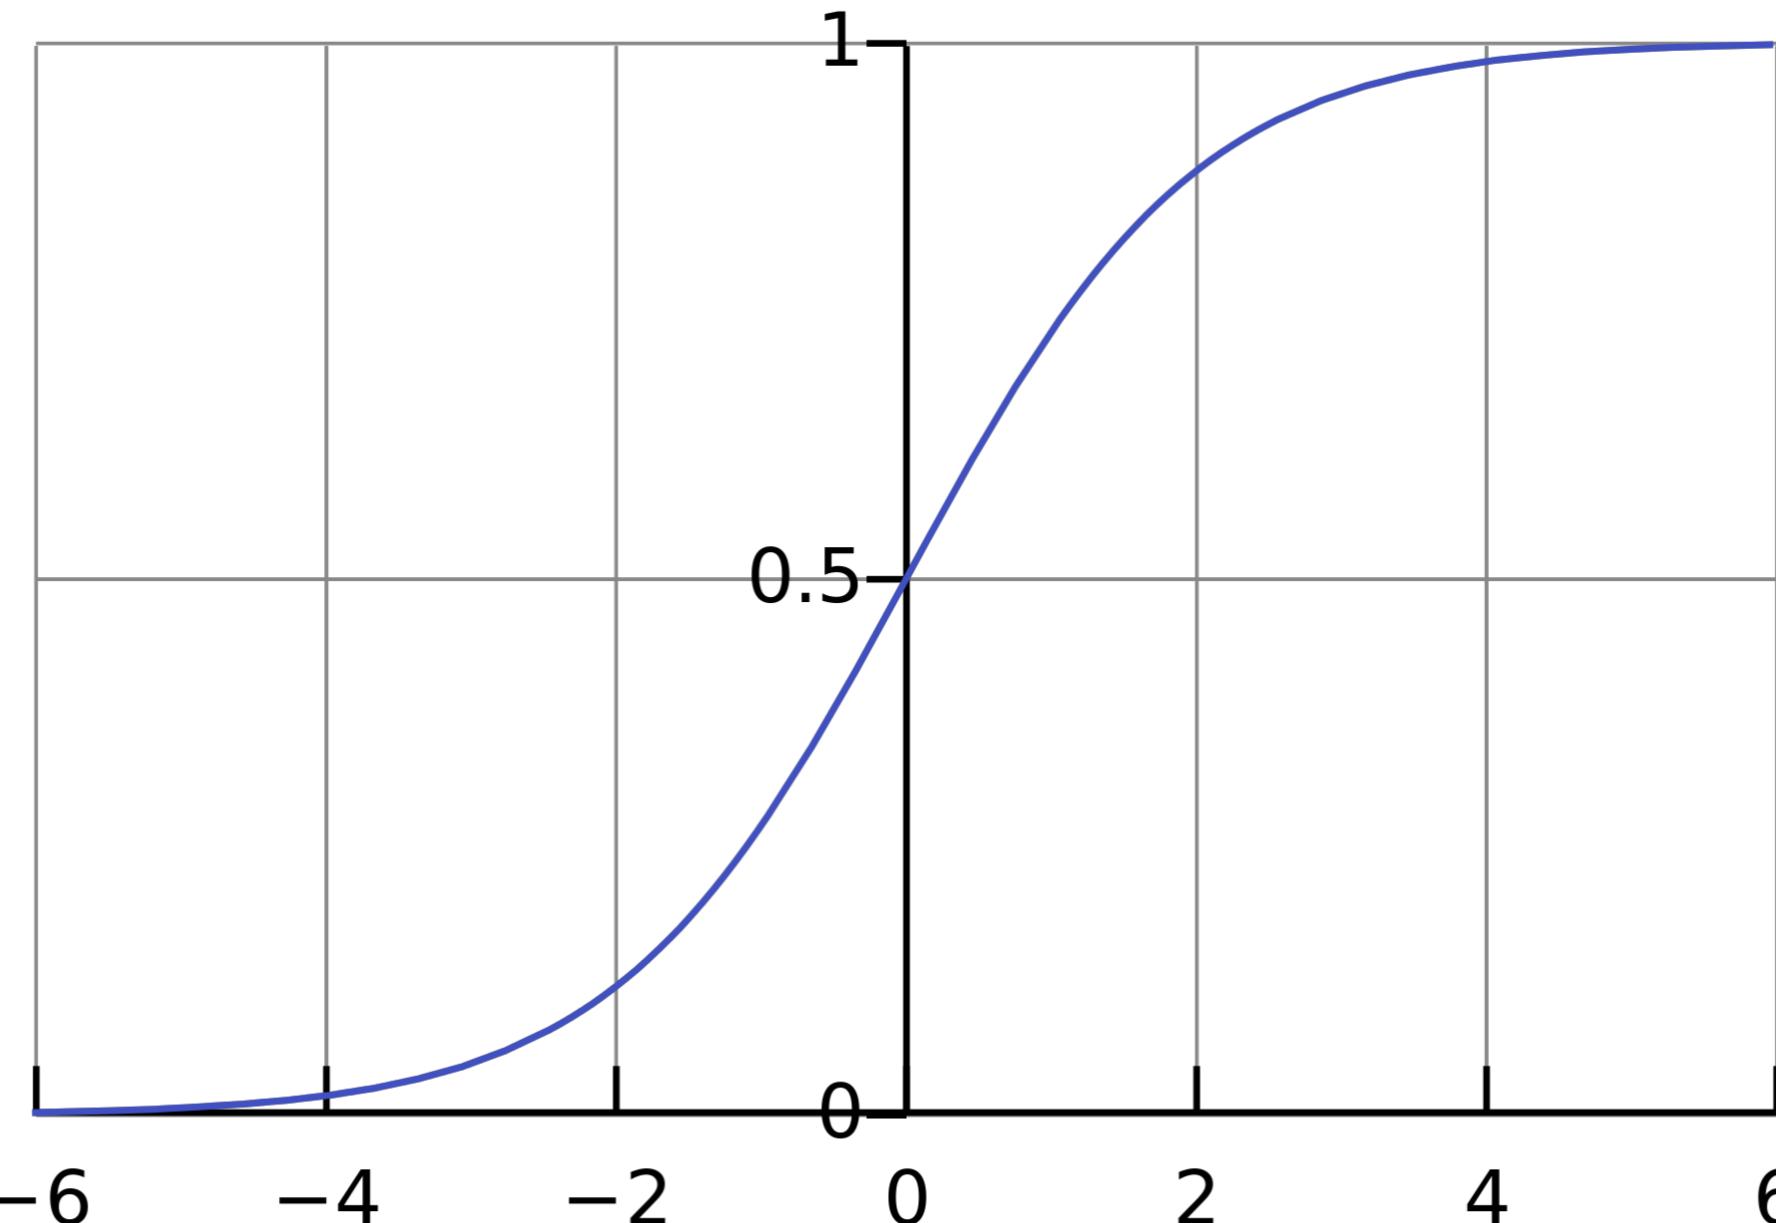
def sigmoid(x):
    return 1 / (1 + math.exp(-x))

def predict(dp):
    return sigmoid(a+b*dp.MaxTemp1+c*dp.MinTemp1+d*dp.MaxVibration1)
```

Neural Networks

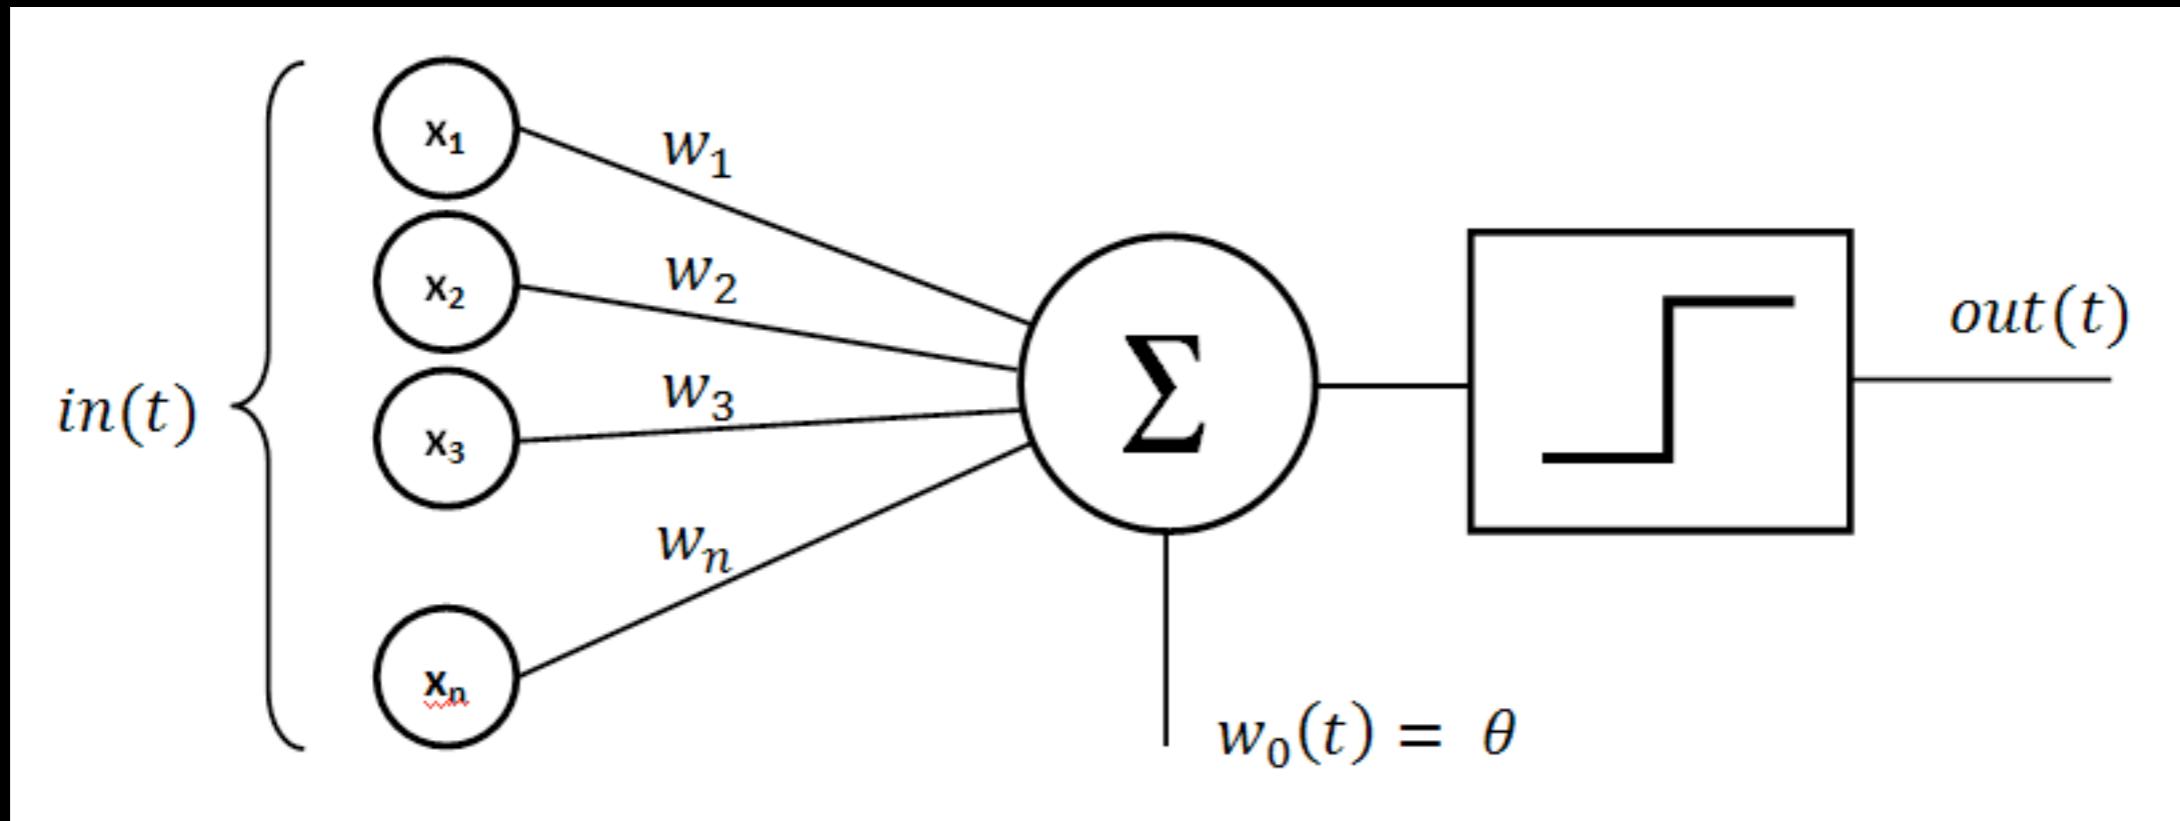
Part No.
100
101
130

```
import math  
  
def sigmoid(x):  
    return 1 / (1 + math.exp(-x))  
  
def predict(dp):  
    return sigmoid(a+b*dp.MaxTemp1+c*dp.MinTemp1+d*dp.MaxVibration1)
```

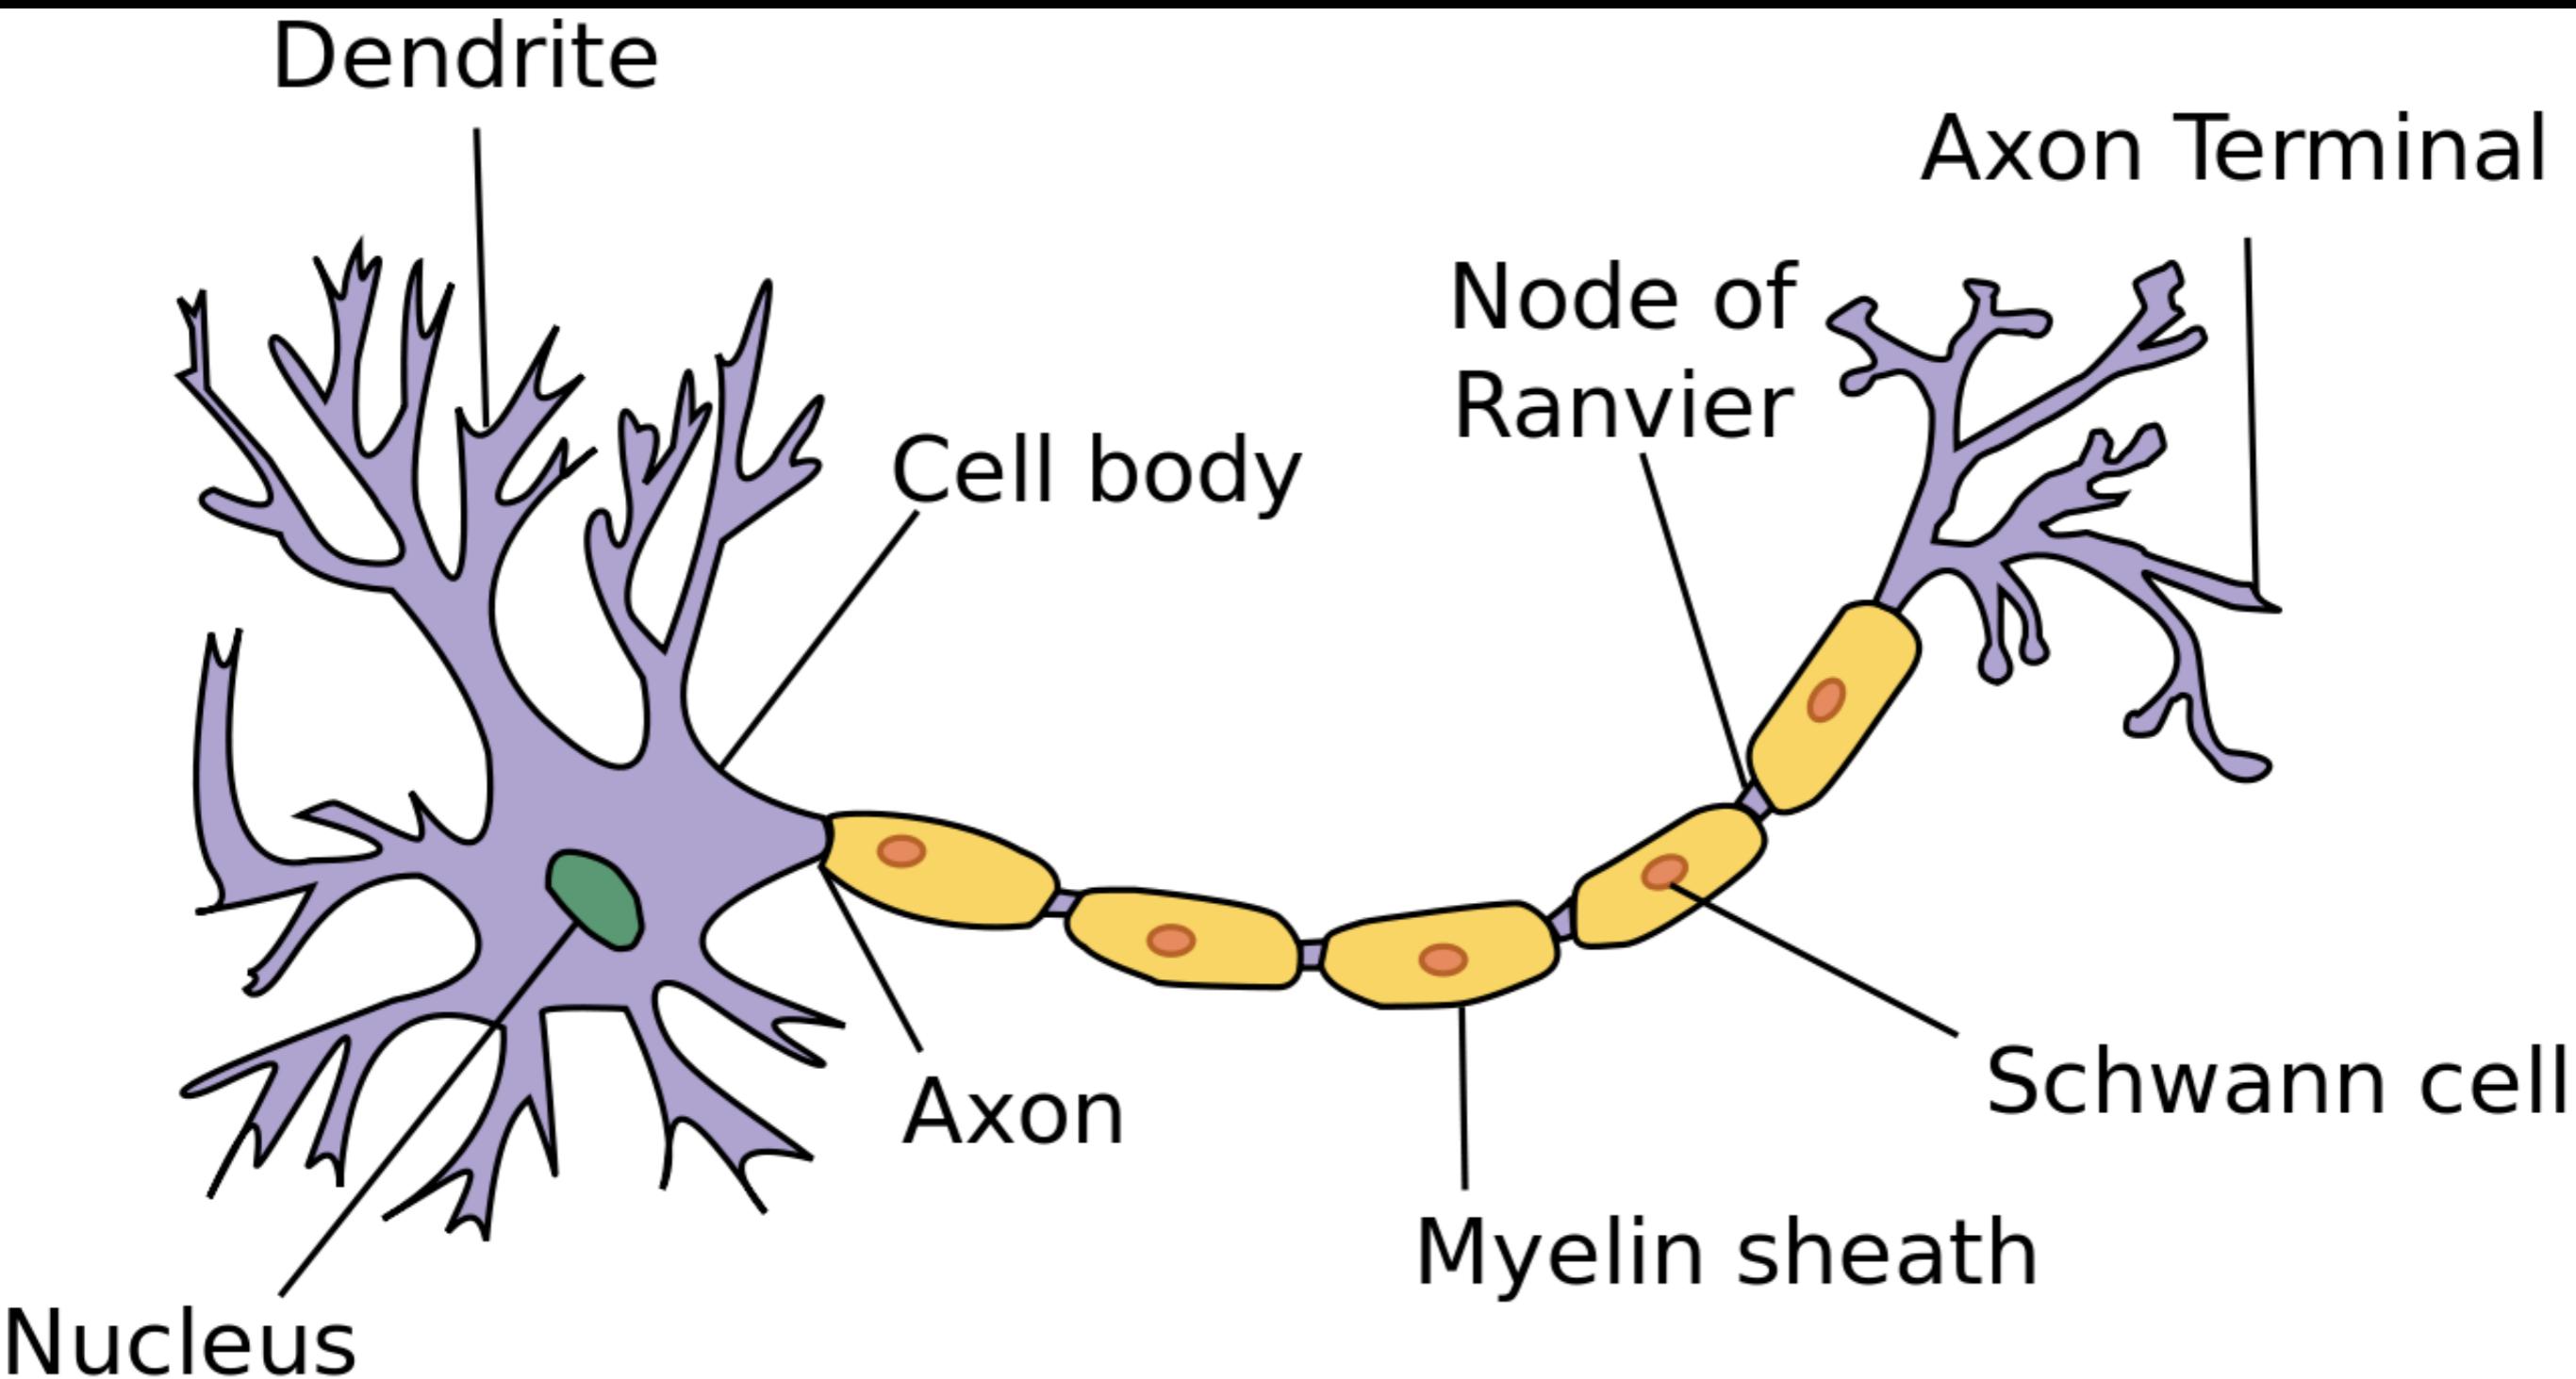


ne

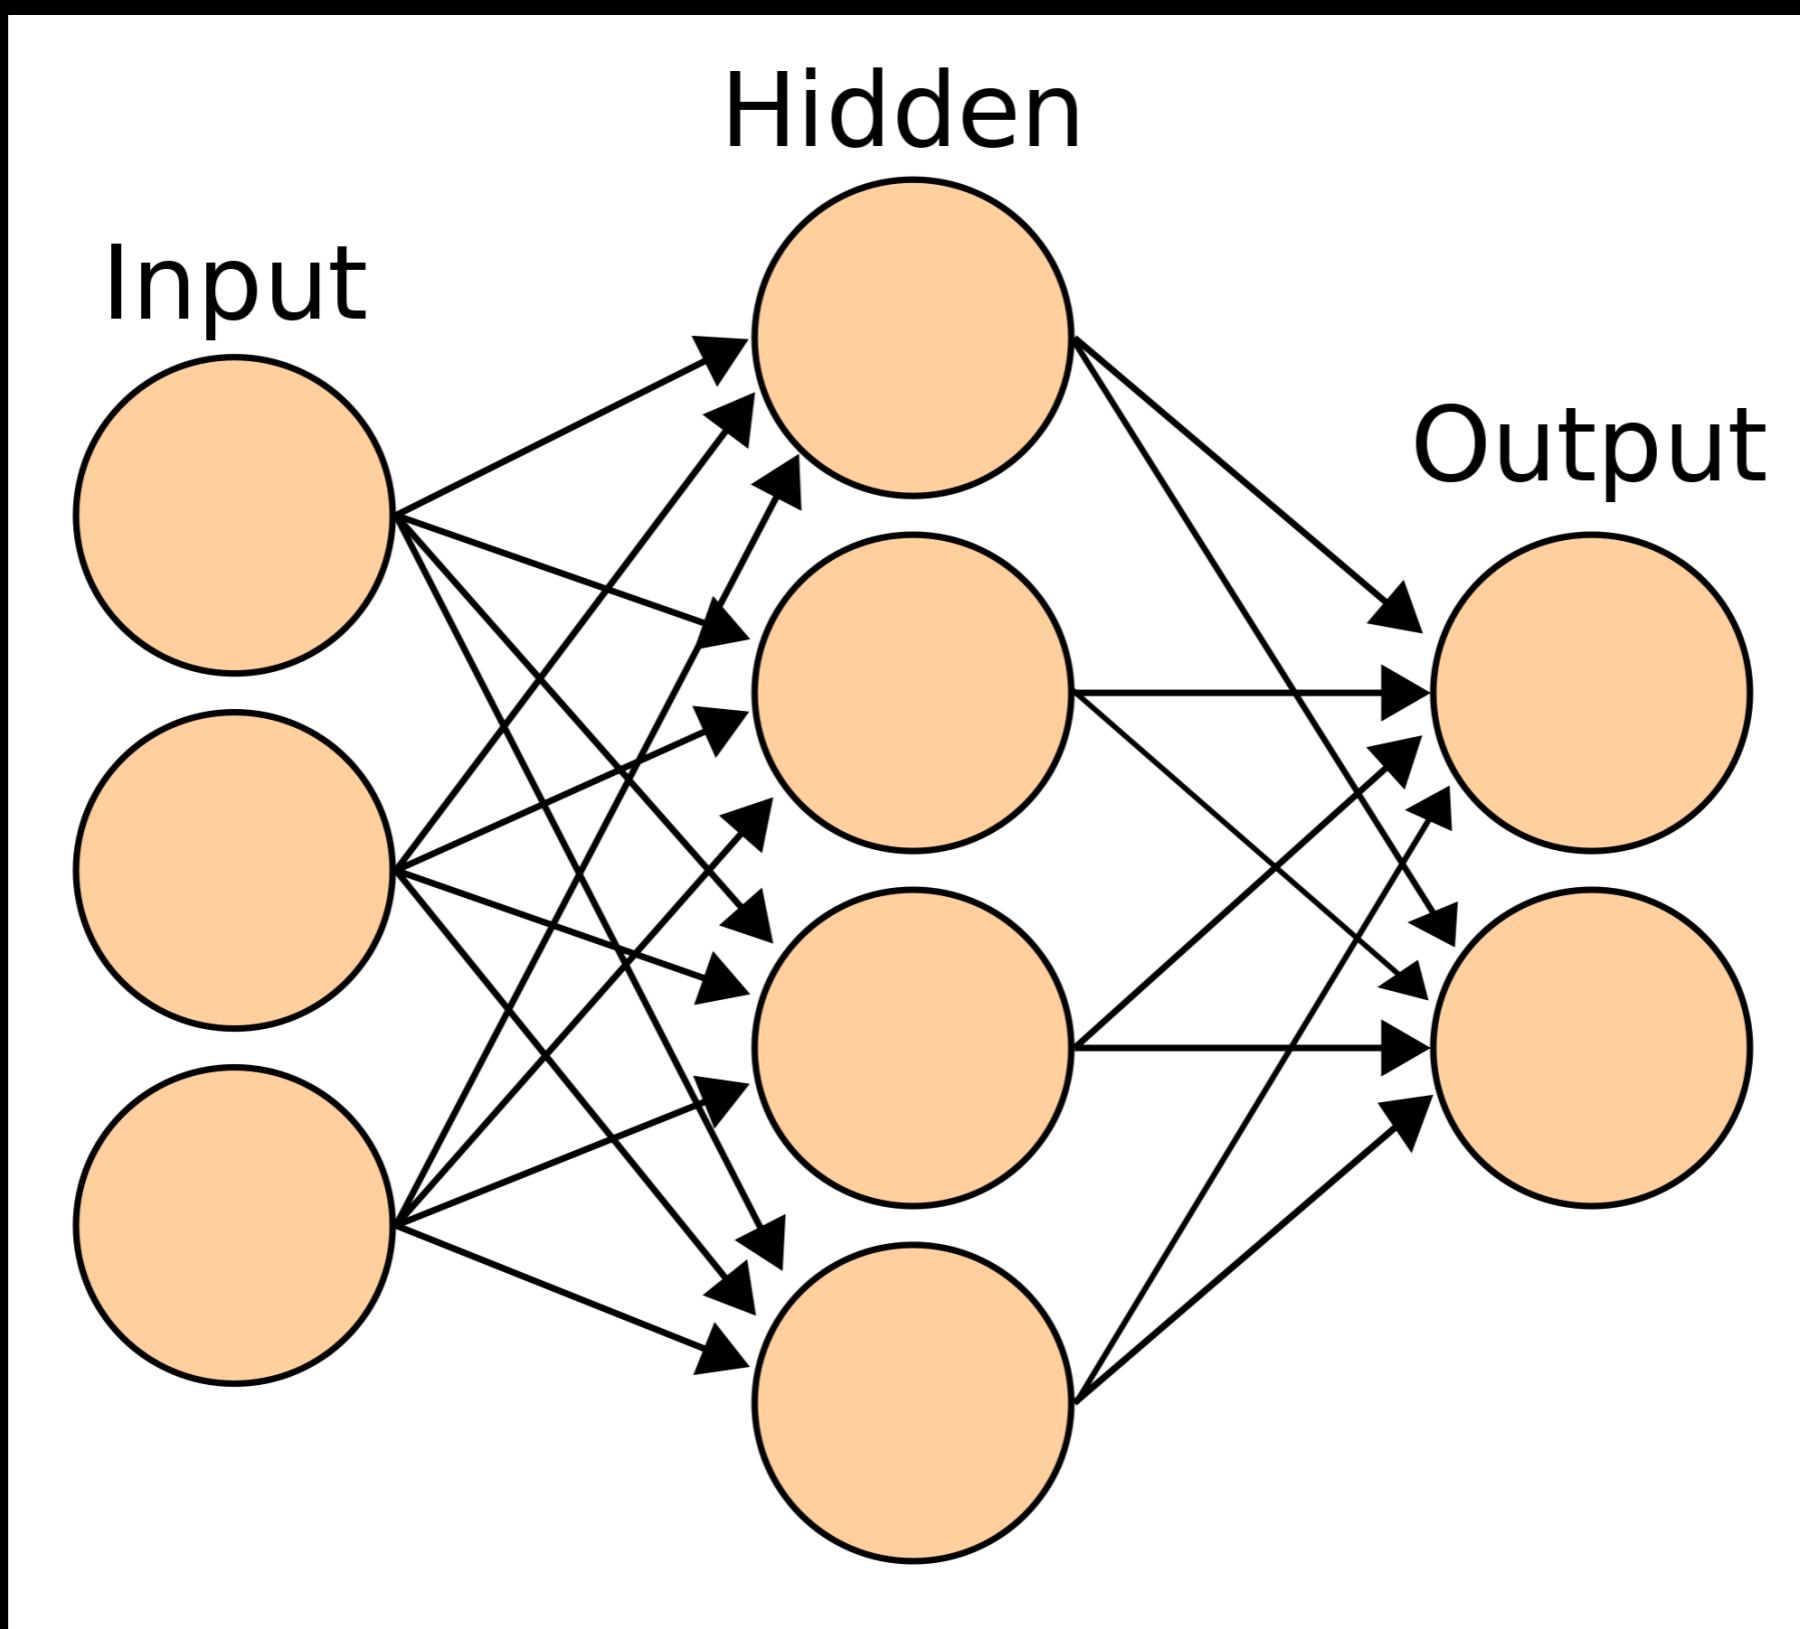
Neural Networks



Neural Networks

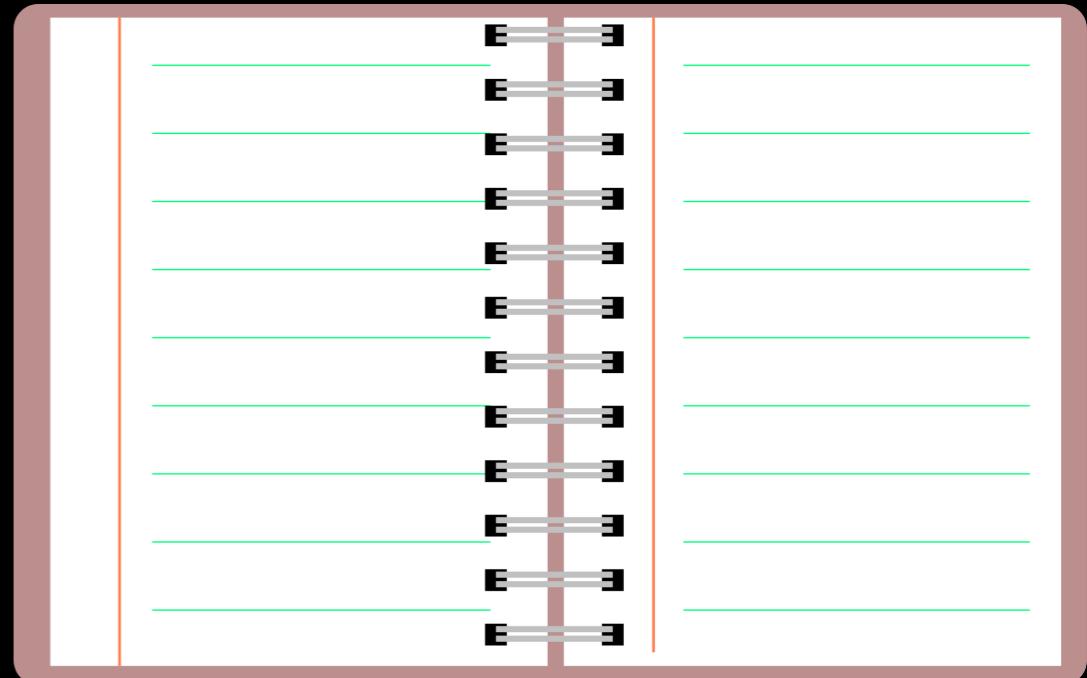


Neural Networks



Agenda

- **Introduction to DeepLearning** ✓
- **Introduction to ApacheSpark**
- **Introduction to DL4J**
- **Introduction to SystemML**
- **Demos**



JVM Process

Driver JVM

Compute Node



Compute Node





Compute Node



Driver JVM

Compute Node

Executor
JVM

Executor
JVM

Executor
JVM



Compute Node



Compute Node



Compute Node



Compute Node

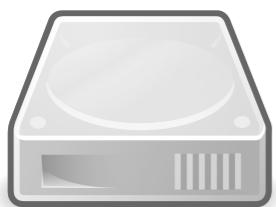


Compute Node





Compute Node



Compute Node



Compute Node



Compute Node



Compute Node



HDFS

Compute Node



Part of File

Compute Node



Part of File

Compute Node



Part of File

Compute Node



Part of File

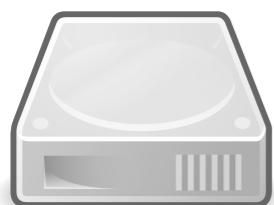
Compute Node



Part of File

HDFS

Compute Node



Part of File

Compute Node



Part of File

Compute Node



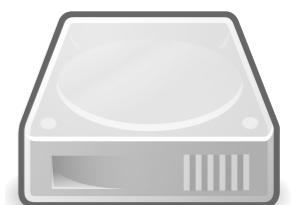
Part of File

Compute Node



Part of File

Compute Node



Part of File

Virtual File

RDD

- “Resilient Distributed Dataset”
- Distributed, immutable collection of data
- Created from HDFS, ObjectStore, Cloudant NoSQL, dashDB SQL, simple files, ...
- In-memory, but spillable to disk
- lazy

Summary

- ApacheSpark programs are implicitly parallel
- Same code can process 1 KB or 1 PB
- RDD central API
- Data and task distribution transparent

Programming Language options on ApacheSpark

Summary

	Scala	Java	R	Python
Spark API support	complete	complete	very limited	limited
ease of use	low	very low	high	very high
Speed	very high	high	very low	low
3rd party libraries	few	few	many	many

Functional Programming basics

Functional Programming (FP)

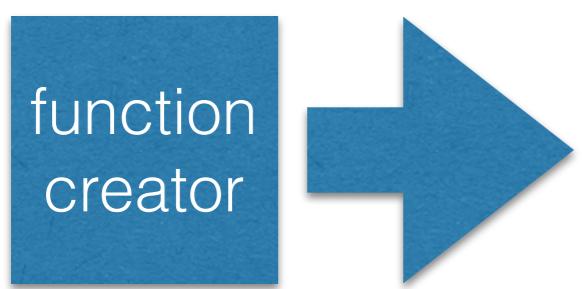
- mathematical concept of “Lambda Calculus
- 1st implementation was LISP in the 50s
- Haskell is part of every computer science curriculum since the 90s
- Scala most recent representative
- Python, R and Java also rudimentary support FP

Idea

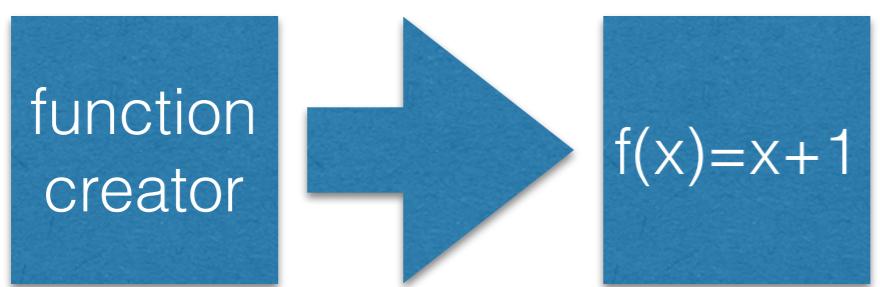
Idea

function
creator

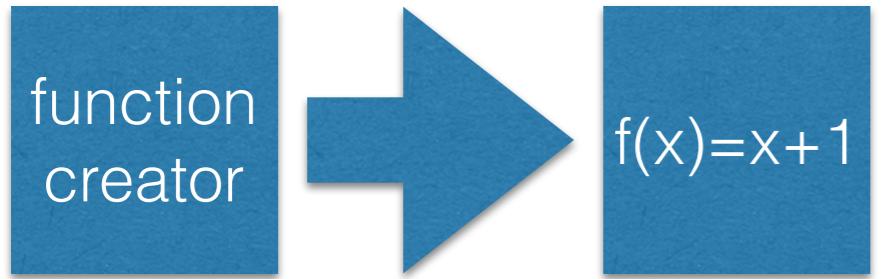
Idea



Idea



Idea



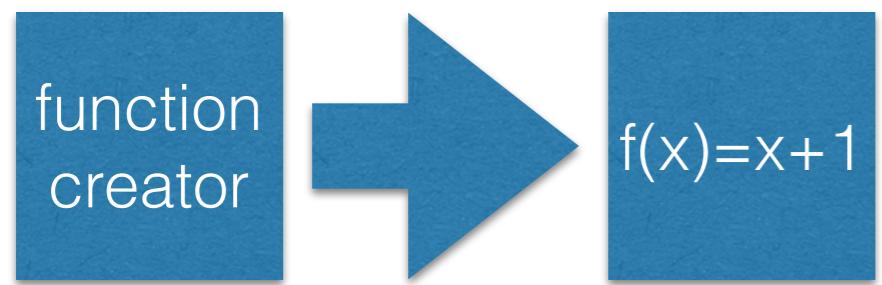
1

2

3

4

Idea



$f(x)=x+1$

1

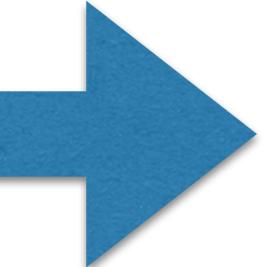
2

3

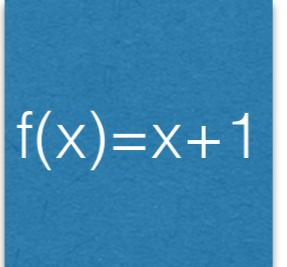
4

Idea

function
creator



$f(x)=x+1$

apply( , )

1

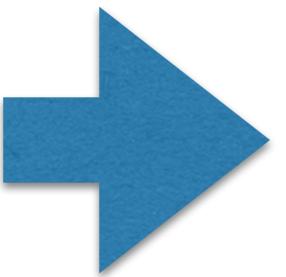
2

3

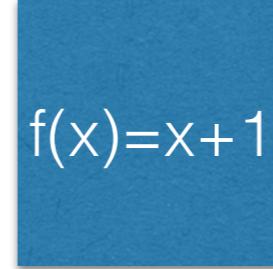
4

Idea

function
creator



$f(x) = x + 1$

apply( , ) = 

$f(x) = x + 1$

1

2

2

3

3

4

4

5

Parallelisation

$$f(x) = x + 1$$

Parallelisation

$f(x)=x+1$



Parallelisation

1 2 3 4 5 6 7 8 9

Parallelisation

1 2 3

4 5 6

7 8 9

$$f(x)=x+1$$

$$f(x)=x+1$$

$$f(x)=x+1$$

Parallelisation

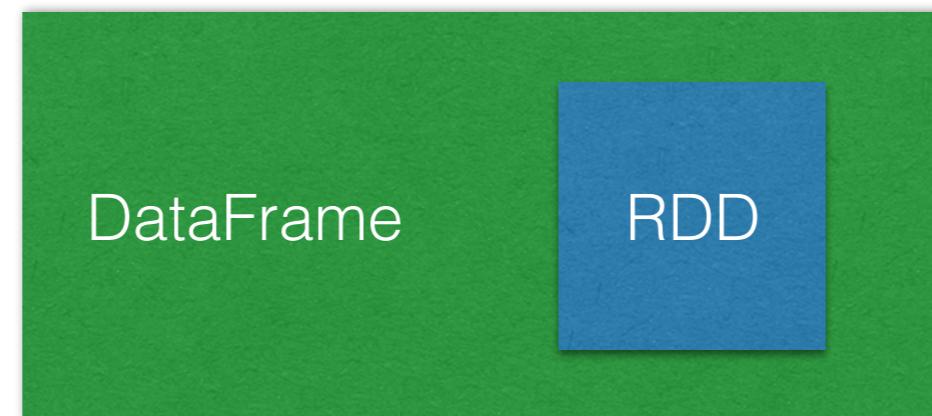
1	2	3	4	5	6	7	8	9
			$f(x)=x+1$					
2	3	4	5	6	7	8	9	10

ApacheSparkSQL

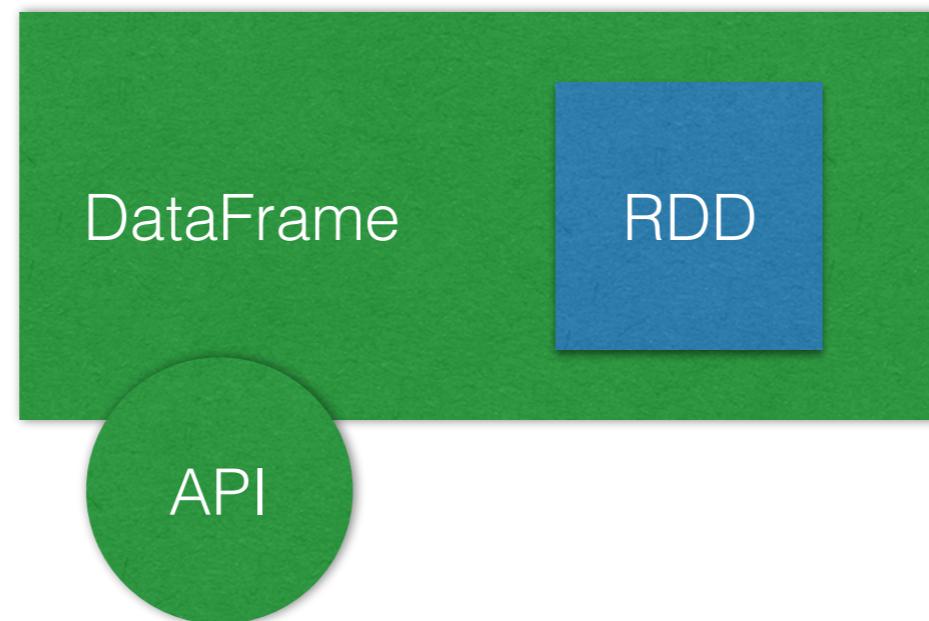
ApacheSparkSQL



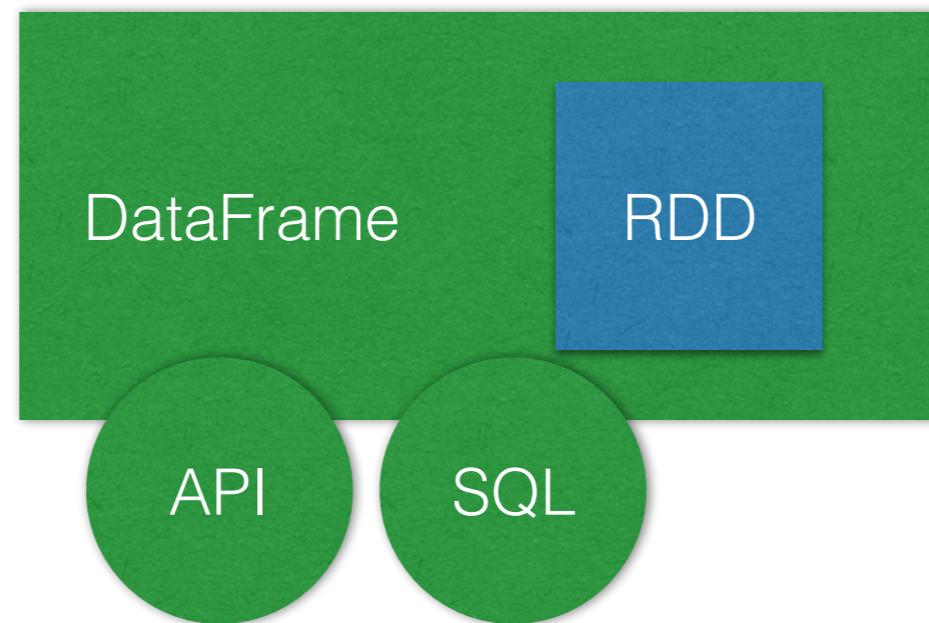
ApacheSparkSQL



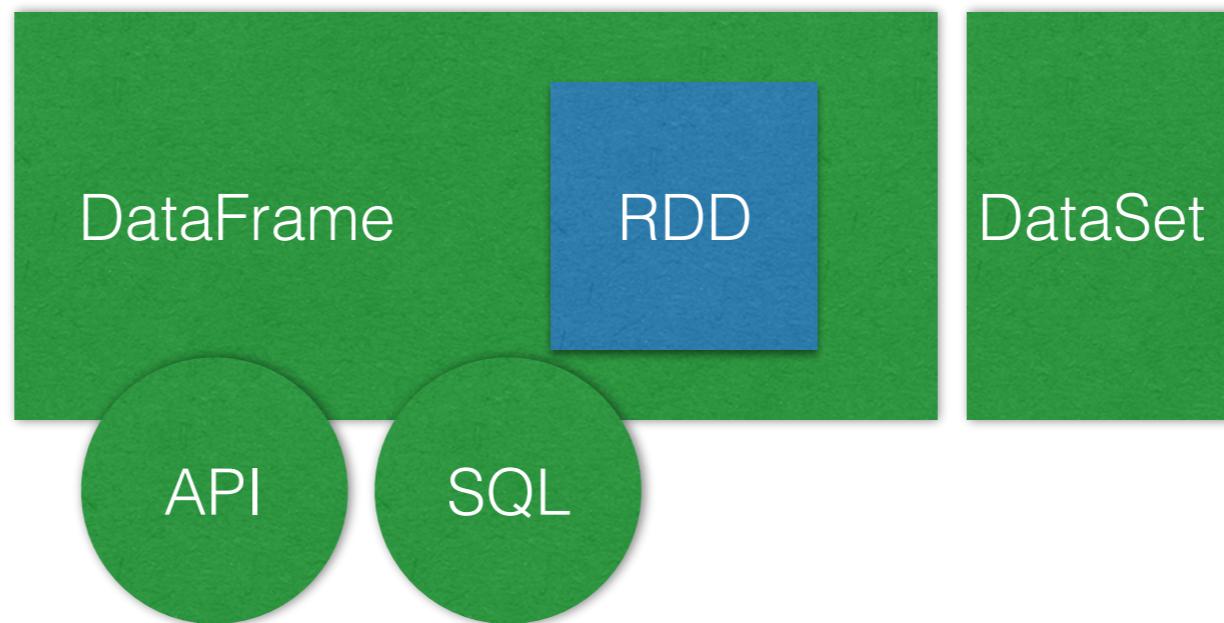
ApacheSparkSQL



ApacheSparkSQL



ApacheSparkSQL



Schemas

- RDDs are schema less (schema on read)
- DataFrames have a schema
 - lazy, inferred
 - explicitly defined

The “Catalyst”

- Creates “logical execution plan” (LEP) from SQL
- Optimises LEP to “physical execution plans” (PEPs)
- based on statistics chooses best PEP to execute
- similar to cost based optimisers in RDBMs

Project Tungsten

- Java Virtual Machine (JVM) is an art piece
- General purpose byte code execution engine
- JVM objects & Garbage Collection (GC) overhead
 - 4 byte string is 48 byte on the JVM
 - GC optimises on object life time estimation
 - Spark knows this better than JVM

Project Tungsten

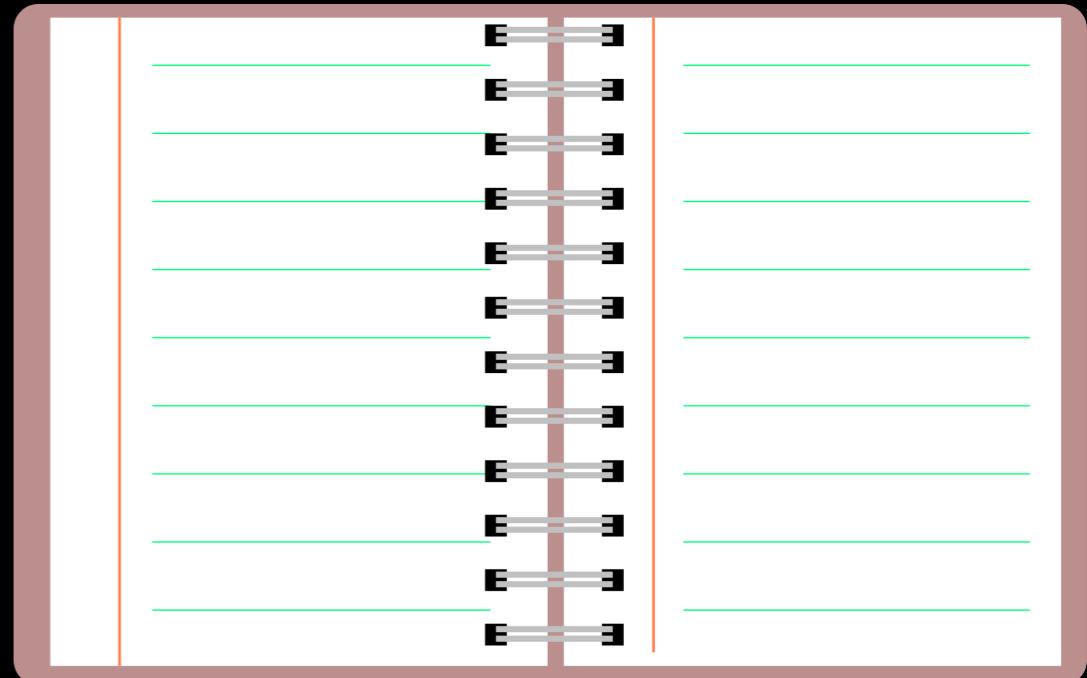
- L1/L2/L3 Cache friendly data structures
- Code generation to remove
 - boxing of primitive types
 - polymorphic function dispatching

Summary

- ApacheSpark supports SQL via data frame API
- Internally still RDDs are used
- Makes writing ApacheSpark jobs easier
- Performance benefits through Catalyst & Tungsten

Agenda

- **Introduction to DeepLearning**
- **Introduction to ApacheSpark**
- **Introduction to DL4J**
- **Introduction to SystemML**
- **Demos**



DeepLearning4J

Components

- DeepLearning4J
Enterprise Grade DeepLearning Library
- DataVec
CSV/Audio/Video/Image/... => Vector
- ND4J / ND4S (NumPy for the JVM)

ND4J

- Tensor support (Linear Buffer + Stride)
- Multiple implementations, one interface
 - vectorized c++ code (JavaCPP), off-heap data storage, BLAS (OpenBLAS, Intel MKL, cuBLAS)
 - GPU (CUDA 7.5)

turn on GPU

```
<name>DeepLearning4j Examples Parent</name>
<description>Examples of training different data sets</description>
<properties>
    <nd4j.backend>nd4j-native-platform</nd4j.backend>
```

```
<nd4j.backend>nd4j-cuda-7.5-platform</nd4j.backend>
```

DL4J parallelisation

- TensorFlow on ApacheSpark =>
 - Scoring
 - Multi-model hyper-parameter tuning
 - Parallel training since V r0.8
- DeepLearning4J =>
 - Scoring, Multi-model hyper-parameter tuning
 - Parallel training
 - “Jeff Dean style parameter averaging”

“Code local vs spark”

`new MultiLayerNetwork(conf);`

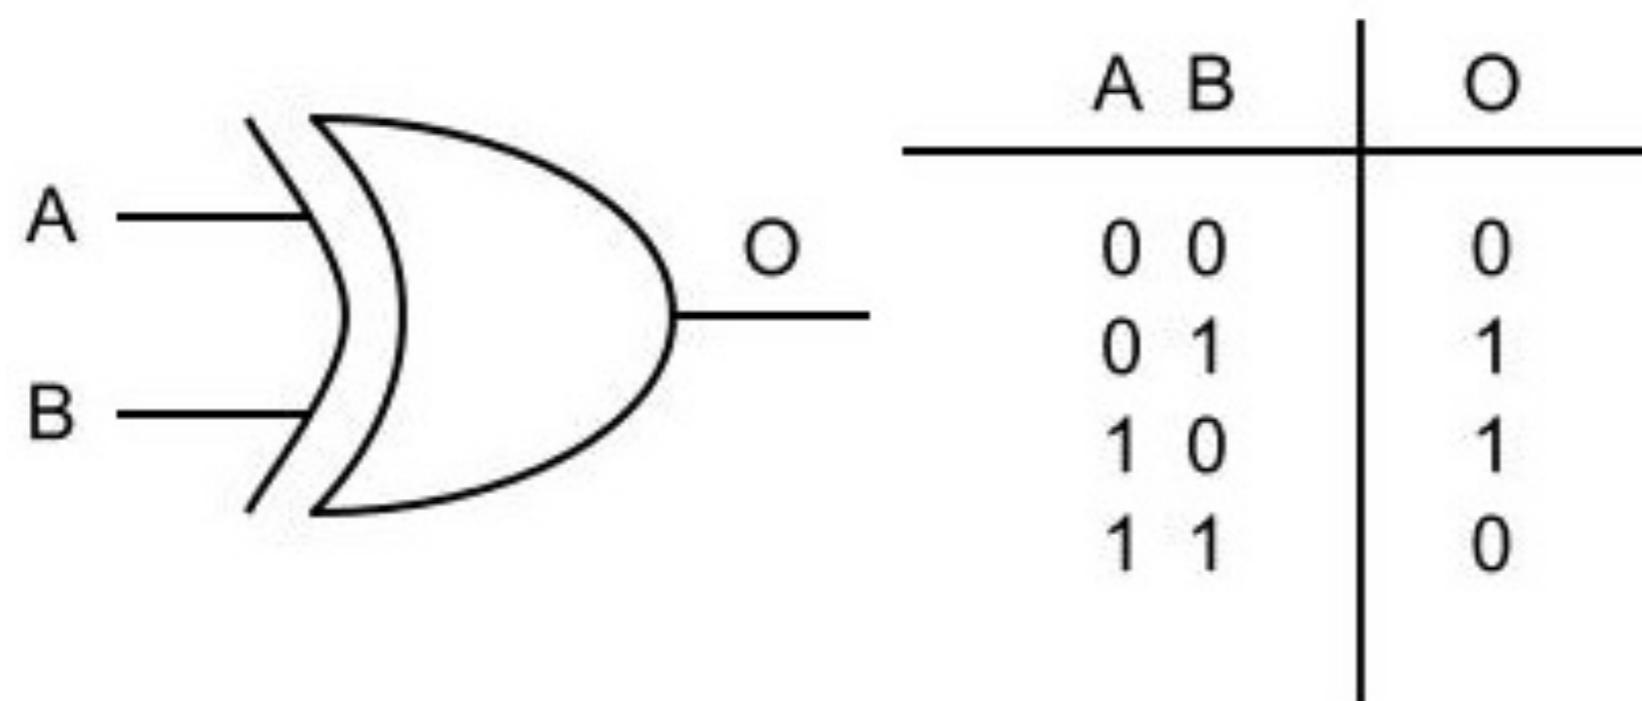
vs.

`new SparkDl4jMultiLayer(sc, conf, tm);`

```
TrainingMaster tm = new ParameterAveragingTrainingMaster.Builder(int dataSet0bjectSize)
    ... (your configuration here)
    .build();
```

XOR

<https://github.com/romeokienzler/dl4j-examples/blob/master/dl4j-examples-scala/src/main/scala/org/deeplearning4j/examples/feedforward/xor/XOrExampleScala.scala>



Agenda

- **Introduction to DeepLearning**
- **Introduction to ApacheSpark**
- **Introduction to DL4J**
- **Introduction to SystemML**
- **Demos**



Why another lib?

- Custom machine learning algorithms
- Declarative ML
- Transparent distribution on data-parallel framework
 - Scale-up
 - Scale-out
- Cost-based optimiser generates low level execution plans

Why on Spark?

- Unification of SQL, Graph, Stream, ML
- Common RDD structure
- General DAG execution engine
 - lazy evaluation
 - distributed in-memory caching

2007-2008: Multiple projects at IBM Research – Almaden involving machine learning on Hadoop.

2009: We form a dedicated team for scalable ML

2009-2010: Through engagements with customers, we observe how data scientists create **ML solutions**.

2007

2008

2009

2010

Research

2011

2012

2013

2014

June 2015: IBM Announces open-source SystemML

November 2015: SystemML enters Apache incubation

June 2016: Second Apache release (0.10)

September 2015: Code available on Github

February 2016: First release (0.9) of Apache SystemML

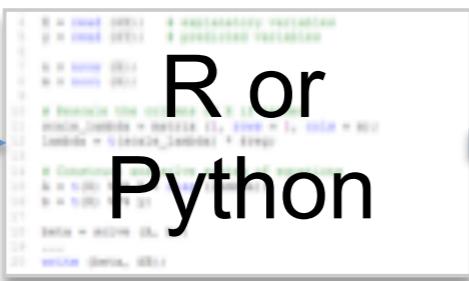
2015

2016

SystemML at IBM Watson Health

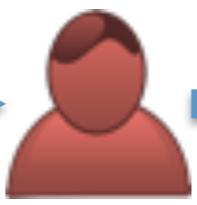
Moved from Hadoop MapReduce to Spark
SystemML supports both frameworks
Exact same code
300X faster on 1/40th as many nodes

Data Scientist



R or Python

Systems Programmer



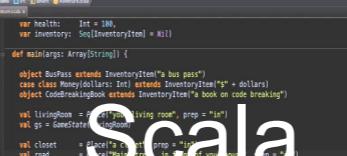
```
class Money(dollars: Int) extends InventoryItem(s"$ - dollars")
class Book extends InventoryItem("a book on code breaking")
class BreakingBook extends InventoryItem("a book on code breaking")
```

LivingRoom = Place("the living room", prep = "in")
= GameState("in living room")

Poset = Place("a poset", prep = "in")
= GameState("in poset")

ad = Place("a cat", prep = "in")
= GameState("in cat")

ad = Place("a book", prep = "on")
= GameState("on book")



The screenshot shows a Java-based IDE interface with multiple tabs open. The main code editor window displays Scala code for a game. The code defines a `Game` class with methods like `start`, `getRoom`, and `getInventory`. It also includes a `Player` class with methods for moving between rooms and interacting with objects. A `Trans` trait is defined for transitioning between rooms. The code uses `Seq[InventoryItem]` for inventories and `Map[Place, Place]` for room transitions. The IDE's status bar at the bottom shows the file path as `Adventure.scala (local) [adventure] - [devkit local]`.

```
object Game {
    def start(): Unit = {
        val player = Player("Alice")
        val room = getRoom("Living Room")
        player.enterRoom(room)
        while (true) {
            val action = readLine("What would you like to do? ")
            if (action == "quit") {
                break
            }
            val target = readLine("Enter target: ")
            if (target == "quit") {
                break
            }
            player.interact(target)
        }
    }

    def getRoom(name: String): Room = {
        val room = Room(name)
        room.setExit("North", Room("Bedroom"))
        room.setExit("South", Room("Bathroom"))
        room.setExit("West", Room("Kitchen"))
        room.setExit("East", Room("Living Room"))
        room.setInventory(Seq[InventoryItem](new Key("House Key")))
        room
    }

    def getInventory(itemName: String): InventoryItem = {
        val item = new InventoryItem(itemName)
        item.setDesc("A shiny metal key used to unlock doors")
        item.setUse("unlock door")
        item
    }

    trait Trans {
        def transitionFrom(place: Place): Map[Place, Place]
    }

    class Player(name: String) {
        var health: Int = 100
        var inventory: Seq[InventoryItem] = Nil

        def move(direction: String): Unit = {
            val targetRoom = getRoom(getExitName(direction))
            if (targetRoom != null) {
                transitionFrom(targetRoom).foreach { case (place, trans) =>
                    if (place == this) {
                        trans.transitionFrom(this)
                    }
                }
                transitionFrom(this).foreach { case (place, trans) =>
                    if (place == targetRoom) {
                        trans.transitionFrom(targetRoom)
                    }
                }
                this.exitRoom()
                targetRoom.enterRoom(this)
            } else {
                println("There is no exit in that direction")
            }
        }

        def interact(itemName: String): Unit = {
            val item = getInventory(itemName)
            if (item != null) {
                if (item.getUse != null) {
                    println(item.getUse)
                } else {
                    println("The item has no use")
                }
            } else {
                println("Item not found")
            }
        }

        def exitRoom(): Unit = {
            val exits = transitionFrom(this)
            if (exits.size > 0) {
                val exitName = readLine("Which way would you like to go? ")
                if (exits.contains(exitName)) {
                    move(exitName)
                } else {
                    println("That exit does not exist")
                }
            } else {
                println("You have nowhere to go")
            }
        }

        def enterRoom(room: Room): Unit = {
            room.setPlayer(this)
            println(s"You entered the ${room.getName} room")
            room.printInventory()
        }

        def printInventory(): Unit = {
            inventory.foreach { item =>
                item.printDetails()
            }
        }
    }
}
```

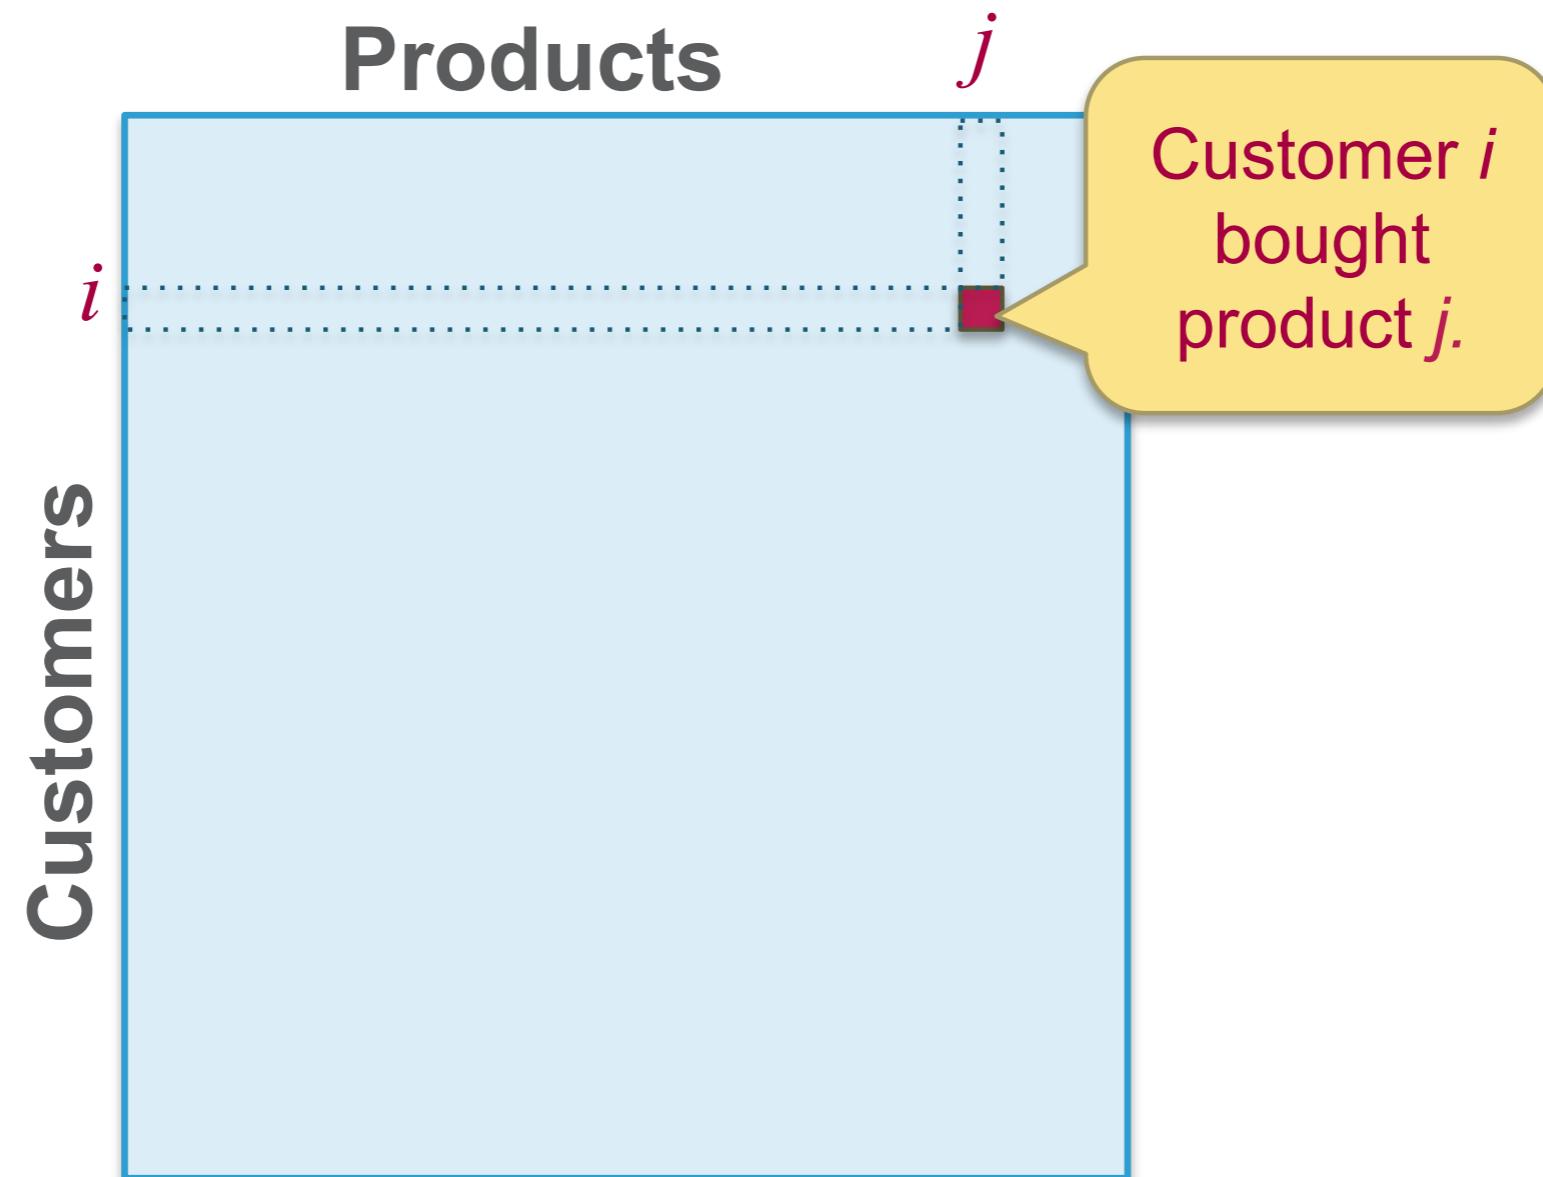
Results

	AAPL	30/05/2008	182.75	188.75
24	AAPL	06/06/2008	188.6	185.64
25	AAPL	13/06/2008	184.79	172.37
26	AAPL	20/06/2008	171.3	175.27
27	AAPL	27/06/2008	174.74	170.09
28	AAPL	03/07/2008	170.19	170.12
29	AAPL	10/07/2008	171.16	172.58
30	AAPL	17/07/2008	175.74	165.15
31	AAPL	25/07/2008	166.9	162.12
32	AAPL	01/08/2008	162.34	156.66
33	AAPL	08/08/2008	158.6	169.55
34	AAPL	15/08/2008	170.07	175.74
35	AAPL	22/08/2008	175.57	176.79
36	AAPL	29/08/2008	176.15	169.53

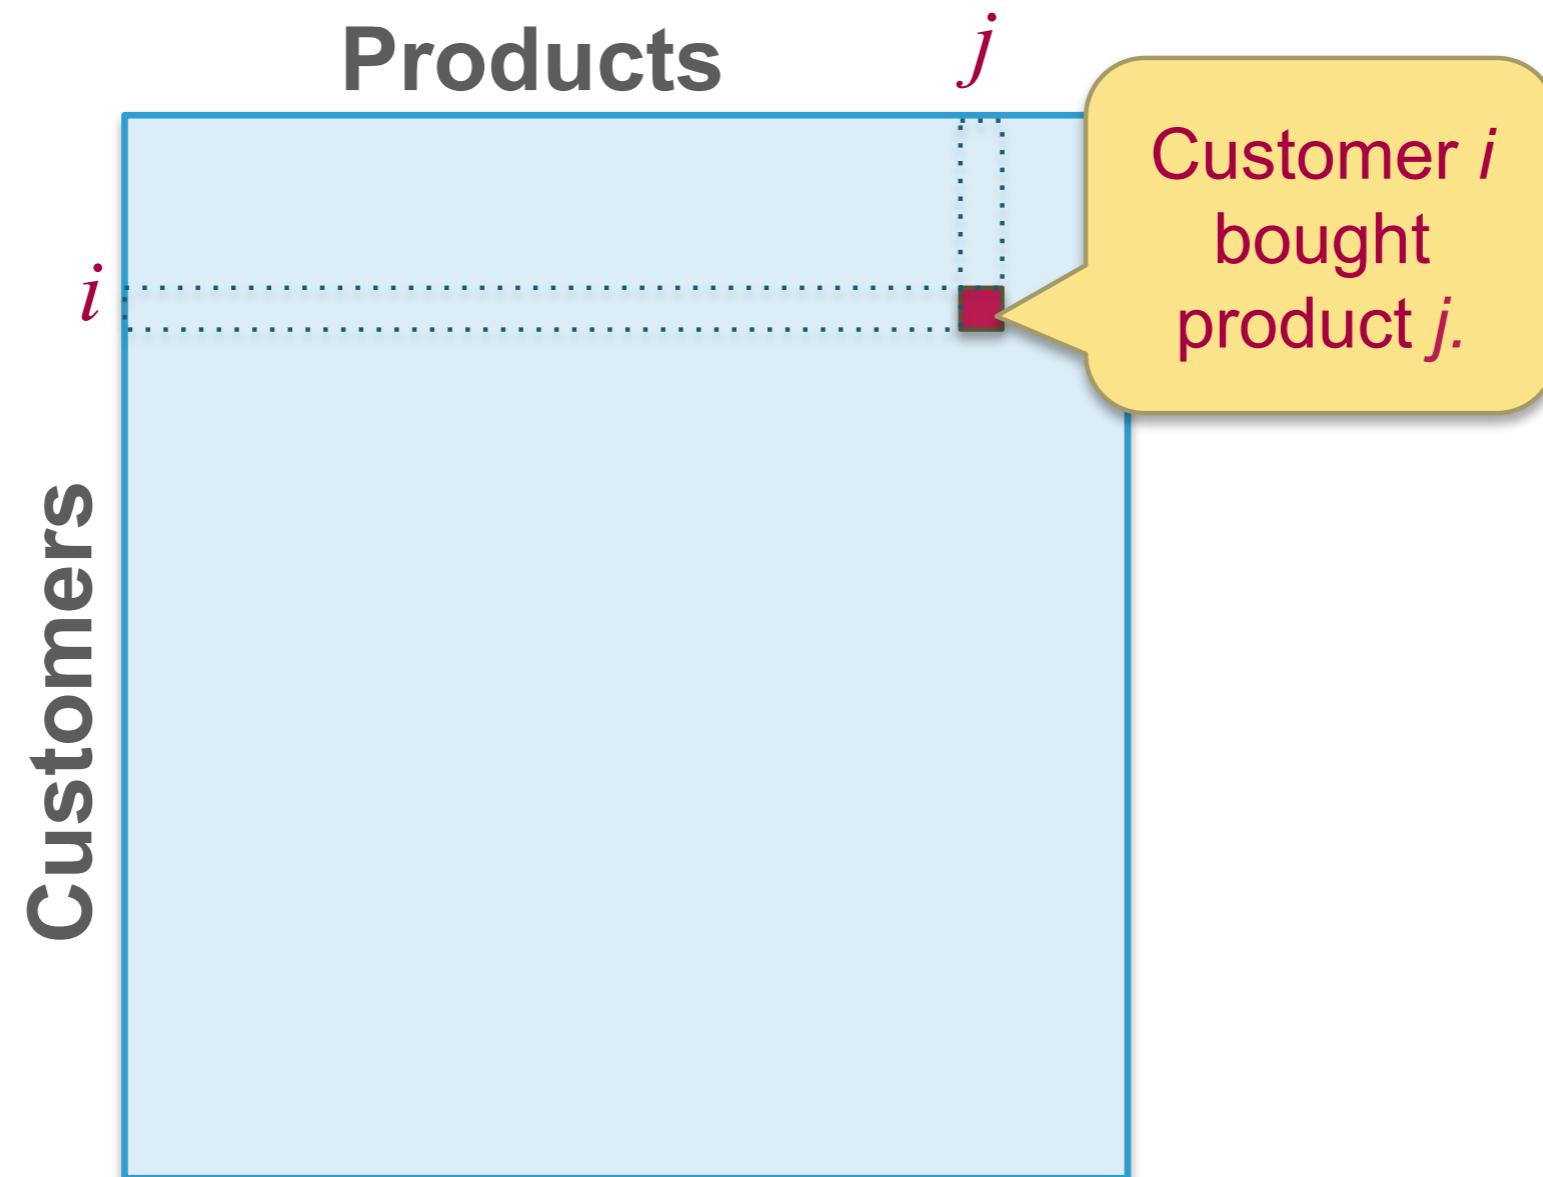


Spark

Alternating Least Squares

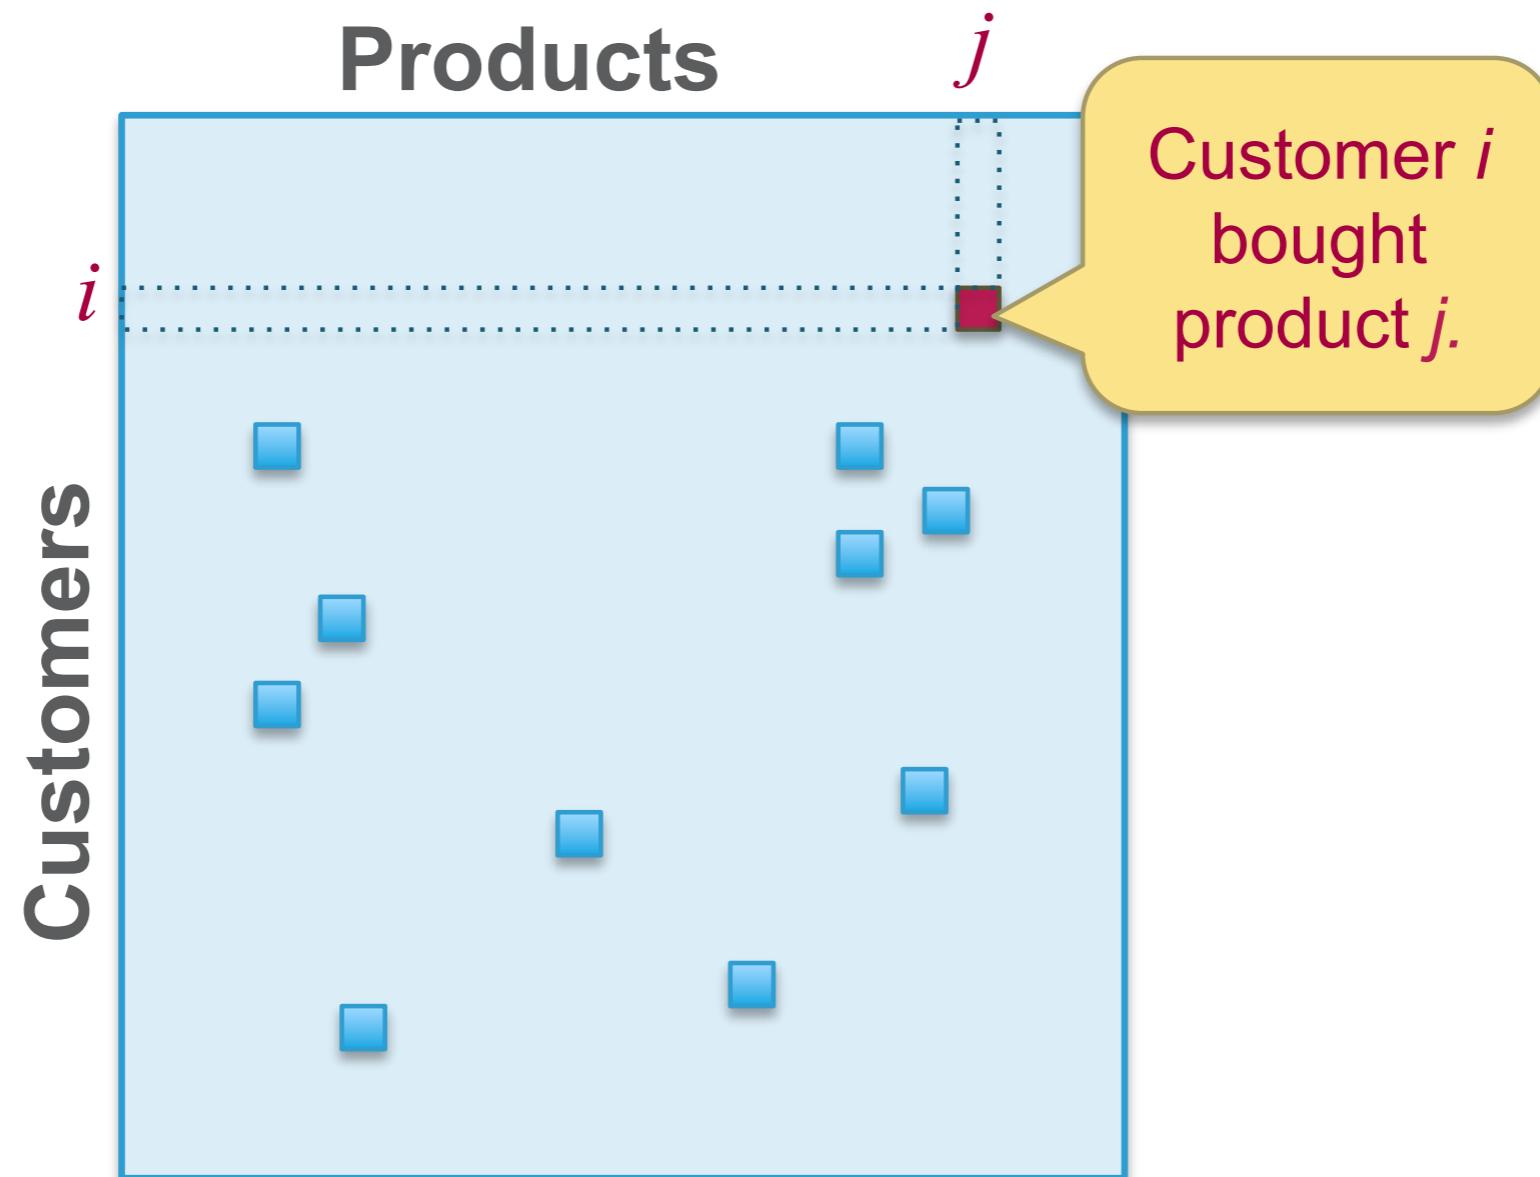


Alternating Least Squares



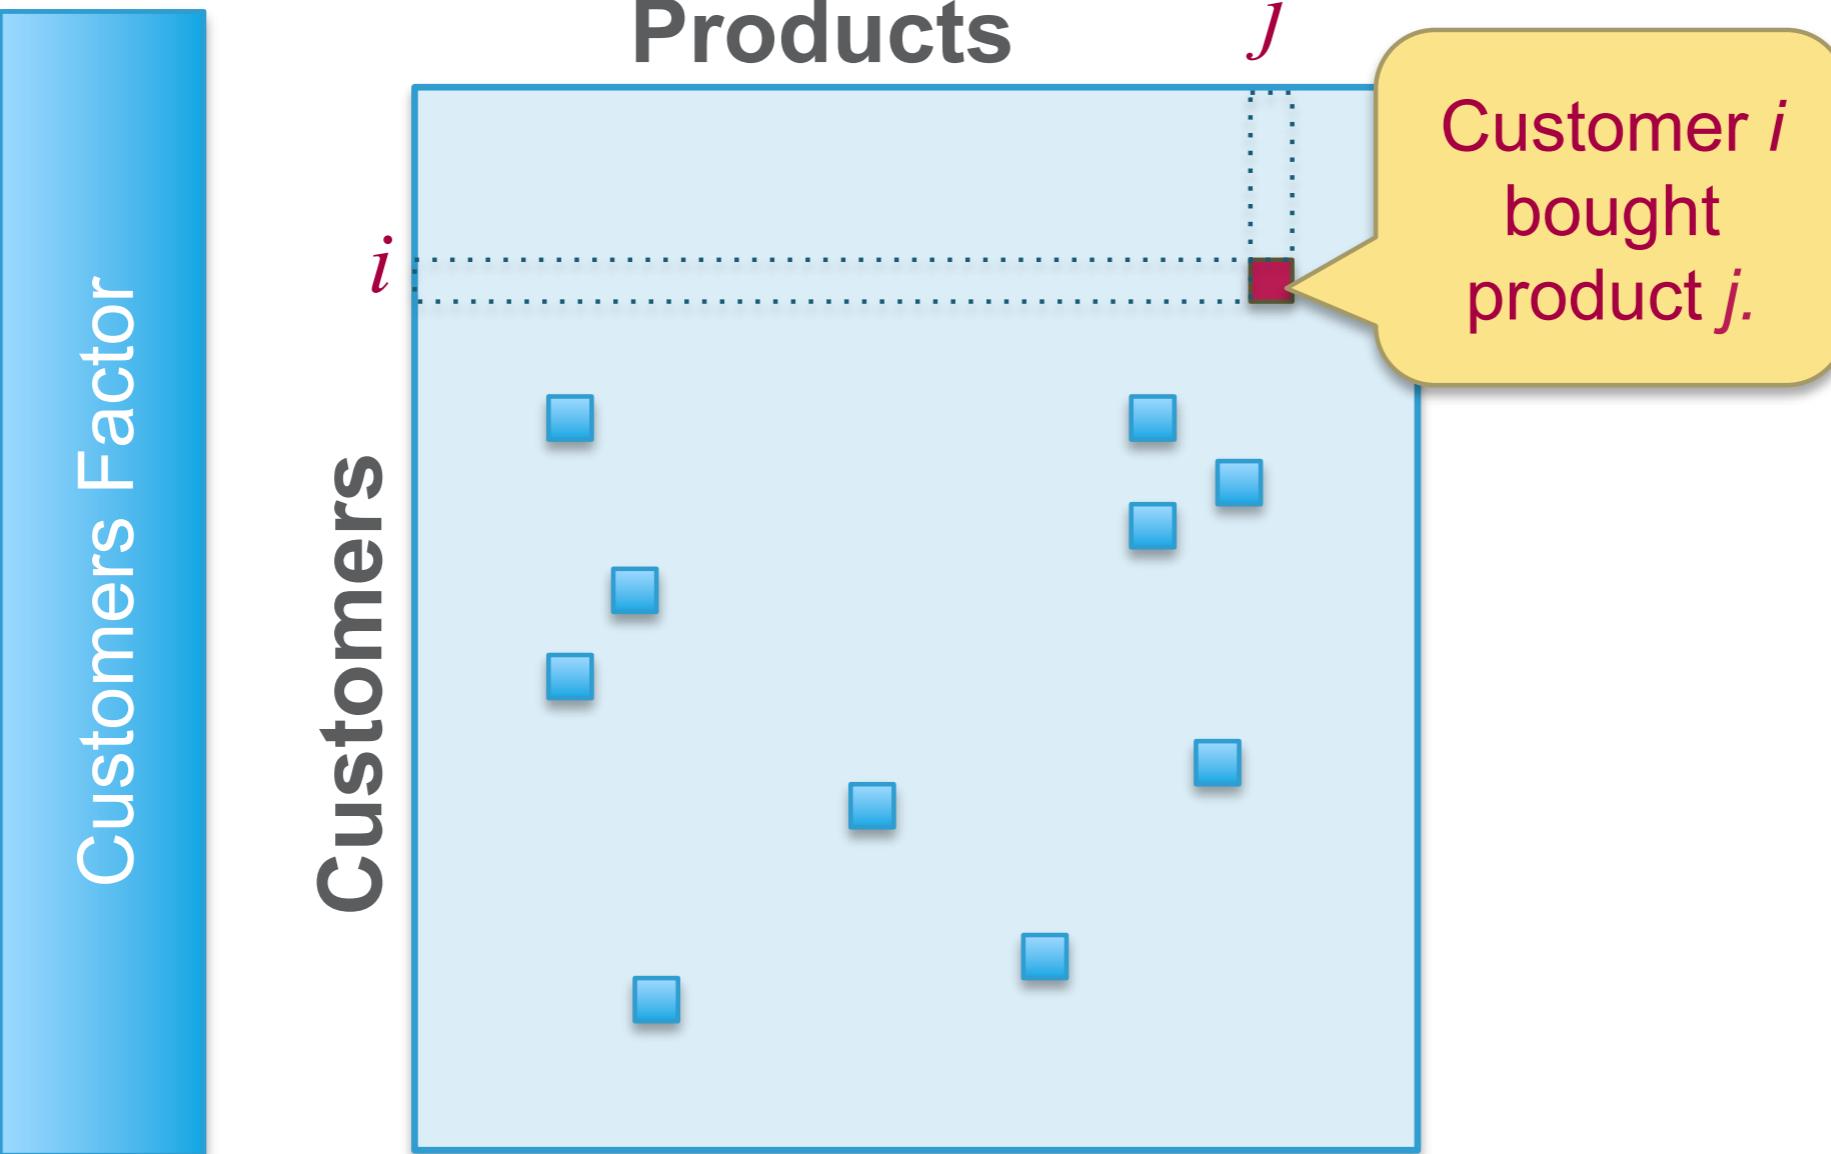
$$r_{ui}$$

Alternating Least Squares

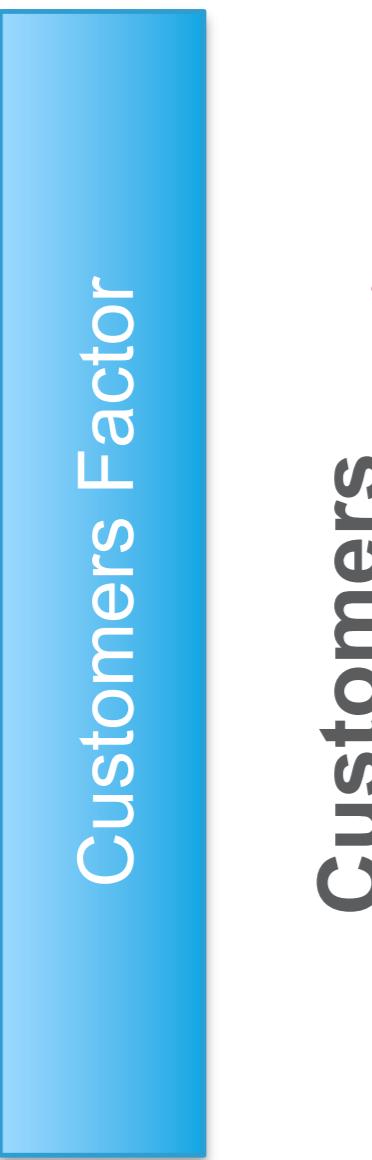


$$r_{ui}$$

p_u

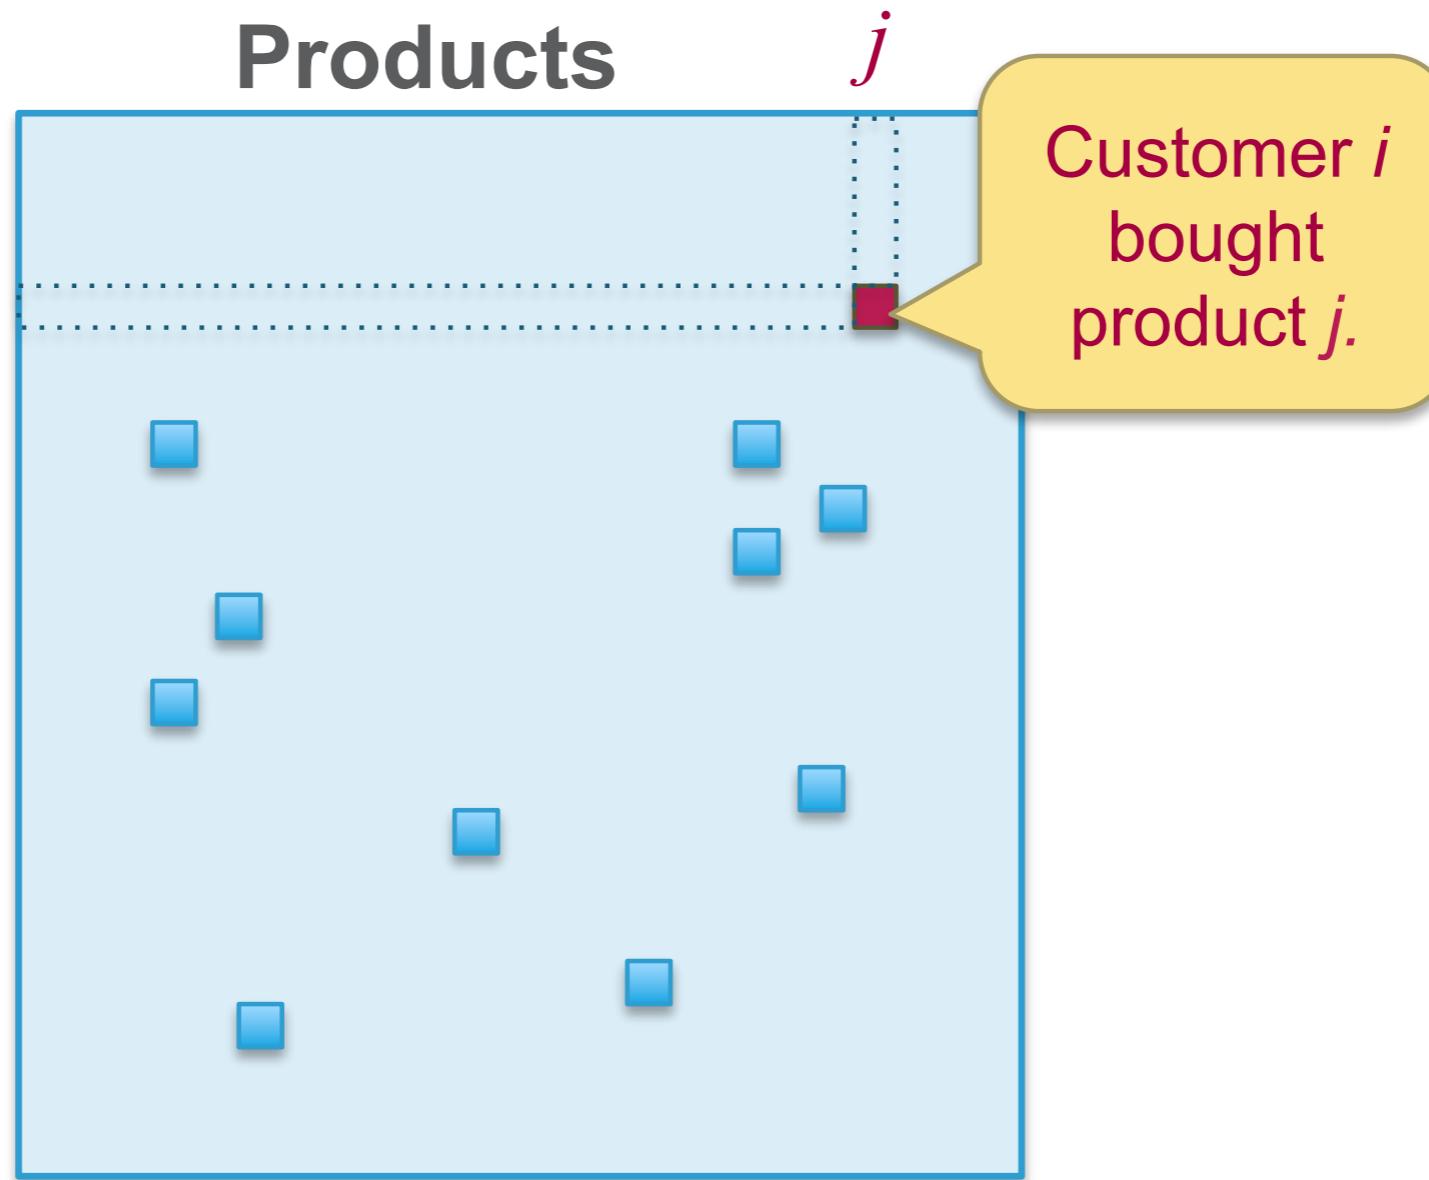


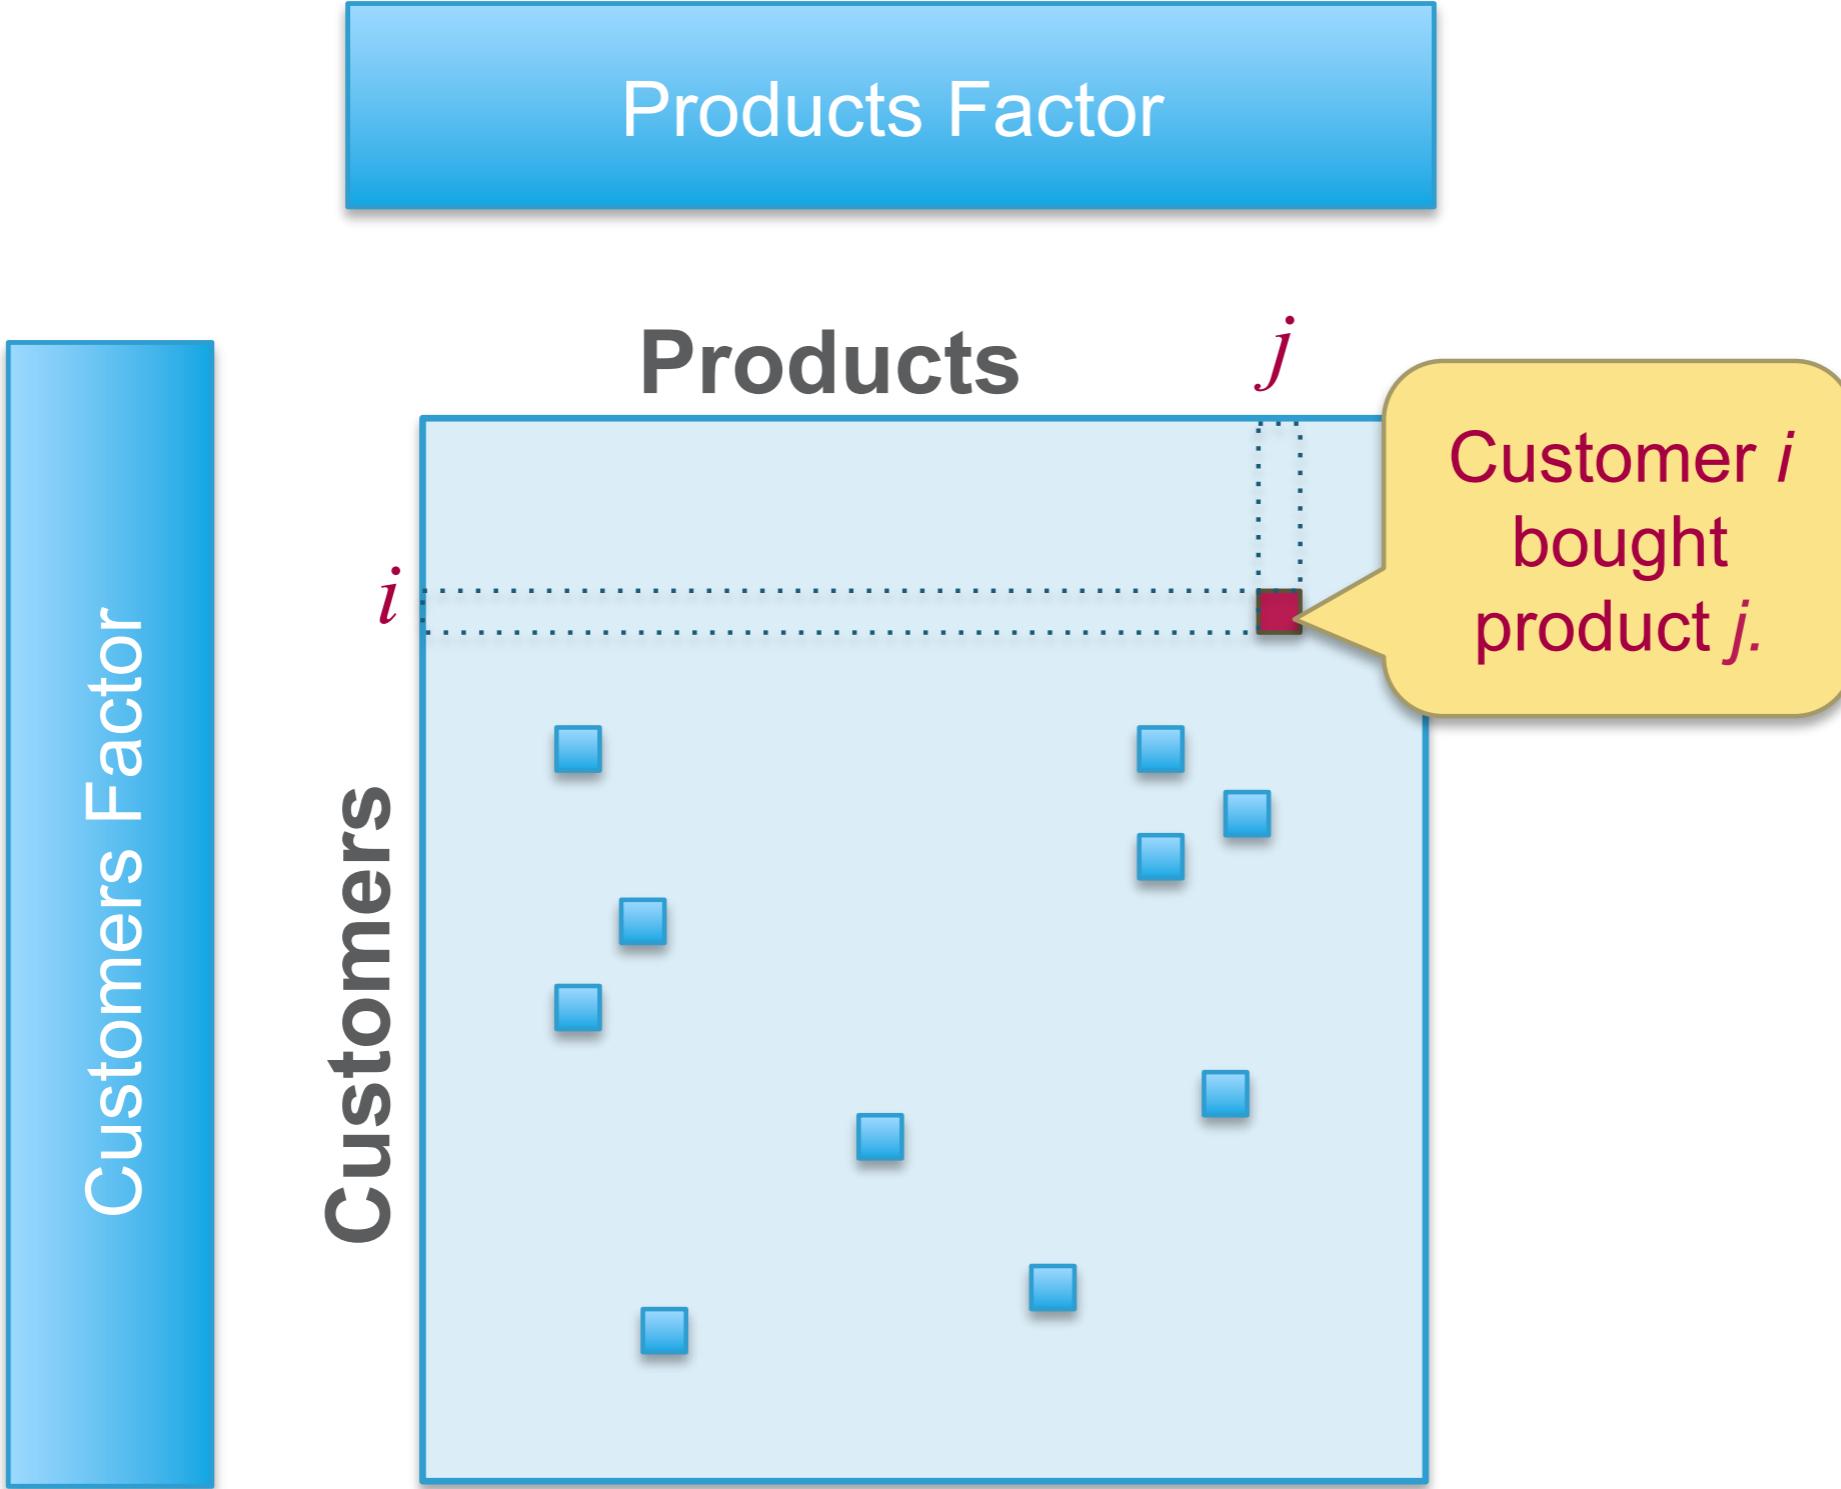
p_u



Products Factor

q_i





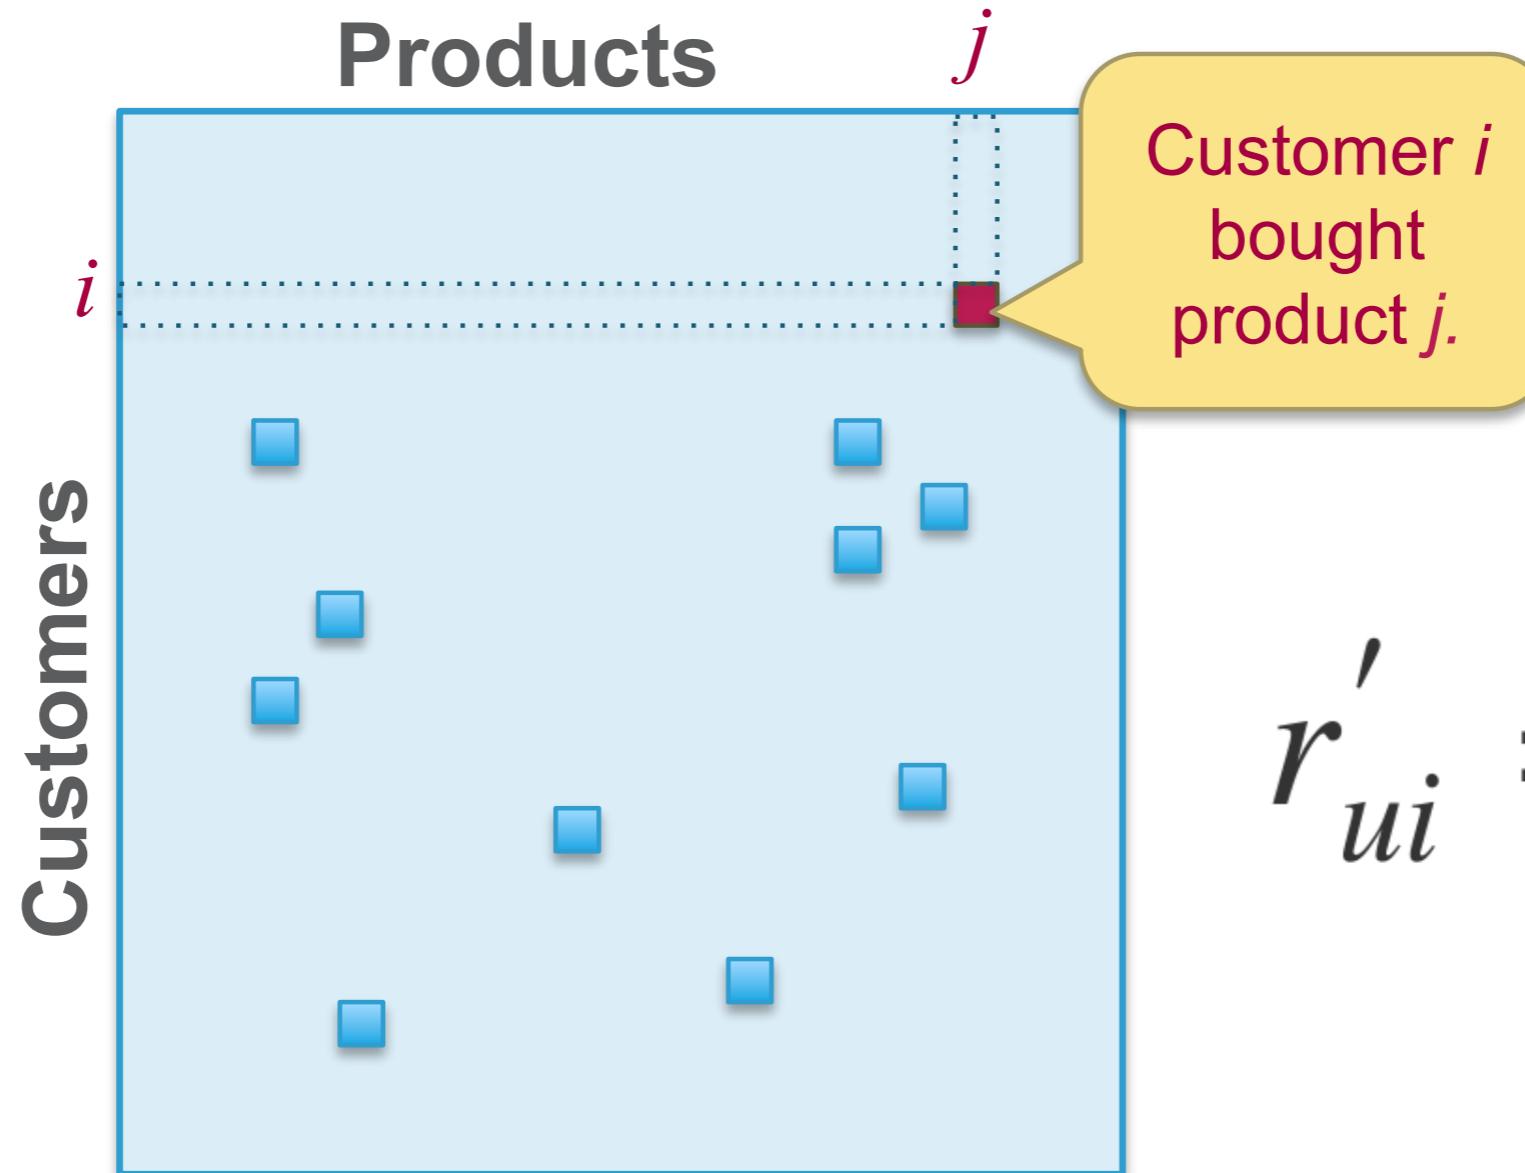
$$\min_{q,p} \sum_{u,i} (r_{ui} - p_u^T q_i)^2$$

Multiply these two factors to produce a less-sparse matrix.



Products Factor

Customers Factor



$$r'_{ui} = p_u^T q_i$$

$$\min_{q,p} \sum_{u,i} (r_{ui} - p_u^T q_i)^2$$

Multiply these two factors to produce a less-sparse matrix.



Products Factor



Products

Customers

j

i

Customer *i* bought product *j*.

New nonzero values become product suggestions.

$$r'_{ui} = p_u^T q_i$$

```
val model = ALS.train(ratings, rank, numIterations, 0.01)
```

```

U = rand(nrow(X), r, min = -1.0, max = 1.0);
V = rand(r, ncol(X), min = -1.0, max = 1.0);
while(i < mi) {
  i = i + 1; ii = 1;
  if (is_U)
    G = (W * (U %*% V - X)) %*% t(V) + lambda * U;
  else
    G = t(U) %*% (W * (U %*% V - X)) + lambda * V;
  norm_G2 = sum(G ^ 2); norm_R2 = norm_G2;
  R = -G; S = R;
  while(norm_R2 > 10E-9 * norm_G2 & ii <= mii) {
    if (is_U) {
      HS = (W * (S %*% V)) %*% t(V) + lambda * S;
      alpha = norm_R2 / sum(S * HS);
      U = U + alpha * S;
    } else {
      HS = t(U) %*% (W * (U %*% S)) + lambda * S;
      alpha = norm_R2 / sum(S * HS);
      V = V + alpha * S;
    }
    R = R - alpha * HS;
    old_norm_R2 = norm_R2; norm_R2 = sum(R ^ 2);
    S = R + (norm_R2 / old_norm_R2) * S;
    ii = ii + 1;
  }
  is_U = ! is_U;
}

```

```

U = rand(nrow(X), r, min = -1.0, max = 1.0);
V = rand(r, ncol(X), min = -1.0, max = 1.0);
while(i < mi) {
  i = i + 1; ii = 1;
  if (is_U)
    G = (W * (U %*% V - X)) %*% t(V) + lambda * U;
  else
    G = t(U) %*% (W * (U %*% V - X)) + lambda * V;
  norm_G2 = sum(G ^ 2); norm_R2 = norm_G2;
  R = -G; S = R;
  while(norm_R2 > 10E-9 * norm_G2 & ii <= mii) {
    if (is_U) {
      HS = (W * (S %*% V)) %*% t(V) + lambda * S;
      alpha = norm_R2 / sum(S * HS);
      U = U + alpha * S;
    } else {
      HS = t(U) %*% (W * (U %*% S)) + lambda * S;
      alpha = norm_R2 / sum(S * HS);
      V = V + alpha * S;
    }
    R = R - alpha * HS;
    old_norm_R2 = norm_R2; norm_R2 = sum(R ^ 2);
    S = R + (norm_R2 / old_norm_R2) * S;
    ii = ii + 1;
  }
  is_U = ! is_U;
}

```

```

U = rand(nrow(X), r, min = -1.0, max = 1.0);
V = rand(r, ncol(X), min = -1.0, max = 1.0);
while(i < mi) {
  i = i + 1; ii = 1;
  if (is_U)
    G = (W * (U %*% V - X)) %*% t(V) + lambda * U;
  else
    G = t(U) %*% (W * (U %*% V - X)) + lambda * V;
  norm_G2 = sum(G ^ 2); norm_R2 = norm_G2;
  R = -G; S = R;
  while(norm_R2 > 10E-9 * norm_G2 & ii <= mii) {
    if (is_U) {
      HS = (W * (S %*% V)) %*% t(V) + lambda * S;
      alpha = norm_R2 / sum(S * HS);
      U = U + alpha * S;
    } else {
      HS = t(U) %*% (W * (U %*% S)) + lambda * S;
      alpha = norm_R2 / sum(S * HS);
      V = V + alpha * S;
    }
    R = R - alpha * HS;
    old_norm_R2 = norm_R2; norm_R2 = sum(R ^ 2);
    S = R + (norm_R2 / old_norm_R2) * S;
    ii = ii + 1;
  }
  is_U = ! is_U;
}

```

```

U = rand(nrow(X), r, min = -1.0, max = 1.0);
V = rand(r, ncol(X), min = -1.0, max = 1.0);
while(i < mi) {
  i = i + 1; ii = 1;
  if (is_U)
    G = (W * (U %*% V - X)) %*% t(V) + lambda * U;
  else
    G = t(U) %*% (W * (U %*% V - X)) + lambda * V;
  norm_G2 = sum(G ^ 2); norm_R2 = norm_G2;
  R = -G; S = R;
  while(norm_R2 > 10E-9 * norm_G2 & ii <= mii) {
    if (is_U) {
      HS = (W * (S %*% V)) %*% t(V) + lambda * S;
      alpha = norm_R2 / sum(S * HS);
      U = U + alpha * S;
    } else {
      HS = t(U) %*% (W * (U %*% S)) + lambda * S;
      alpha = norm_R2 / sum(S * HS);
      V = V + alpha * S;
    }
    R = R - alpha * HS;
    old_norm_R2 = norm_R2; norm_R2 = sum(R ^ 2);
    S = R + (norm_R2 / old_norm_R2) * S;
    ii = ii + 1;
  }
  is_U = ! is_U;
}

```

```

U = rand(nrow(X), r, min = -1.0, max = 1.0);
V = rand(r, ncol(X), min = -1.0, max = 1.0);
while(i < mi) {
  i = i + 1; ii = 1;
  if (is_U)
    G = (W * (U %*% V - X)) %*% t(V) + lambda * U;
  else
    G = t(U) %*% (W * (U %*% V - X)) + lambda * V;
  norm_G2 = sum(G ^ 2); norm_R2 = norm_G2;
  R = -G; S = R;
  while(norm_R2 > 10E-9 * norm_G2 & ii <= mii) {
    if (is_U) {
      HS = (W * (S %*% V)) %*% t(V) + lambda * S;
      alpha = norm_R2 / sum(S * HS);
      U = U + alpha * S;
    } else {
      HS = t(U) %*% (W * (U %*% S)) + lambda * S;
      alpha = norm_R2 / sum(S * HS);
      V = V + alpha * S;
    }
    R = R - alpha * HS;
    old_norm_R2 = norm_R2; norm_R2 = sum(R ^ 2);
    S = R + (norm_R2 / old_norm_R2) * S;
    ii = ii + 1;
  }
  is_U = ! is_U;
}

```

Every line has a clear purpose!

<https://github.com/apache/spark/blob/master/mllib/src/main/scala/org/apache/spark/mllib/recommendation/ALS.scala>

<https://github.com/apache/spark/blob/master/mllib/src/main/scala/org/apache/spark/mllib/recommendation/ALS.scala>

25 lines' worth of algorithm...

...mixed with 800 lines of performance code

```

U = rand(nrow(X), r, min = -1.0, max = 1.0);
V = rand(r, ncol(X), min = -1.0, max = 1.0);
while(i < mi) {
  i = i + 1; ii = 1;
  if (is_U)
    G = (W * (U %*% V - X)) %*% t(V) + lambda * U;
  else
    G = t(U) %*% (W * (U %*% V - X)) + lambda * V;
  norm_G2 = sum(G ^ 2); norm_R2 = norm_G2;
  R = -G; S = R;
  while(norm_R2 > 10E-9 * norm_G2 & ii <= mii) {
    if (is_U) {
      HS = (W * (S %*% V)) %*% t(V) + lambda * S;
      alpha = norm_R2 / sum(S * HS);
      U = U + alpha * S;
    } else {
      HS = t(U) %*% (W * (U %*% S)) + lambda * S;
      alpha = norm_R2 / sum(S * HS);
      V = V + alpha * S;
    }
    R = R - alpha * HS;
    old_norm_R2 = norm_R2; norm_R2 = sum(R ^ 2);
    S = R + (norm_R2 / old_norm_R2) * S;
    ii = ii + 1;
  }
  is_U = ! is_U;
}

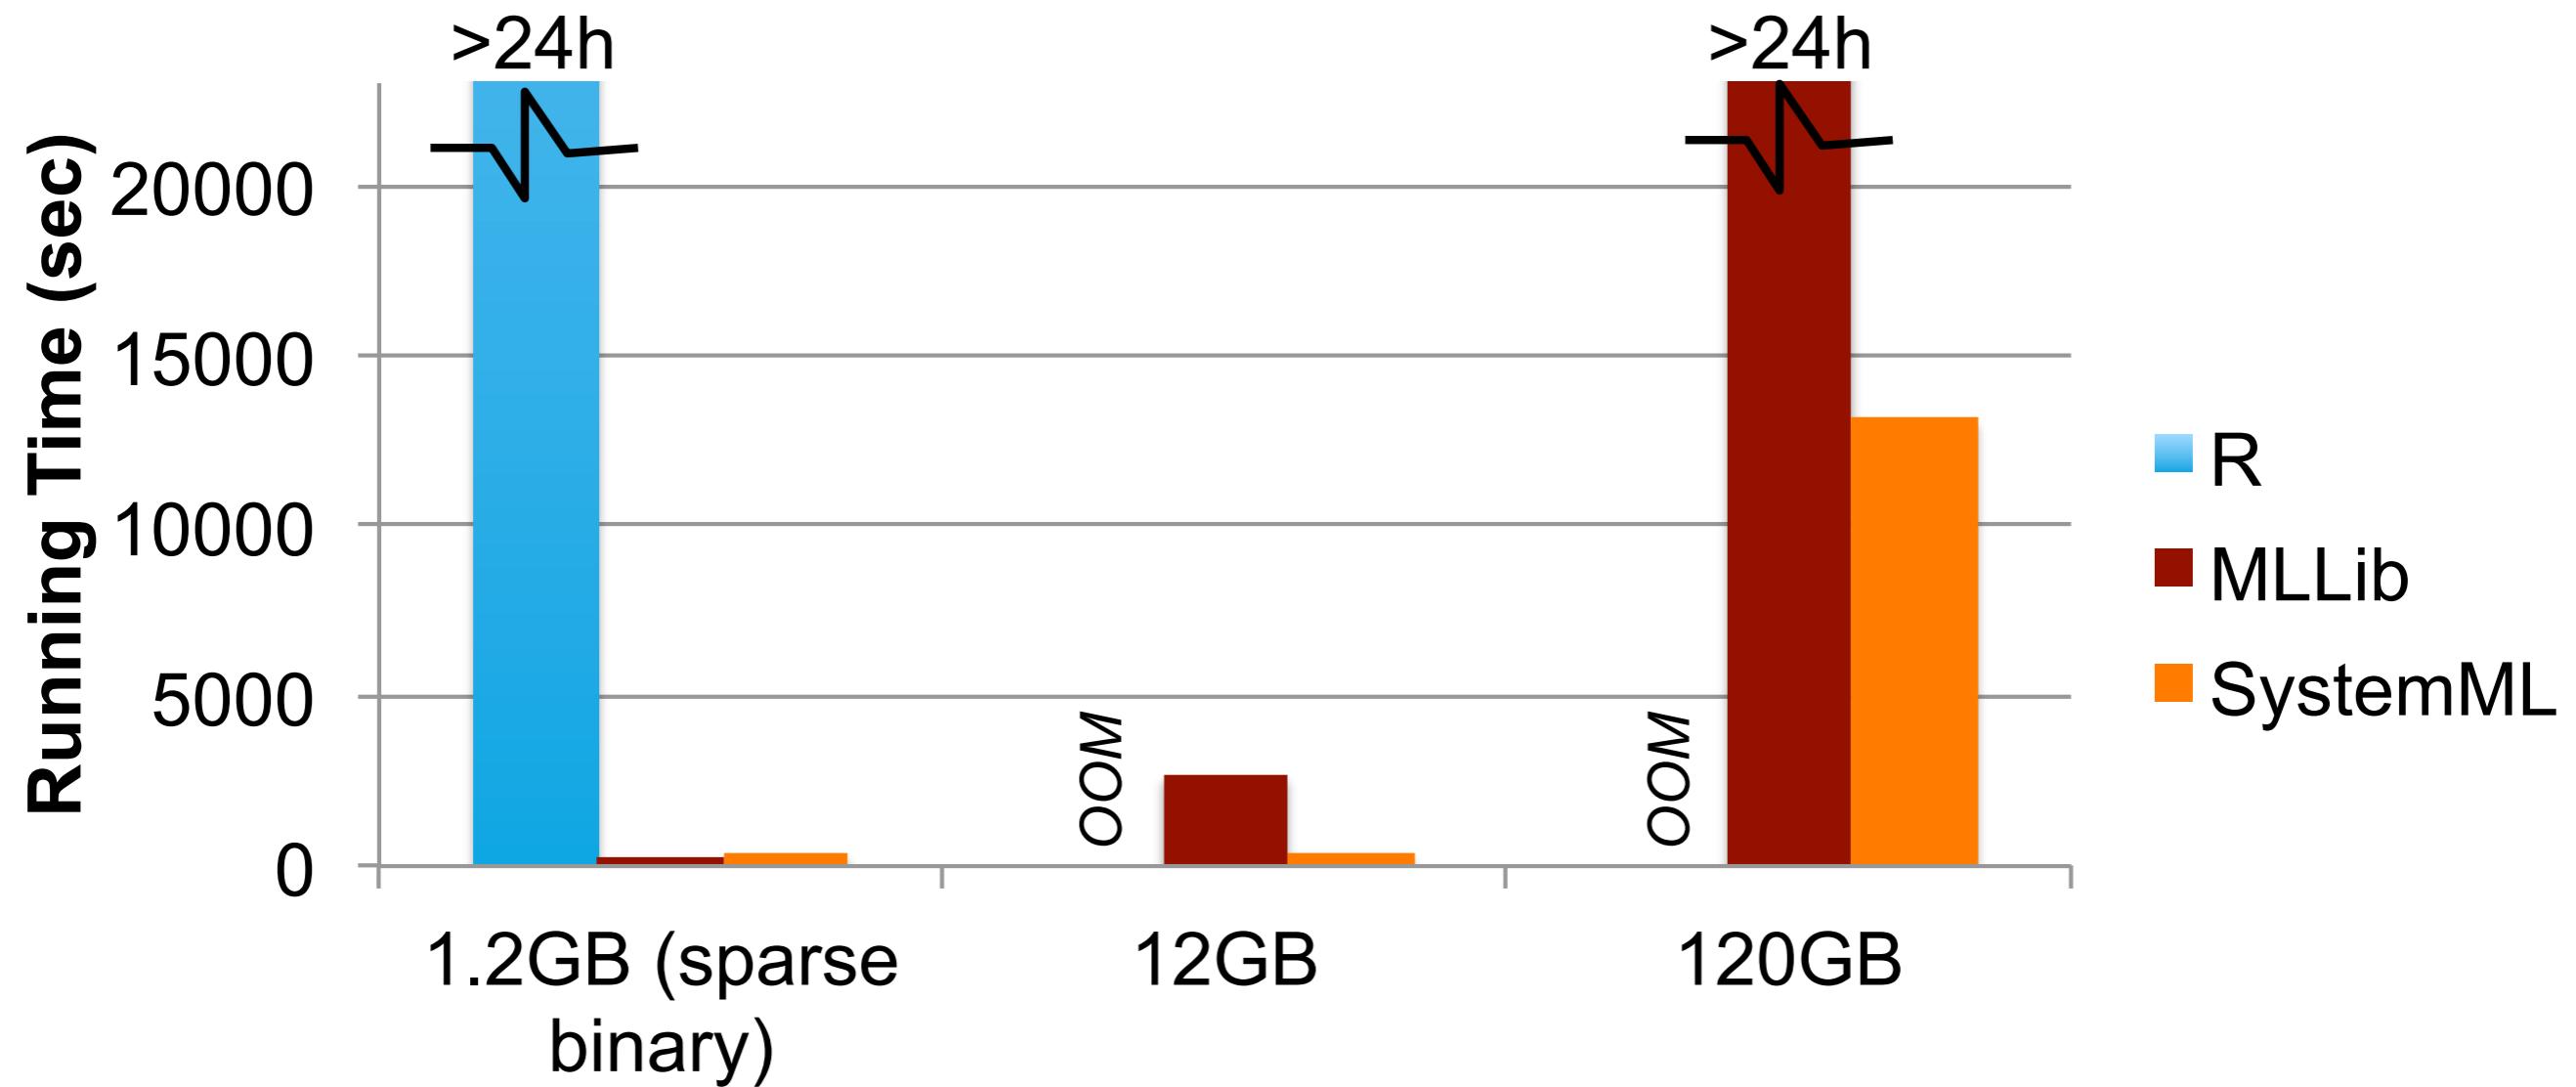
```

```

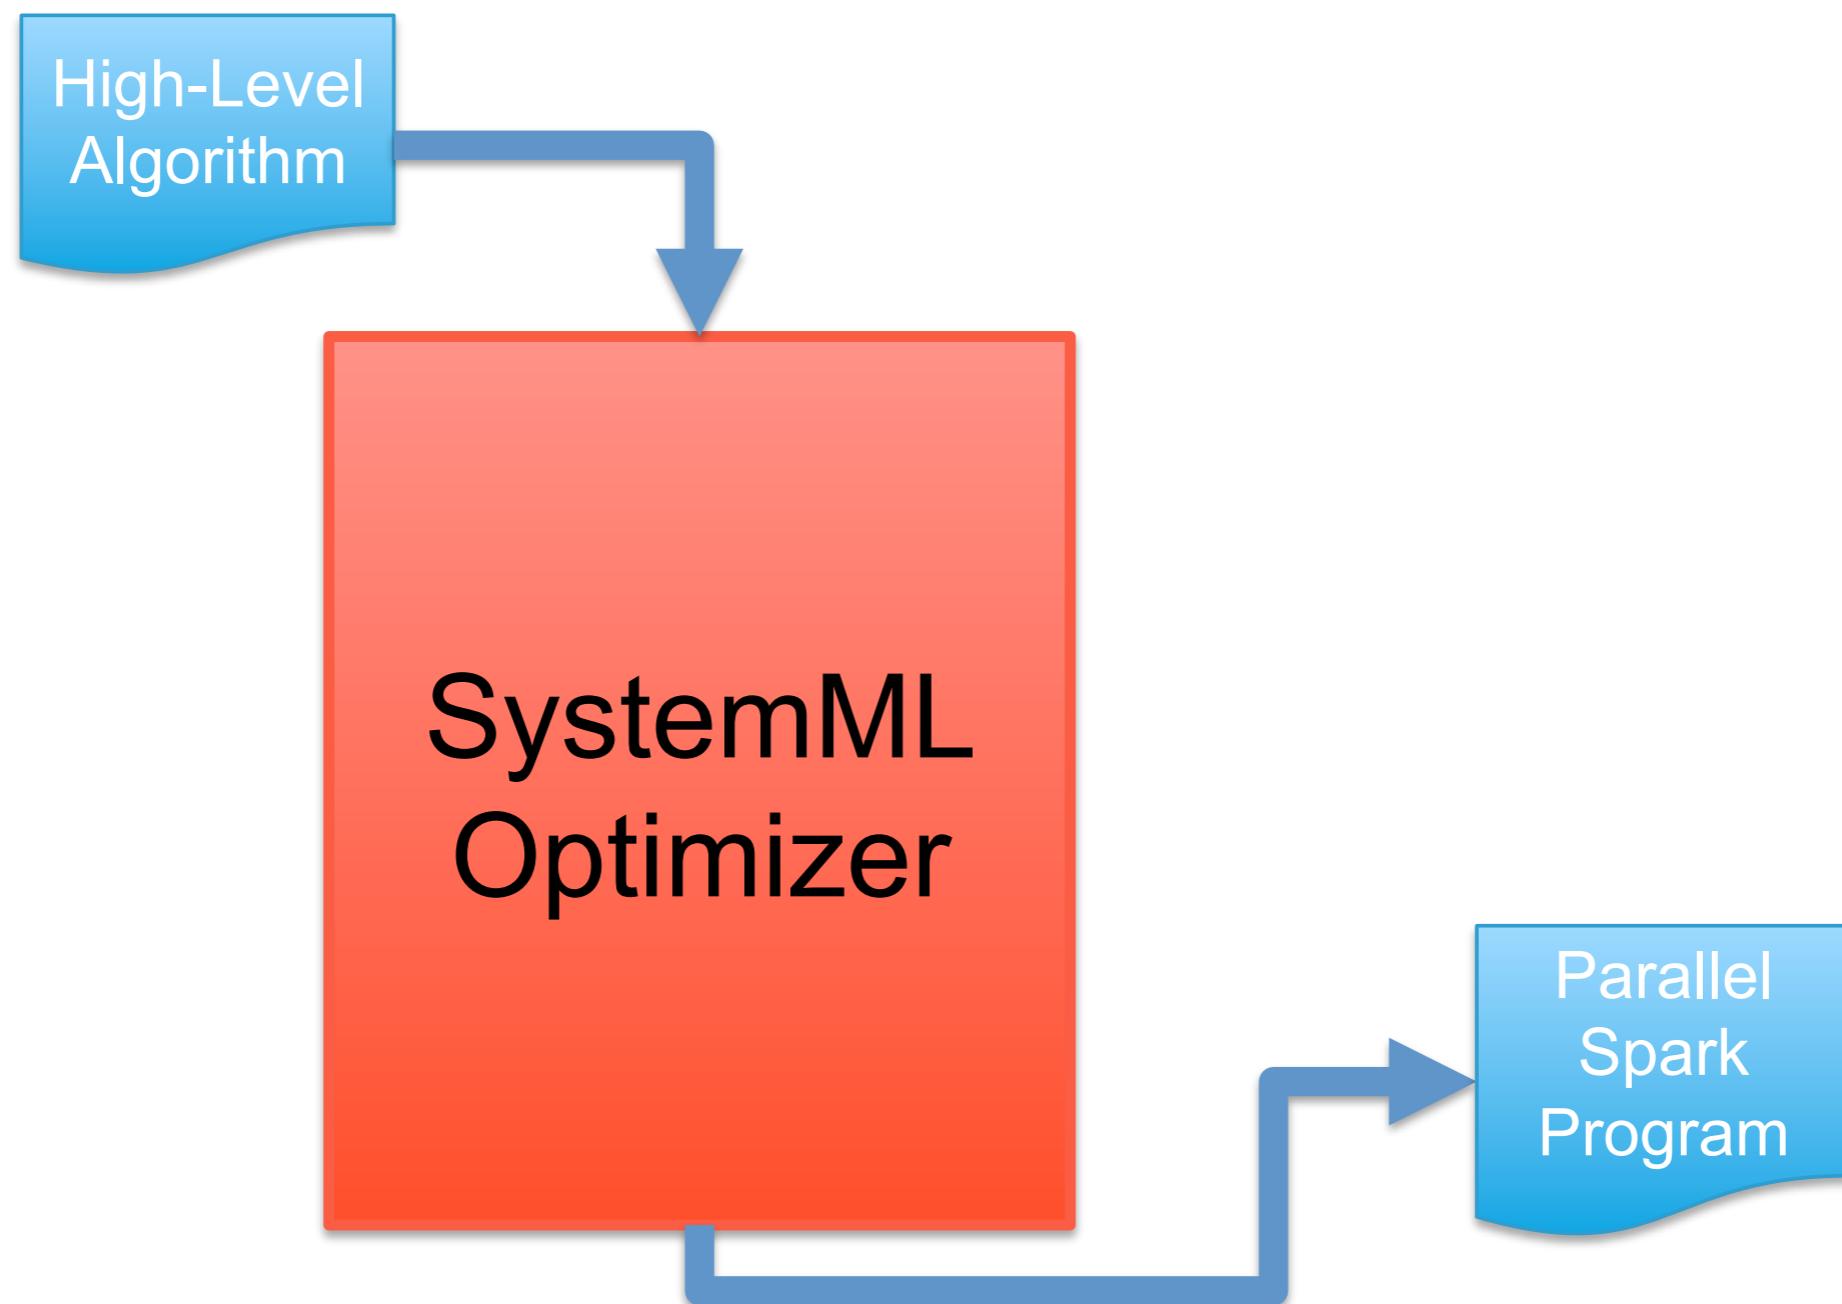
U = rand(nrow(X), r, min = -1.0, max = 1.0);
V = rand(r, ncol(X), min = -1.0, max = 1.0);
while(i < mi) {
  i = i + 1; ii = 1;
  if (is_U)
    G = (W * (U %*% V - X)) %*% t(V) + lambda * U;
  else
    G = t(U) %*% (W * (U %*% V - X)) + lambda * V;
  norm_G2 = sum(G ^ 2); norm_R2 = norm_G2;
  R = -G; S = R;
  while(norm_R2 > 10E-9 * norm_G2 & ii <= mii) {
    if (is_U) {
      HS = (W * (S %*% V)) %*% t(V) + lambda * S;
      alpha = norm_R2 / sum(S * HS);
      U = U + alpha * S;
    } else {
      HS = t(U) %*% (W * (U %*% S)) + lambda * S;
      alpha = norm_R2 / sum(S * HS);
      V = V + alpha * S;
    }
    R = R - alpha * HS;
    old_norm_R2 = norm_R2; norm_R2 = sum(R ^ 2);
    S = R + (norm_R2 / old_norm_R2) * S;
    ii = ii + 1;
  }
  is_U = ! is_U;
}

```

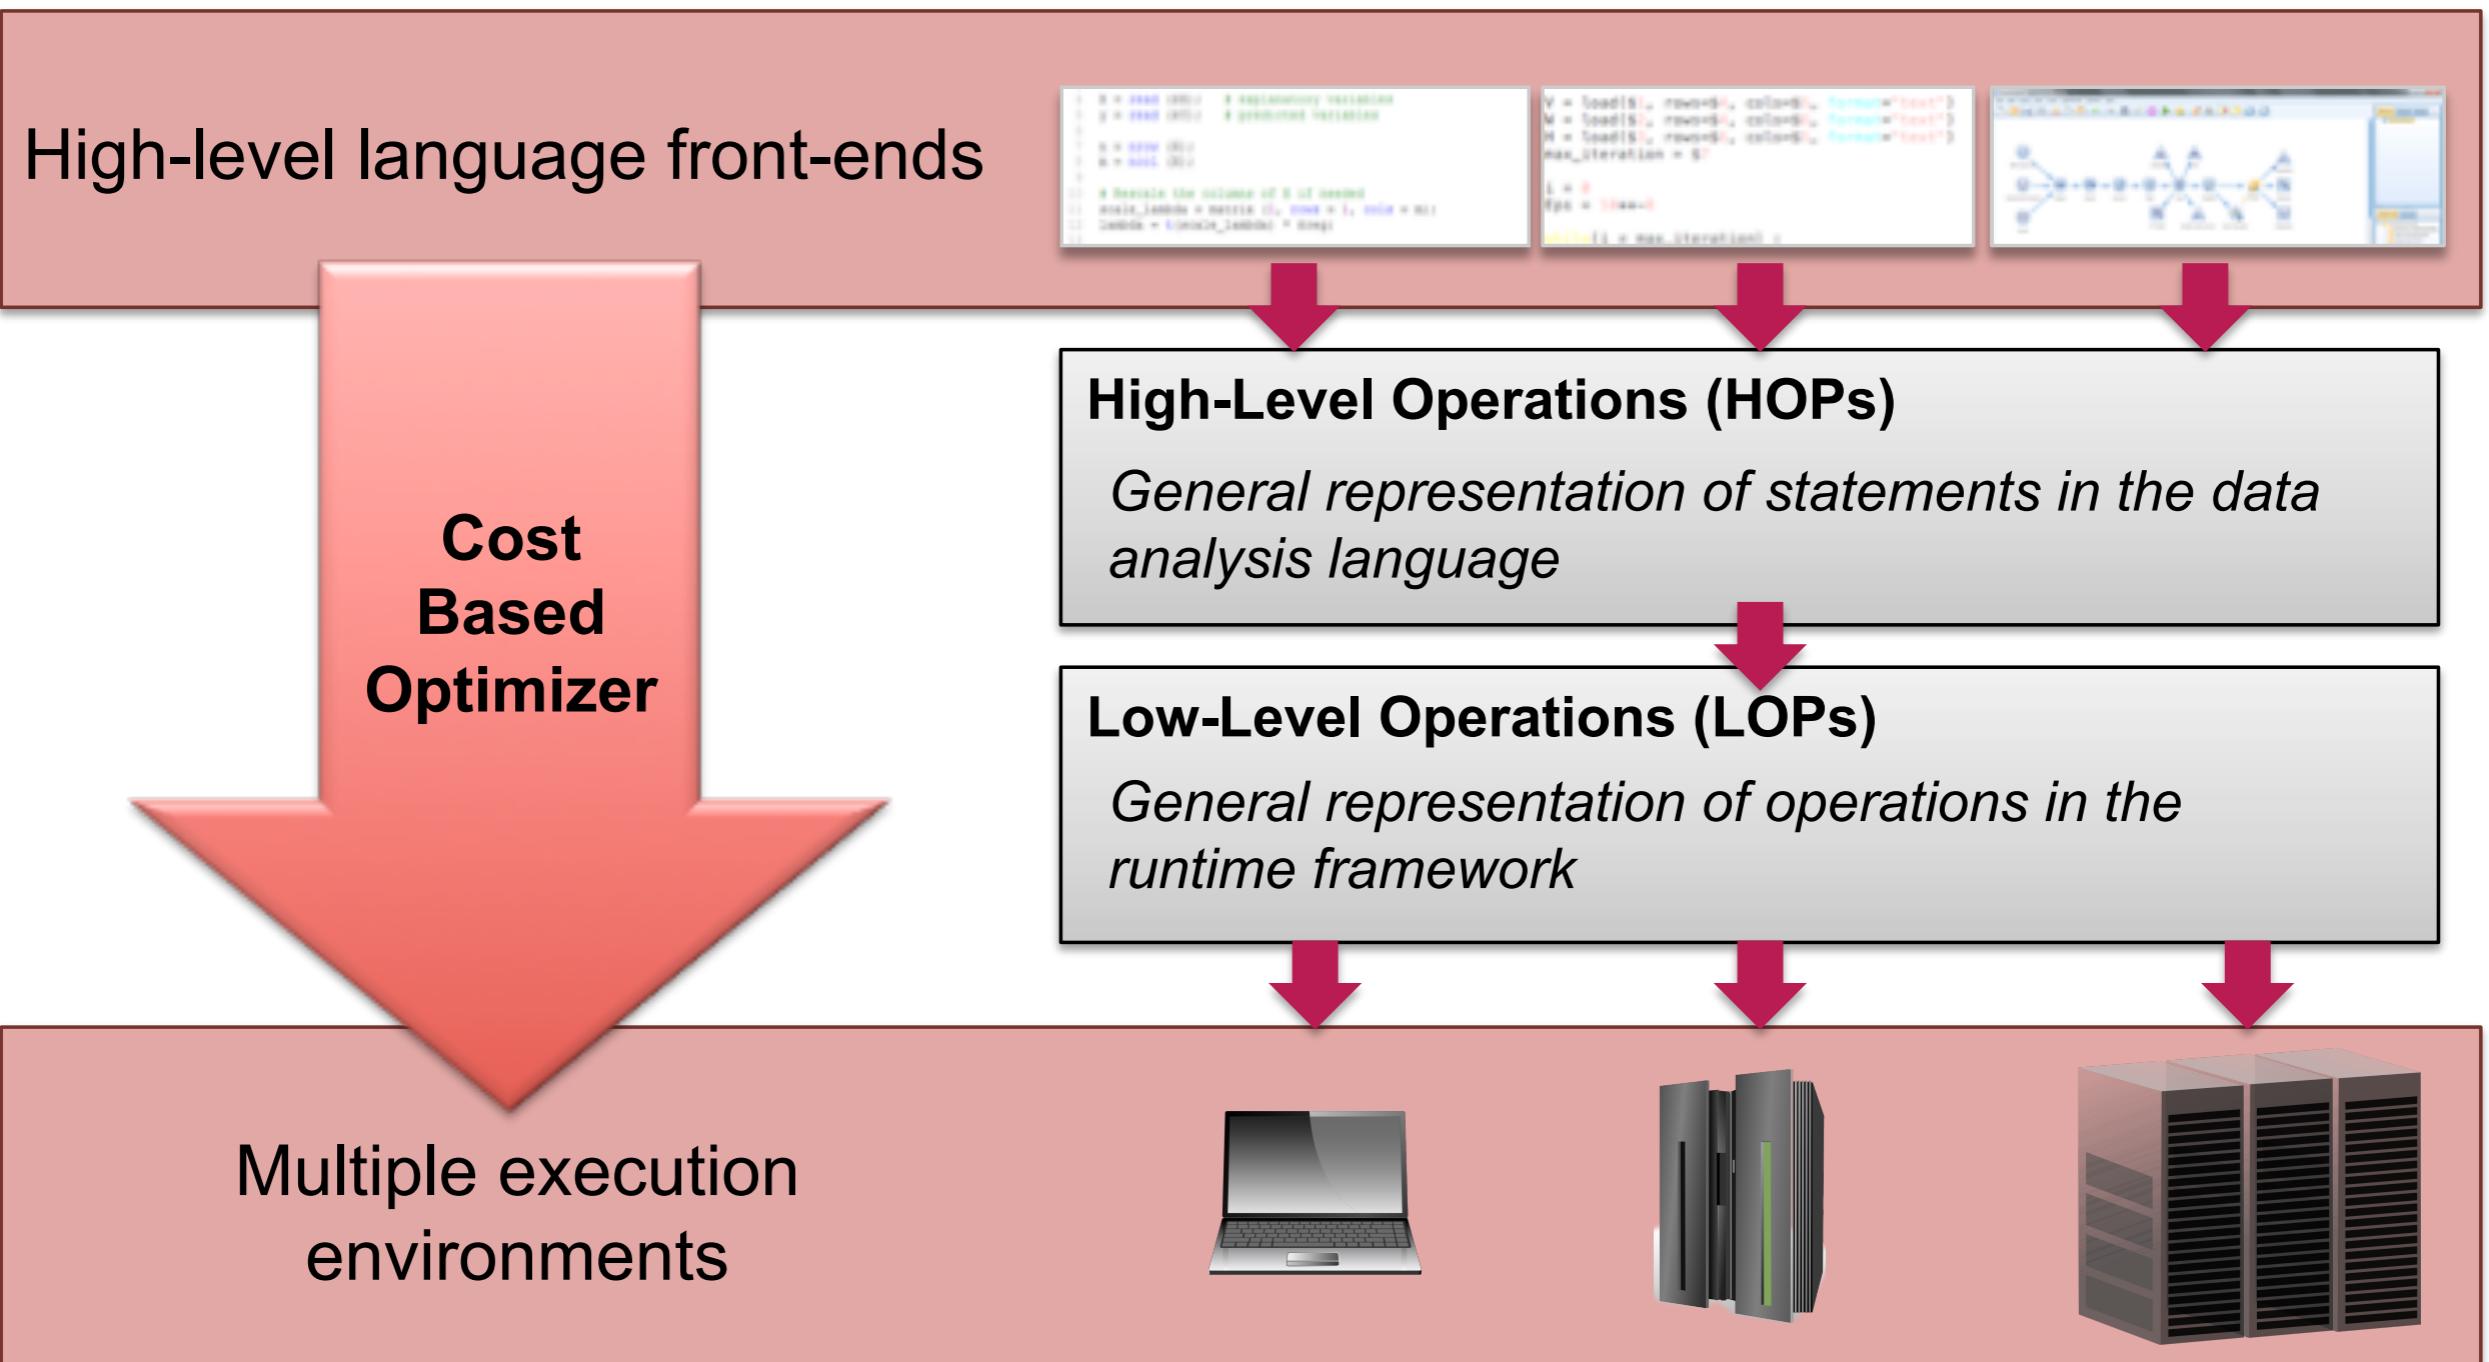
SystemML:
 compile and run at scale
 no performance code needed!

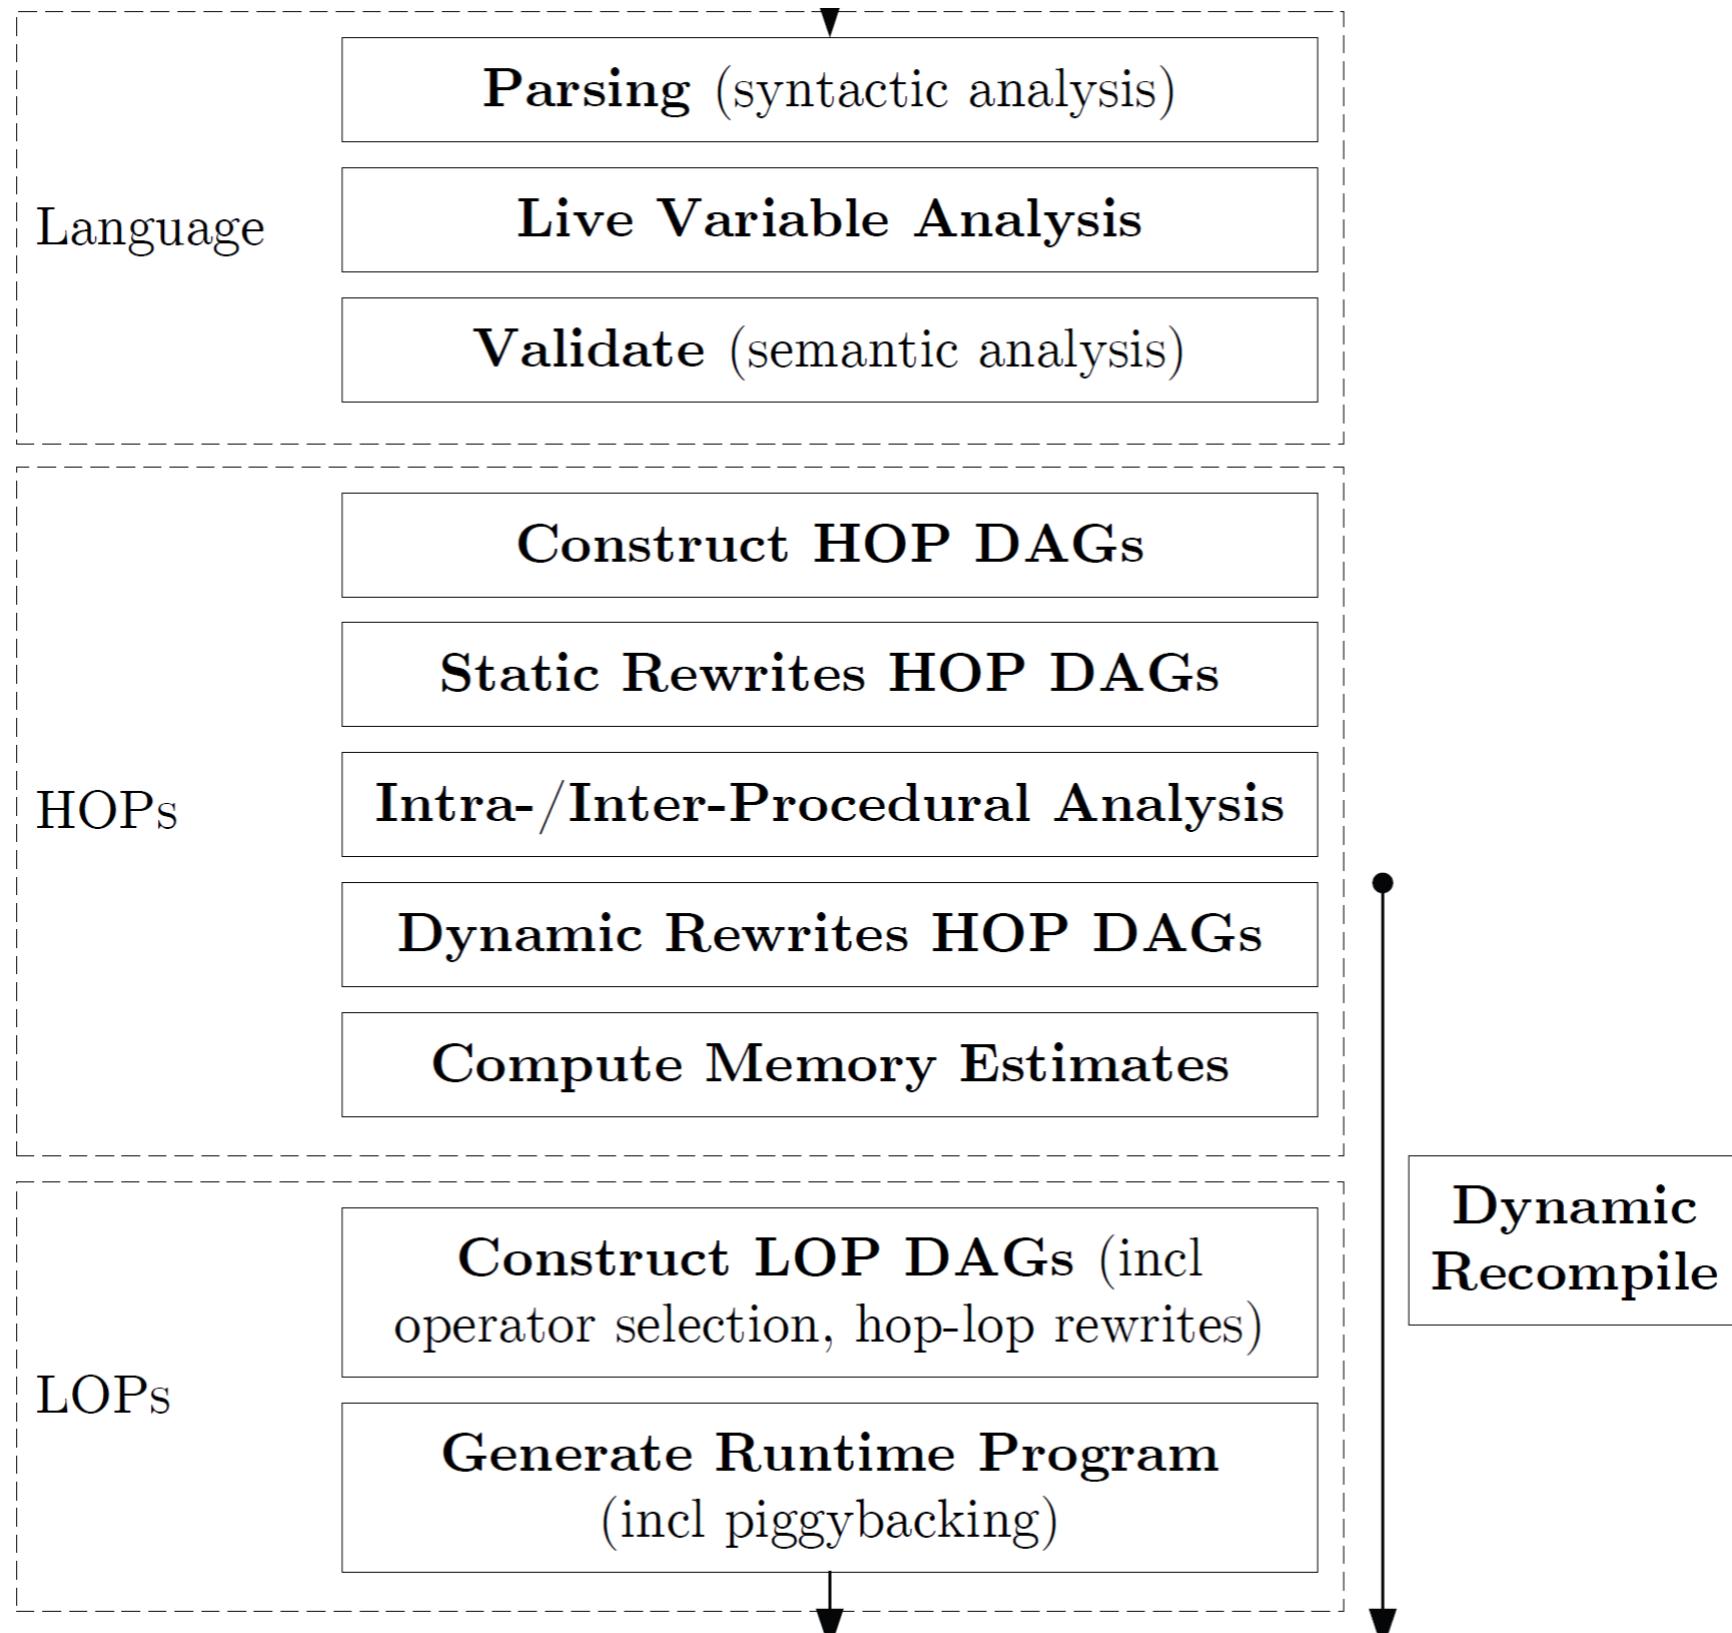


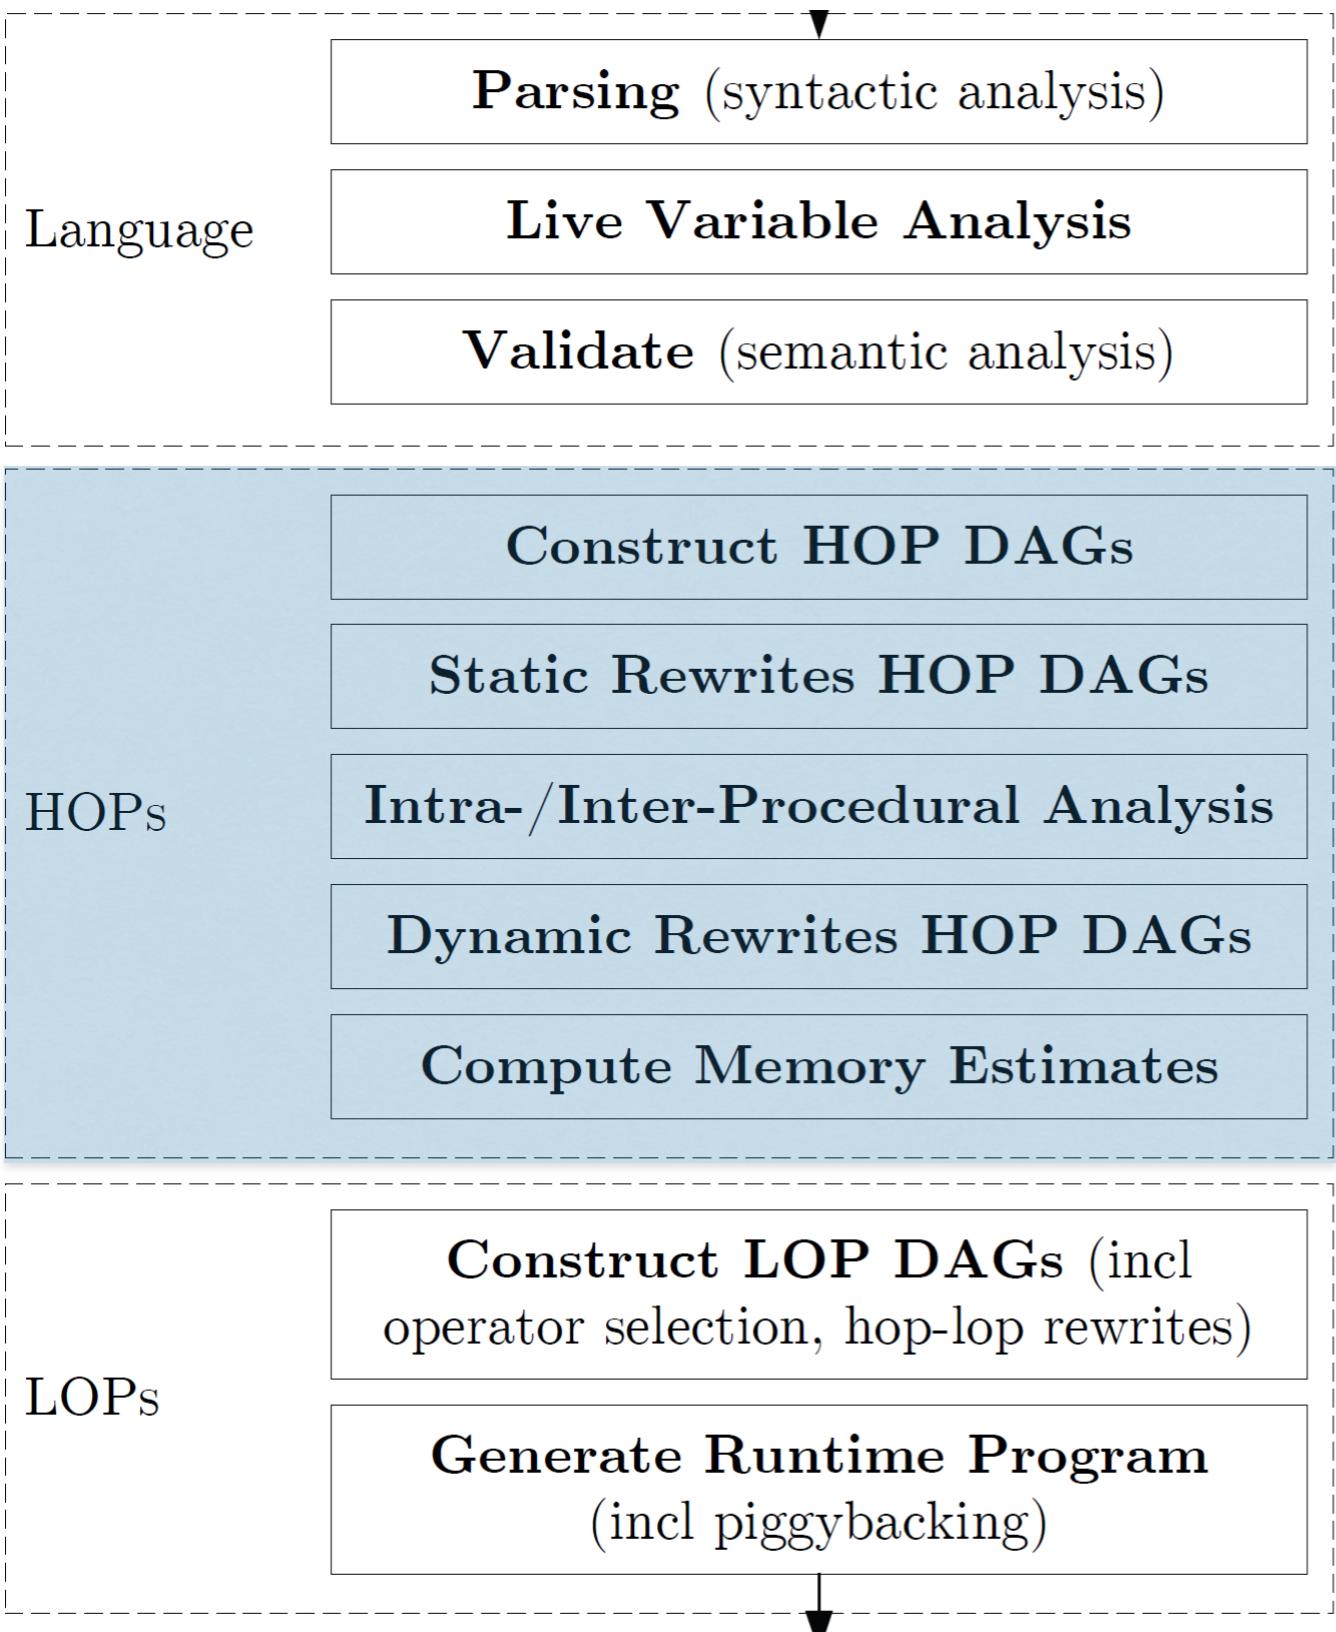
Architecture



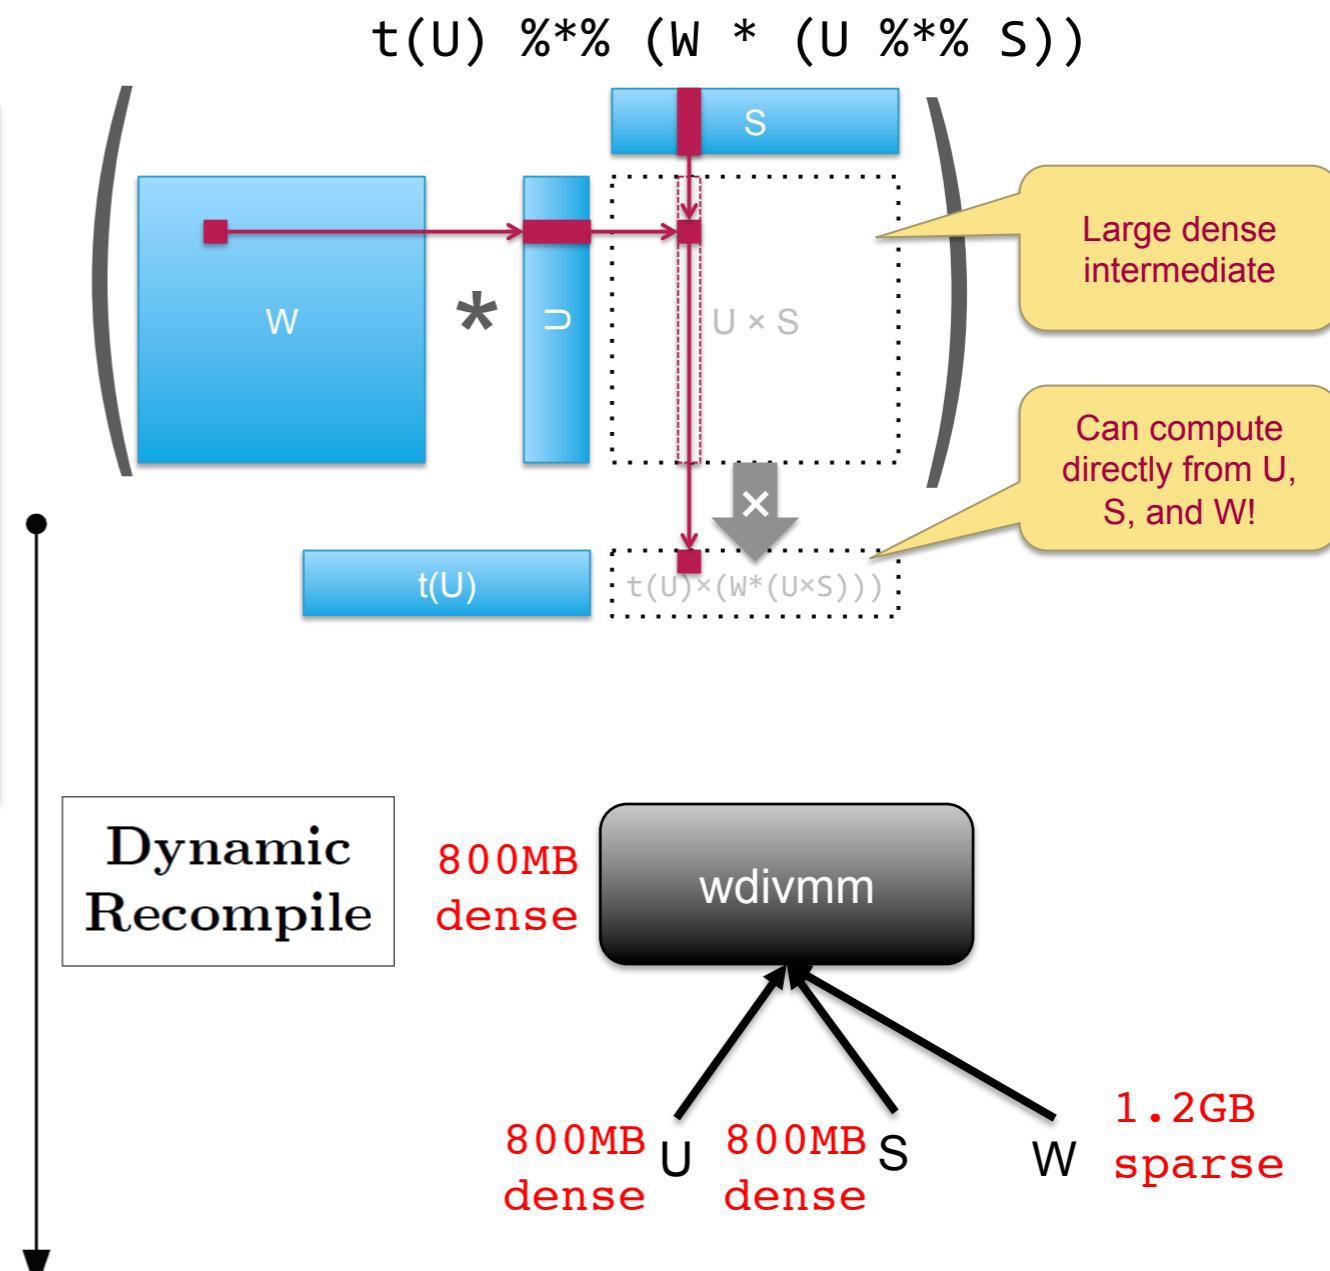
Architecture

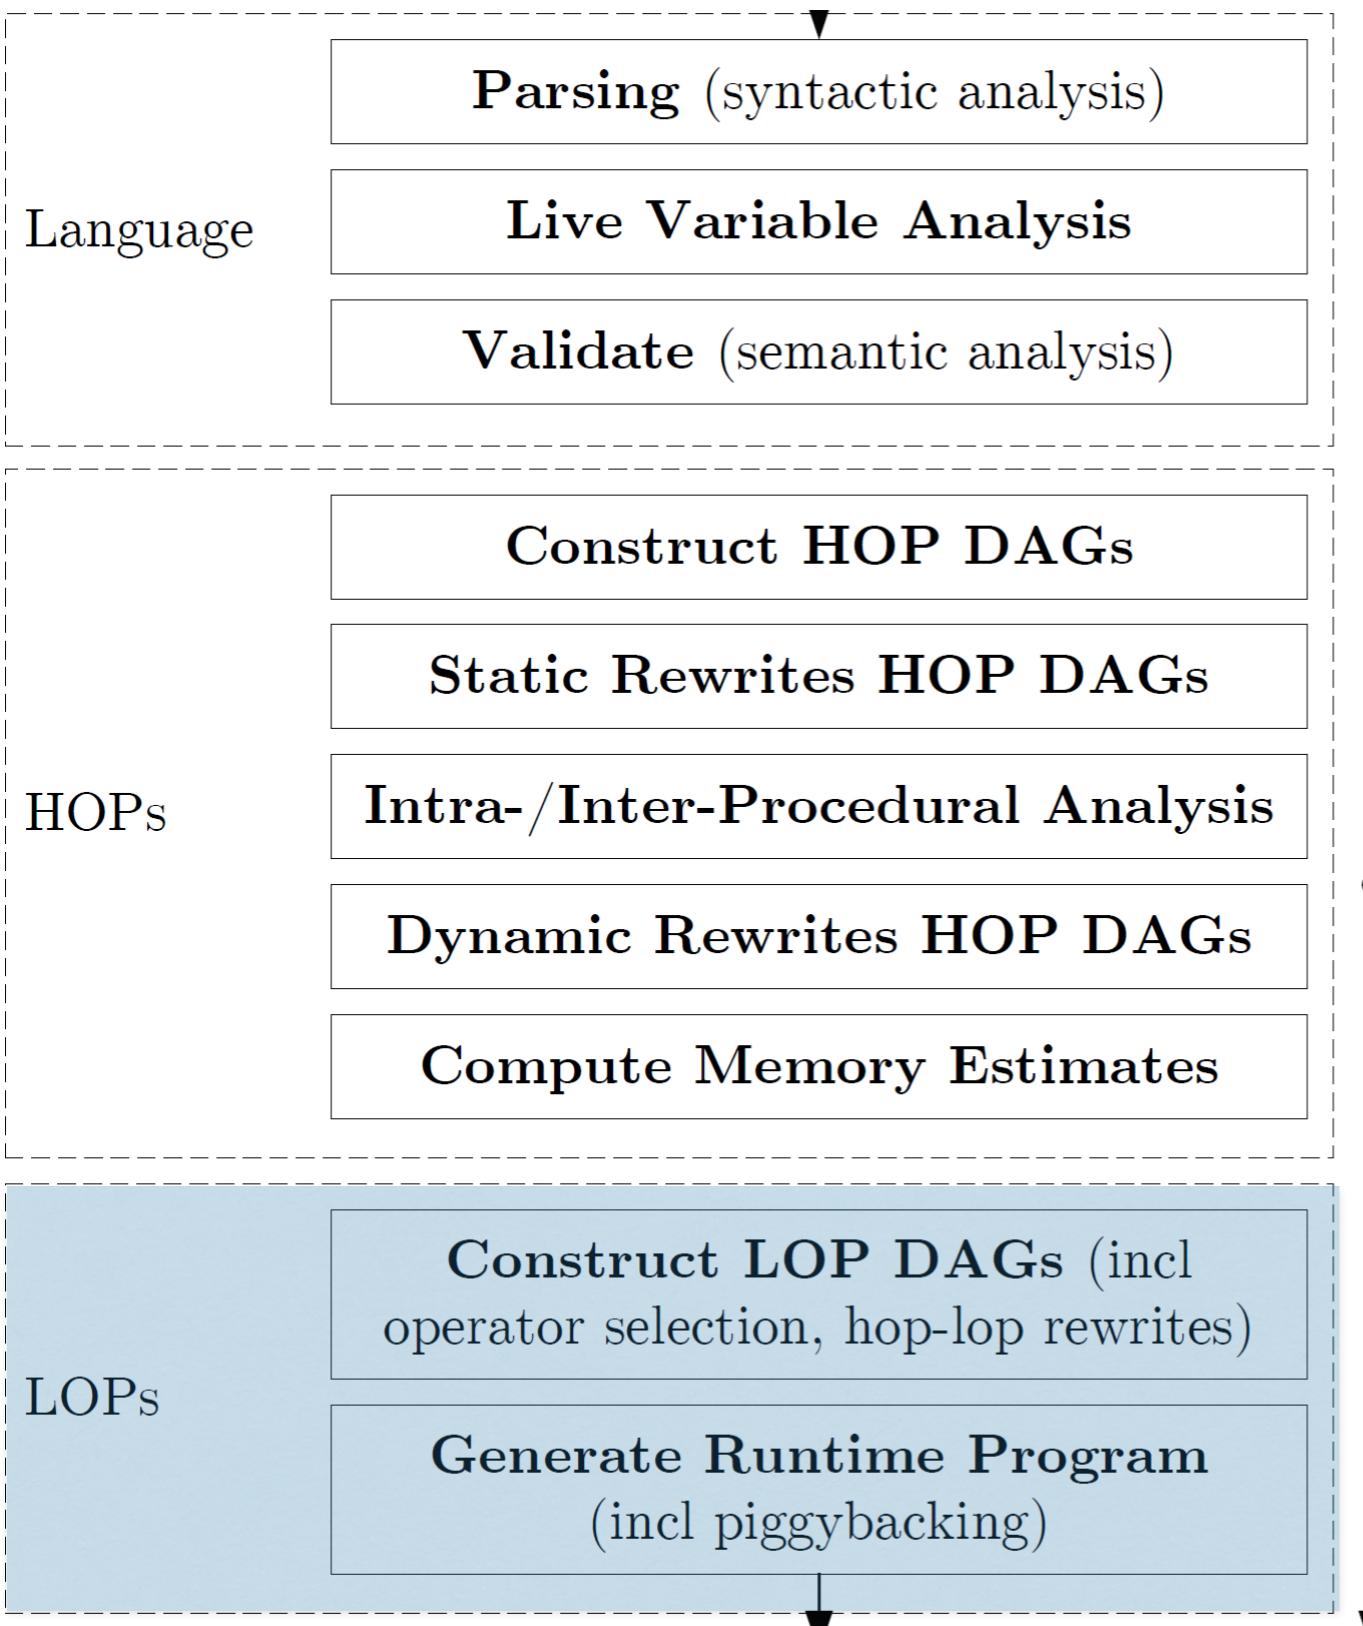




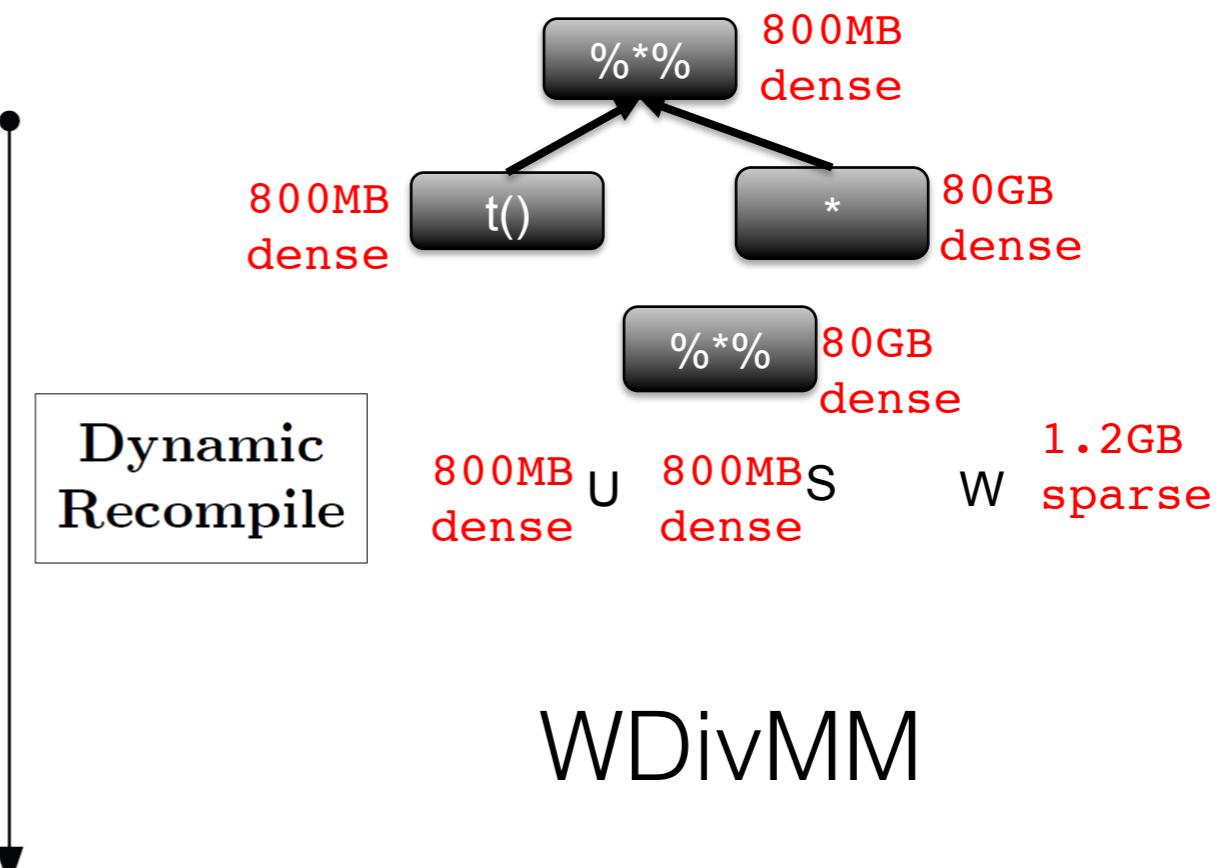


(weighted divide matrix multiplication)





All operands fit into heap
→ use one node





Apache SystemML

Community ▾

GitHub

Documentation

Download

Apache ▾

Apache SystemML

Apache SystemML is a distributed and declarative machine learning platform.

[Get SystemML](#)



Apache SystemML

Community ▾

GitHub

Documentation

Download

Apache ▾

Browse the
source!

Apache SystemML

Apache SystemML is a distributed and declarative machine learning platform.

Get SystemML



Apache SystemML

Apache SystemML

Apache SystemML is a distributed and declarative machine learning platform.

[Get SystemML](#)

Community ▾

GitHub

Documentation

Download

Apache ▾

Browse the
source!

Try out
some
tutorials!



Apache SystemML

Apache System

Apache SystemML is a distributed and declarative machine learning platform.

Get SystemML

Community ▾

Contribute to
the project!

GitHub

Browse the
source!

Documentation

Download

Apache ▾

Try out
some
tutorials!



Apache SystemML

Apache System

Apache SystemML is a distributed and declarative machine learning platform.

[Get SystemML](#)

Download the
binary release!

Contribute to
the project!

Browse the
source!

Try out
some
tutorials!

Community ▾

GitHub

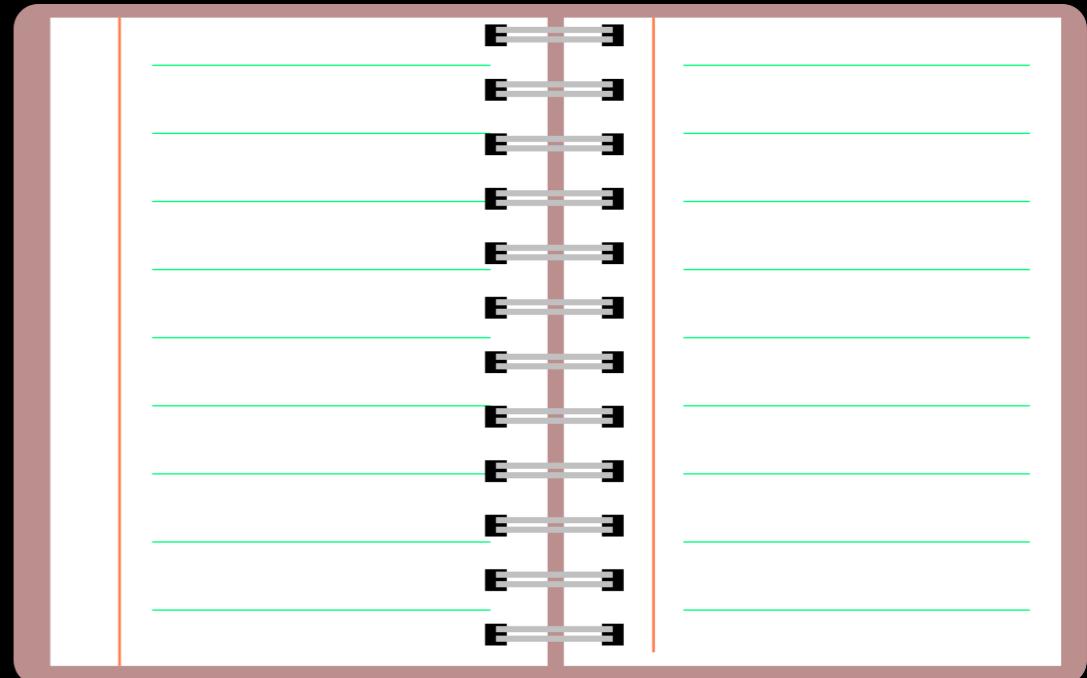
Documentation

Download

Apache ▾

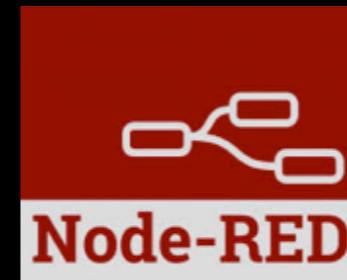
Agenda

- **Introduction to DeepLearning**
- **Introduction to ApacheSpark**
- **Introduction to DL4J**
- **Introduction to SystemML**
- **Demos**





- 6000+ clients
- \$3B investment
- Partners including
 - Avnet, BNP Paribas, EEBus, Capgemini, Tech Mahindra, Vodafone, BMW, Visa, Bosch, Indiegogo, French national railway SNCF, Arrow Electronics, Intel, Cisco

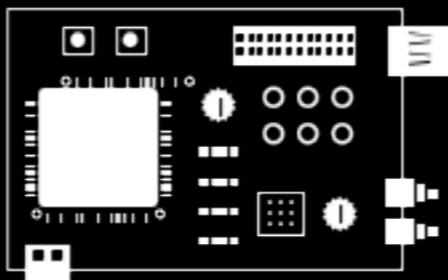


Why IoT (now) ?

- 15 Billion connected devices in 2015
- 40 Billion connected devices in 2020
- World population 7.4 Billion in 2016

Now you can believe the hype

OASIS makes MQTT gold standard for Internet of Things



Your device or gateway

We start with your device, be it a sensor, a gateway or something else. To find out how to get it connected, search our recipes.



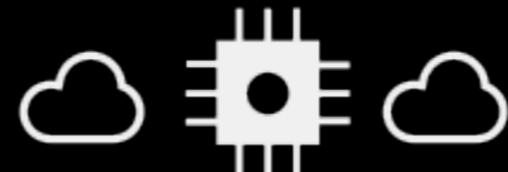
MQTT & HTTP

Your device data is sent securely up to the cloud using the open, lightweight MQTT messaging protocol or HTTP.



REST & Real-time APIs

Use our secure APIs to connect your apps with the data coming from your devices.



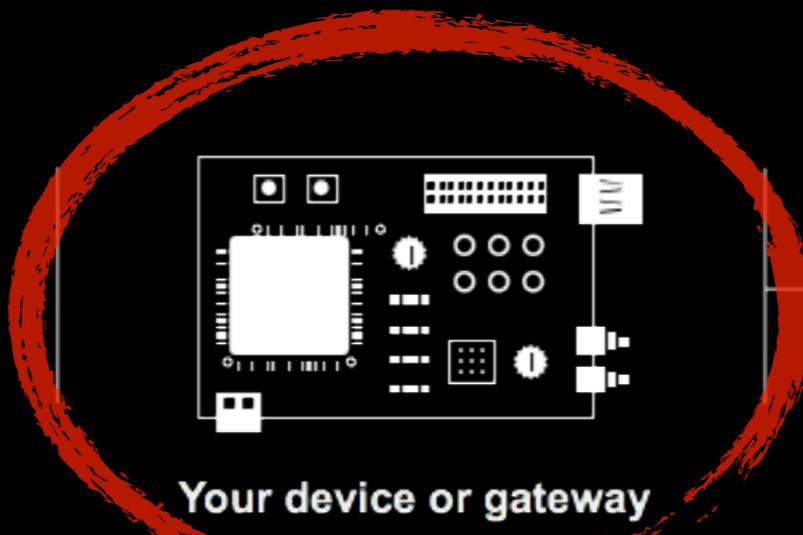
IBM Watson IoT Platform

This is the hub of all things IBM IoT. This is where you can setup and manage your connected devices so that your apps can access their live and historical data.



Your application and analytics

Create applications within IBM Bluemix, another cloud, or your own servers to interpret the data you now have access to!



Your device or gateway

We start with your device, be it a sensor, a gateway or something else.
To find out how to get it connected, search our recipes.



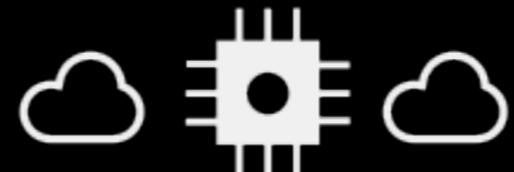
MQTT & HTTP

Your device data is sent securely up to the cloud using the open, lightweight MQTT messaging protocol or HTTP.



REST & Real-time APIs

Use our secure APIs to connect your apps with the data coming from your devices.



IBM Watson IoT Platform

This is the hub of all things IBM IoT. This is where you can setup and manage your connected devices so that your apps can access their live and historical data.



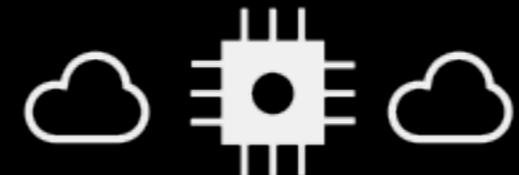
Your application and analytics

Create applications within IBM Bluemix, another cloud, or your own servers to interpret the data you now have access to!



REST & Real-time APIs

Use our secure APIs to connect your apps with the data coming from your devices.



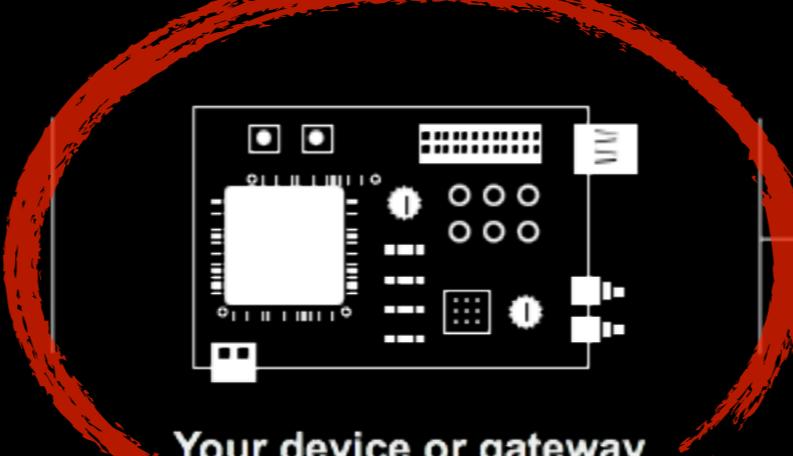
IBM Watson IoT Platform

This is the hub of all things IBM IoT. This is where you can setup and manage your connected devices so that your apps can access their live and historical data.



Your application and analytics

Create applications within IBM Bluemix, another cloud, or your own servers to interpret the data you now have access to!



Your device or gateway

We start with your device, be it a sensor, a gateway or something else. To find out how to get it connected, search our recipes.



MQTT & HTTP

Your device data is sent securely up to the cloud using the open, lightweight MQTT messaging protocol or HTTP.



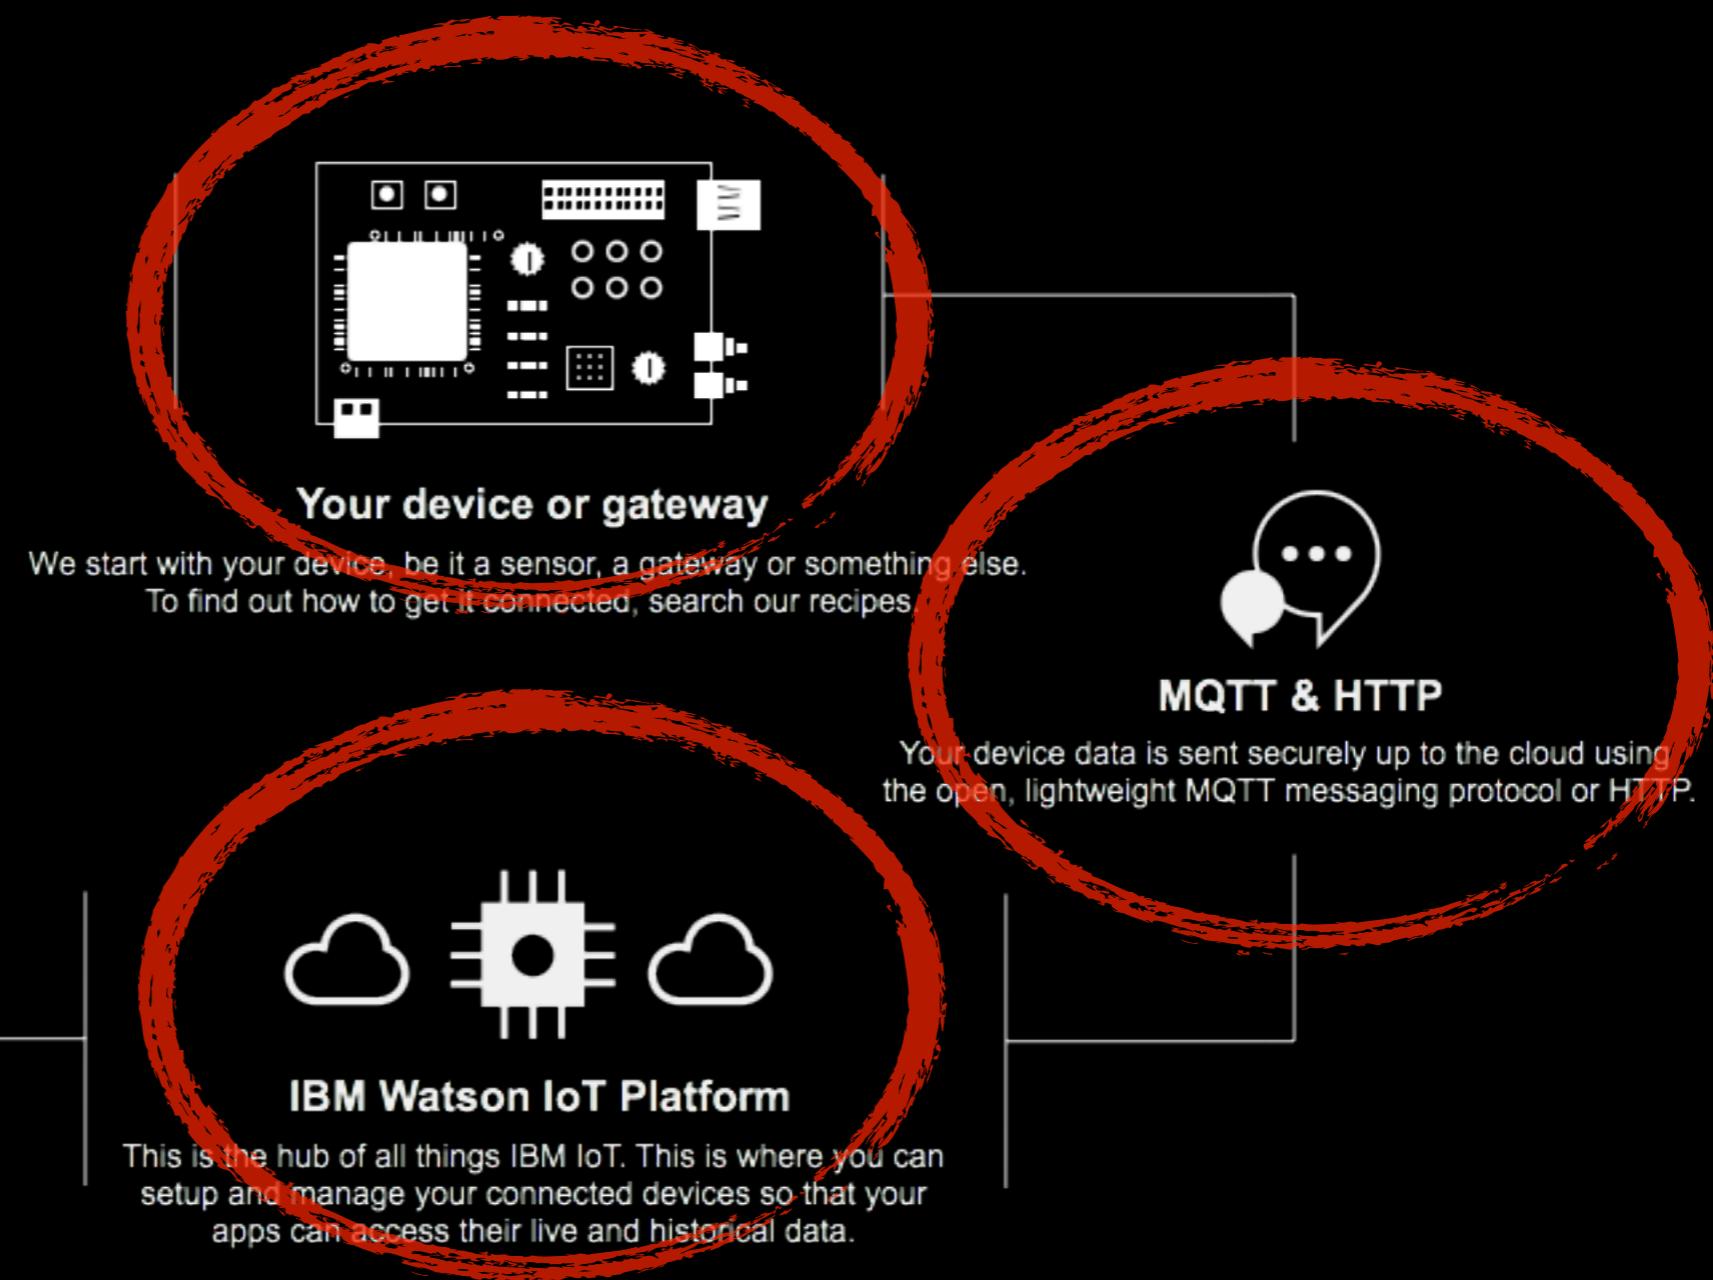
REST & Real-time APIs

Use our secure APIs to connect your apps with the data coming from your devices.



Your application and analytics

Create applications within IBM Bluemix, another cloud, or your own servers to interpret the data you now have access to!





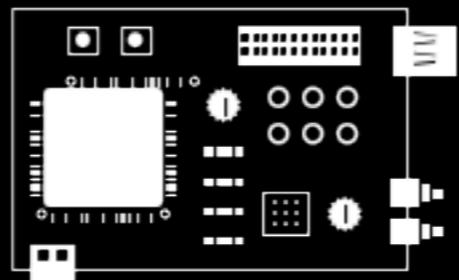
REST & Real-time APIs

Use our secure APIs to connect your apps with the data coming from your devices.



Your application and analytics

Create applications within IBM Bluemix, another cloud, or your own servers to interpret the data you now have access to!



Your device or gateway

We start with your device, be it a sensor, a gateway or something else. To find out how to get it connected, search our recipes.



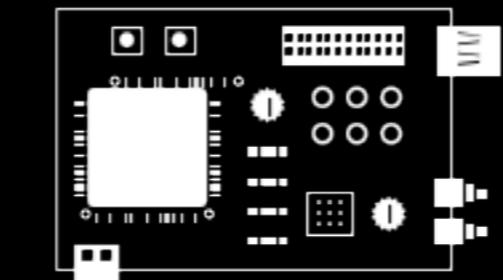
MQTT & HTTP

Your device data is sent securely up to the cloud using the open, lightweight MQTT messaging protocol or HTTP.



REST & Real-time APIs

Use our secure APIs to connect your apps with the data coming from your devices.



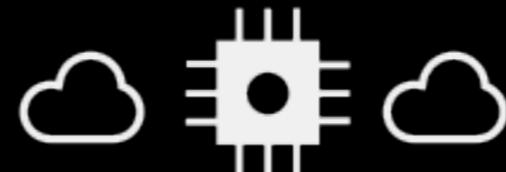
Your device or gateway

We start with your device, be it a sensor, a gateway or something else. To find out how to get it connected, search our recipes.



MQTT & HTTP

Your device data is sent securely up to the cloud using the open, lightweight MQTT messaging protocol or HTTP.



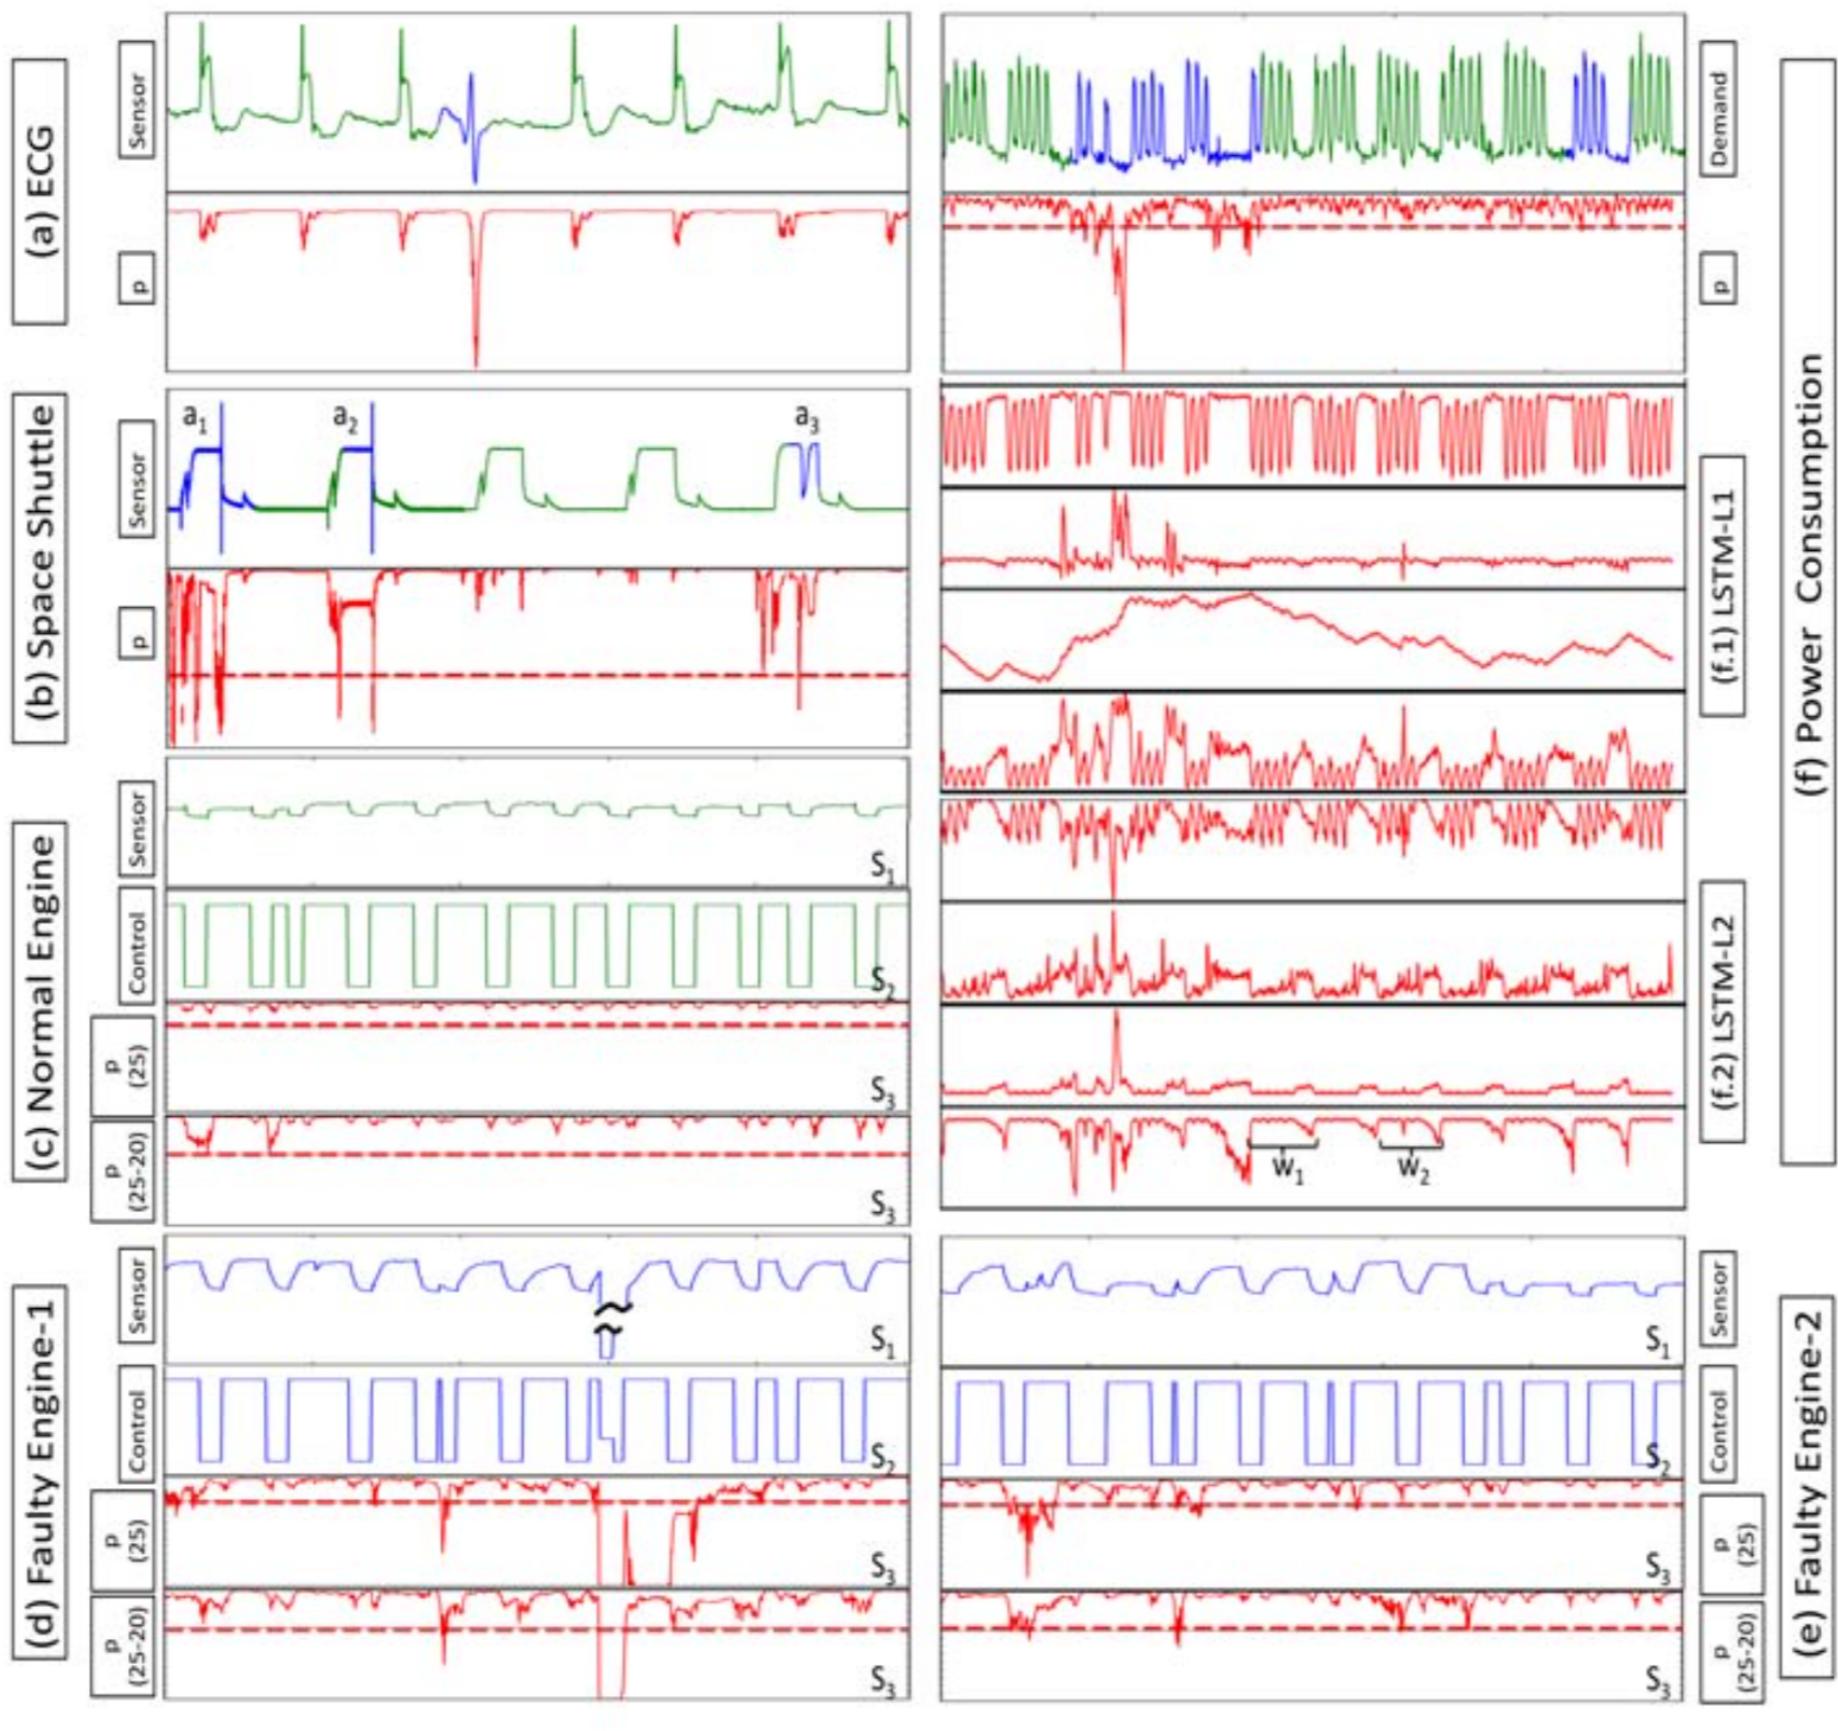
IBM Watson IoT Platform

This is the hub of all things IBM IoT. This is where you can setup and manage your connected devices so that your apps can access their live and historical data.



Your application and analytics

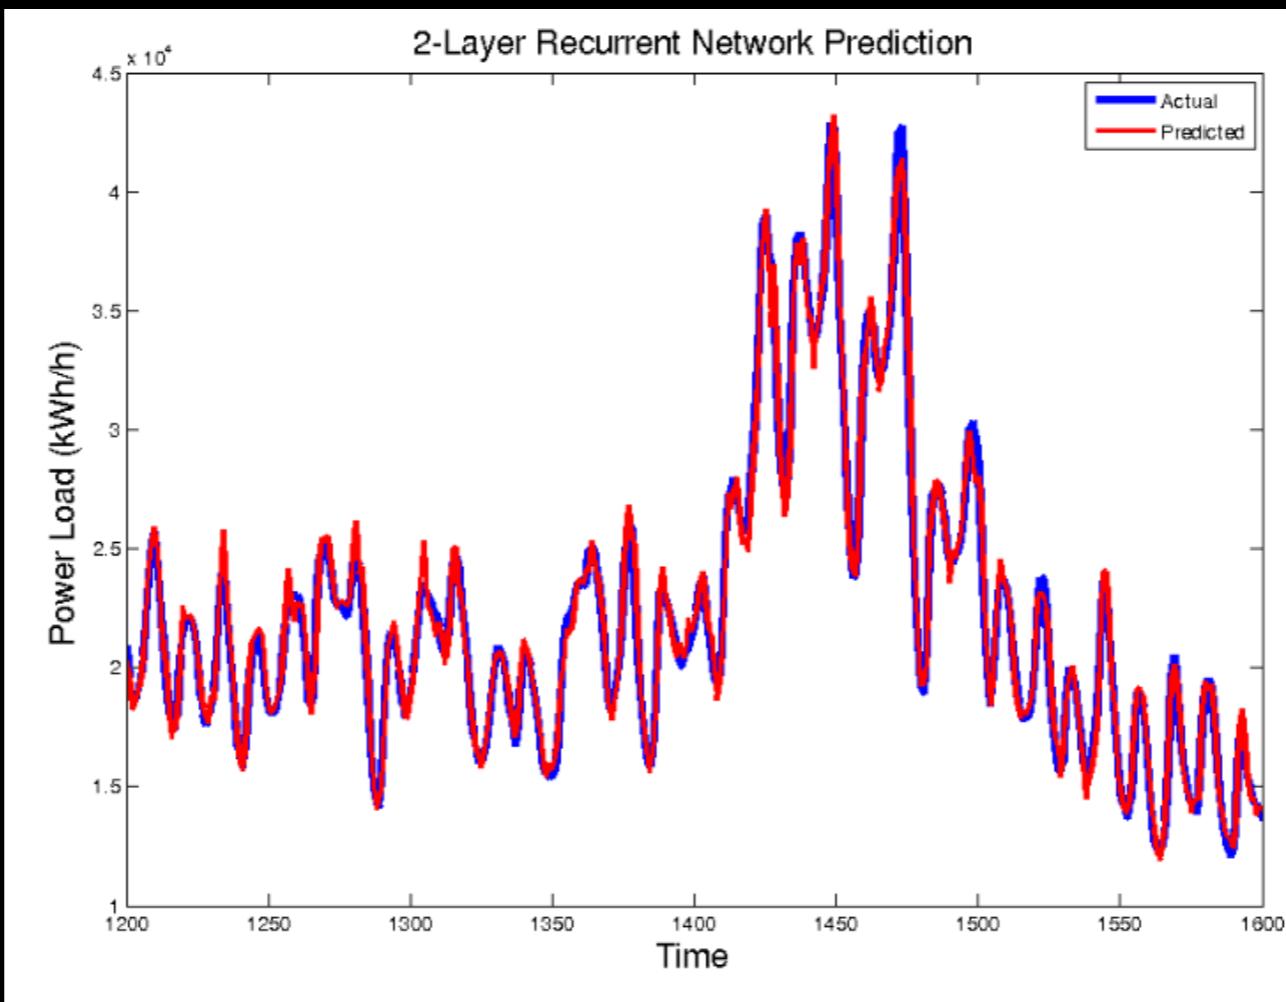
Create applications within IBM Bluemix, another cloud, or your own servers to interpret the data you now have access to!



Deep Learning for Time Series Modeling

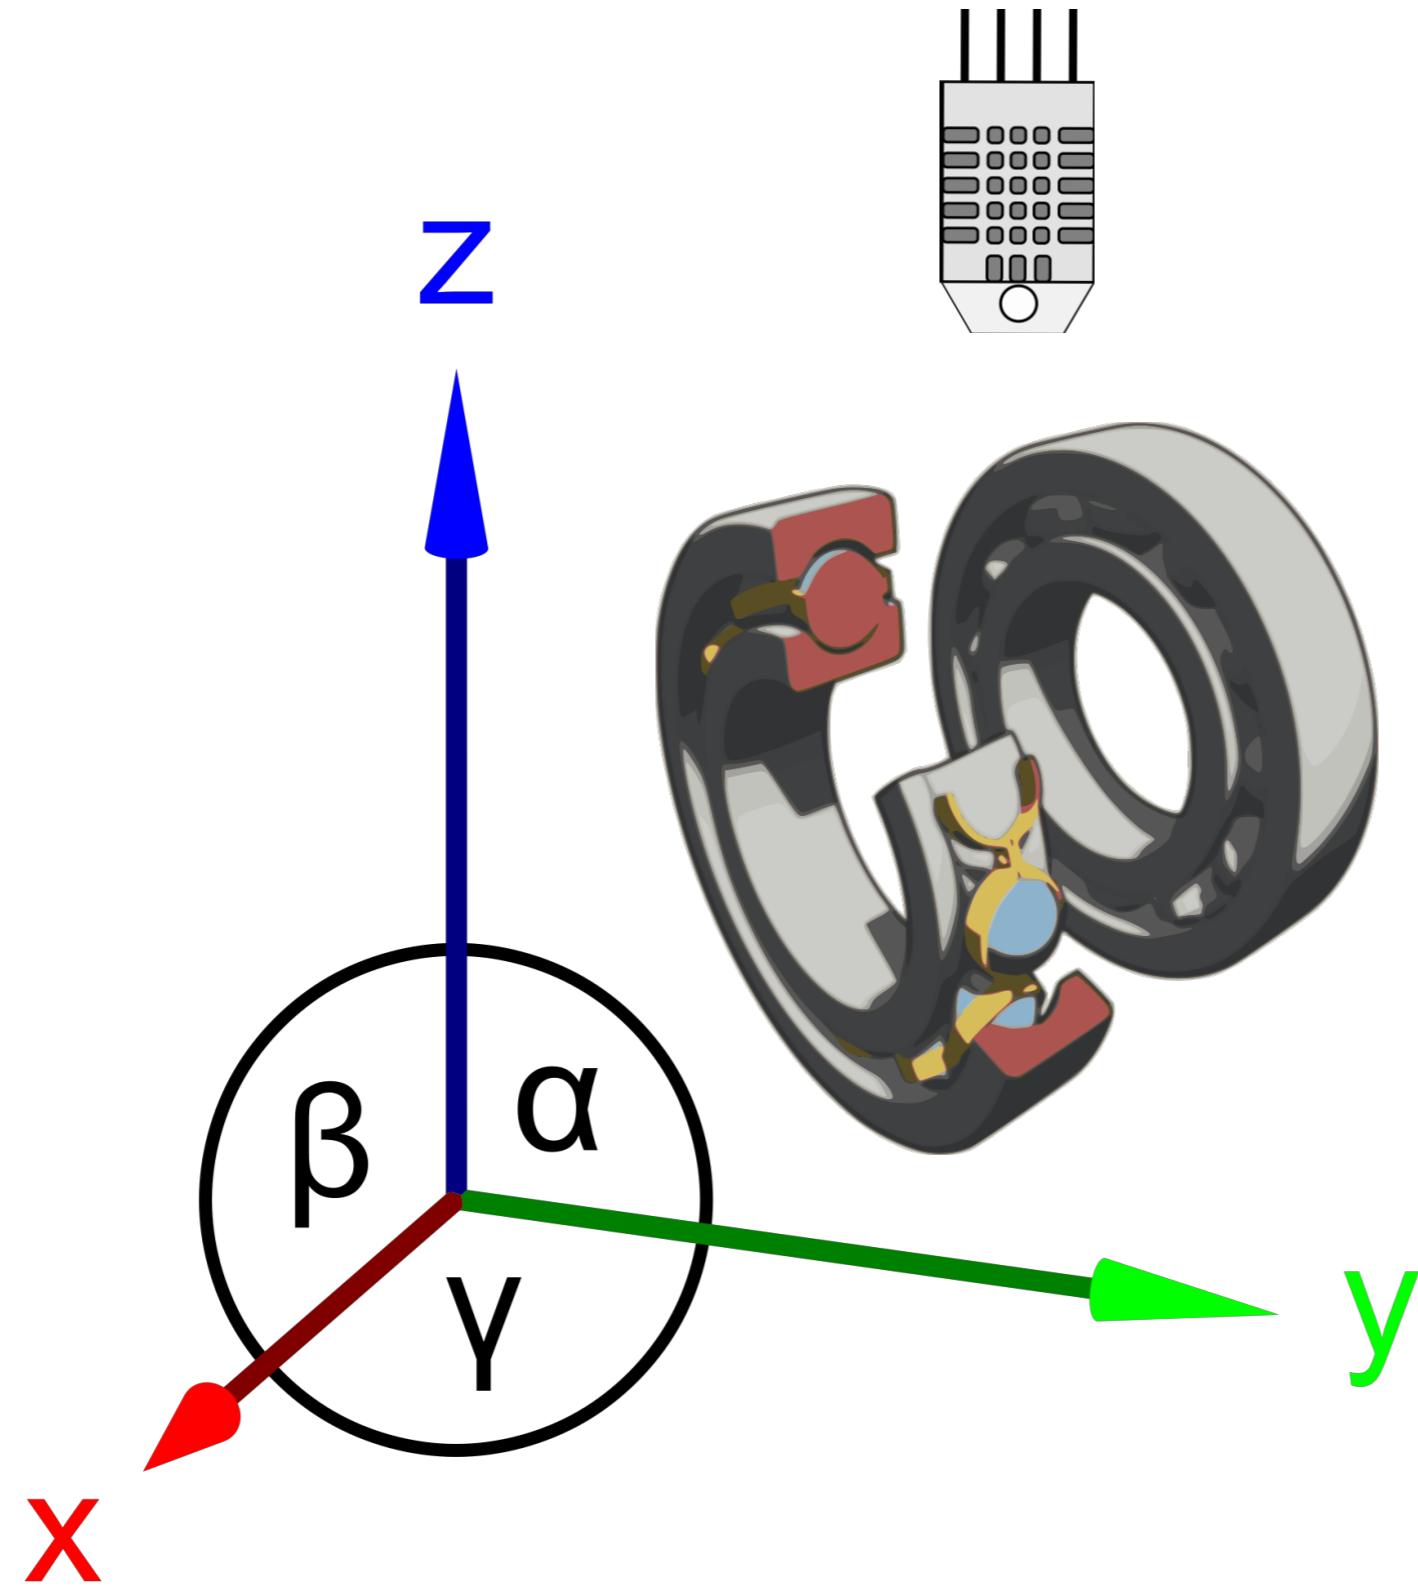
CS 229 Final Project Report

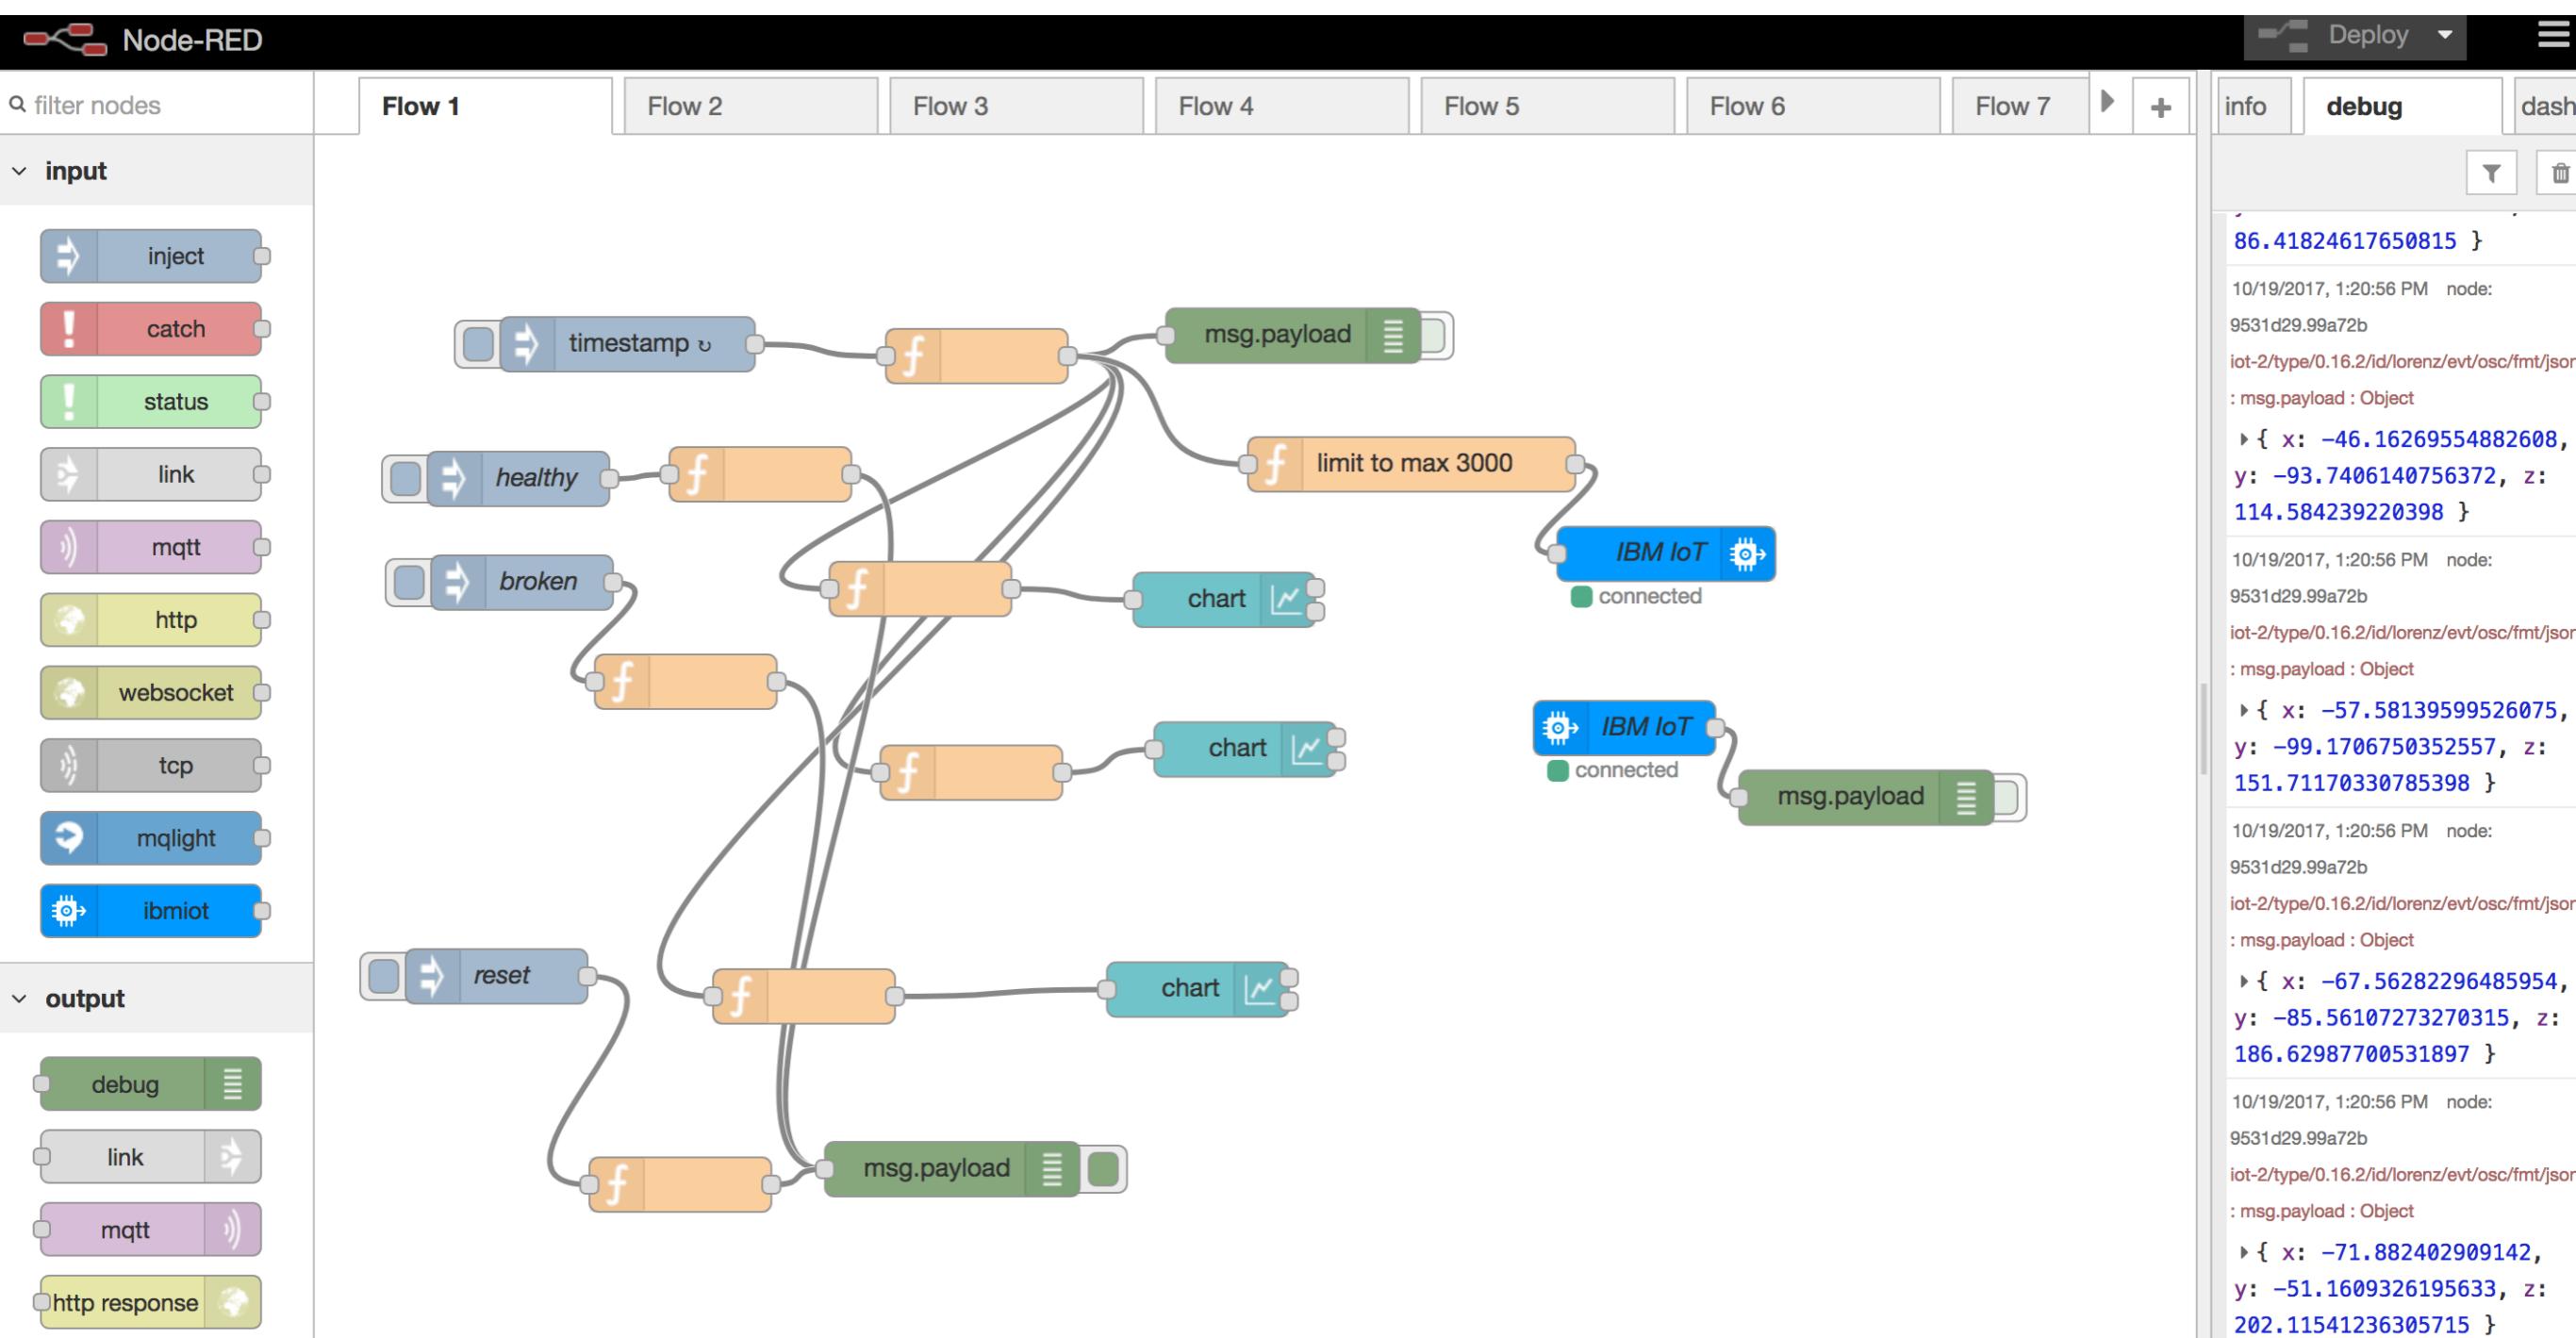
Enzo Busseti, Ian Osband, Scott Wong

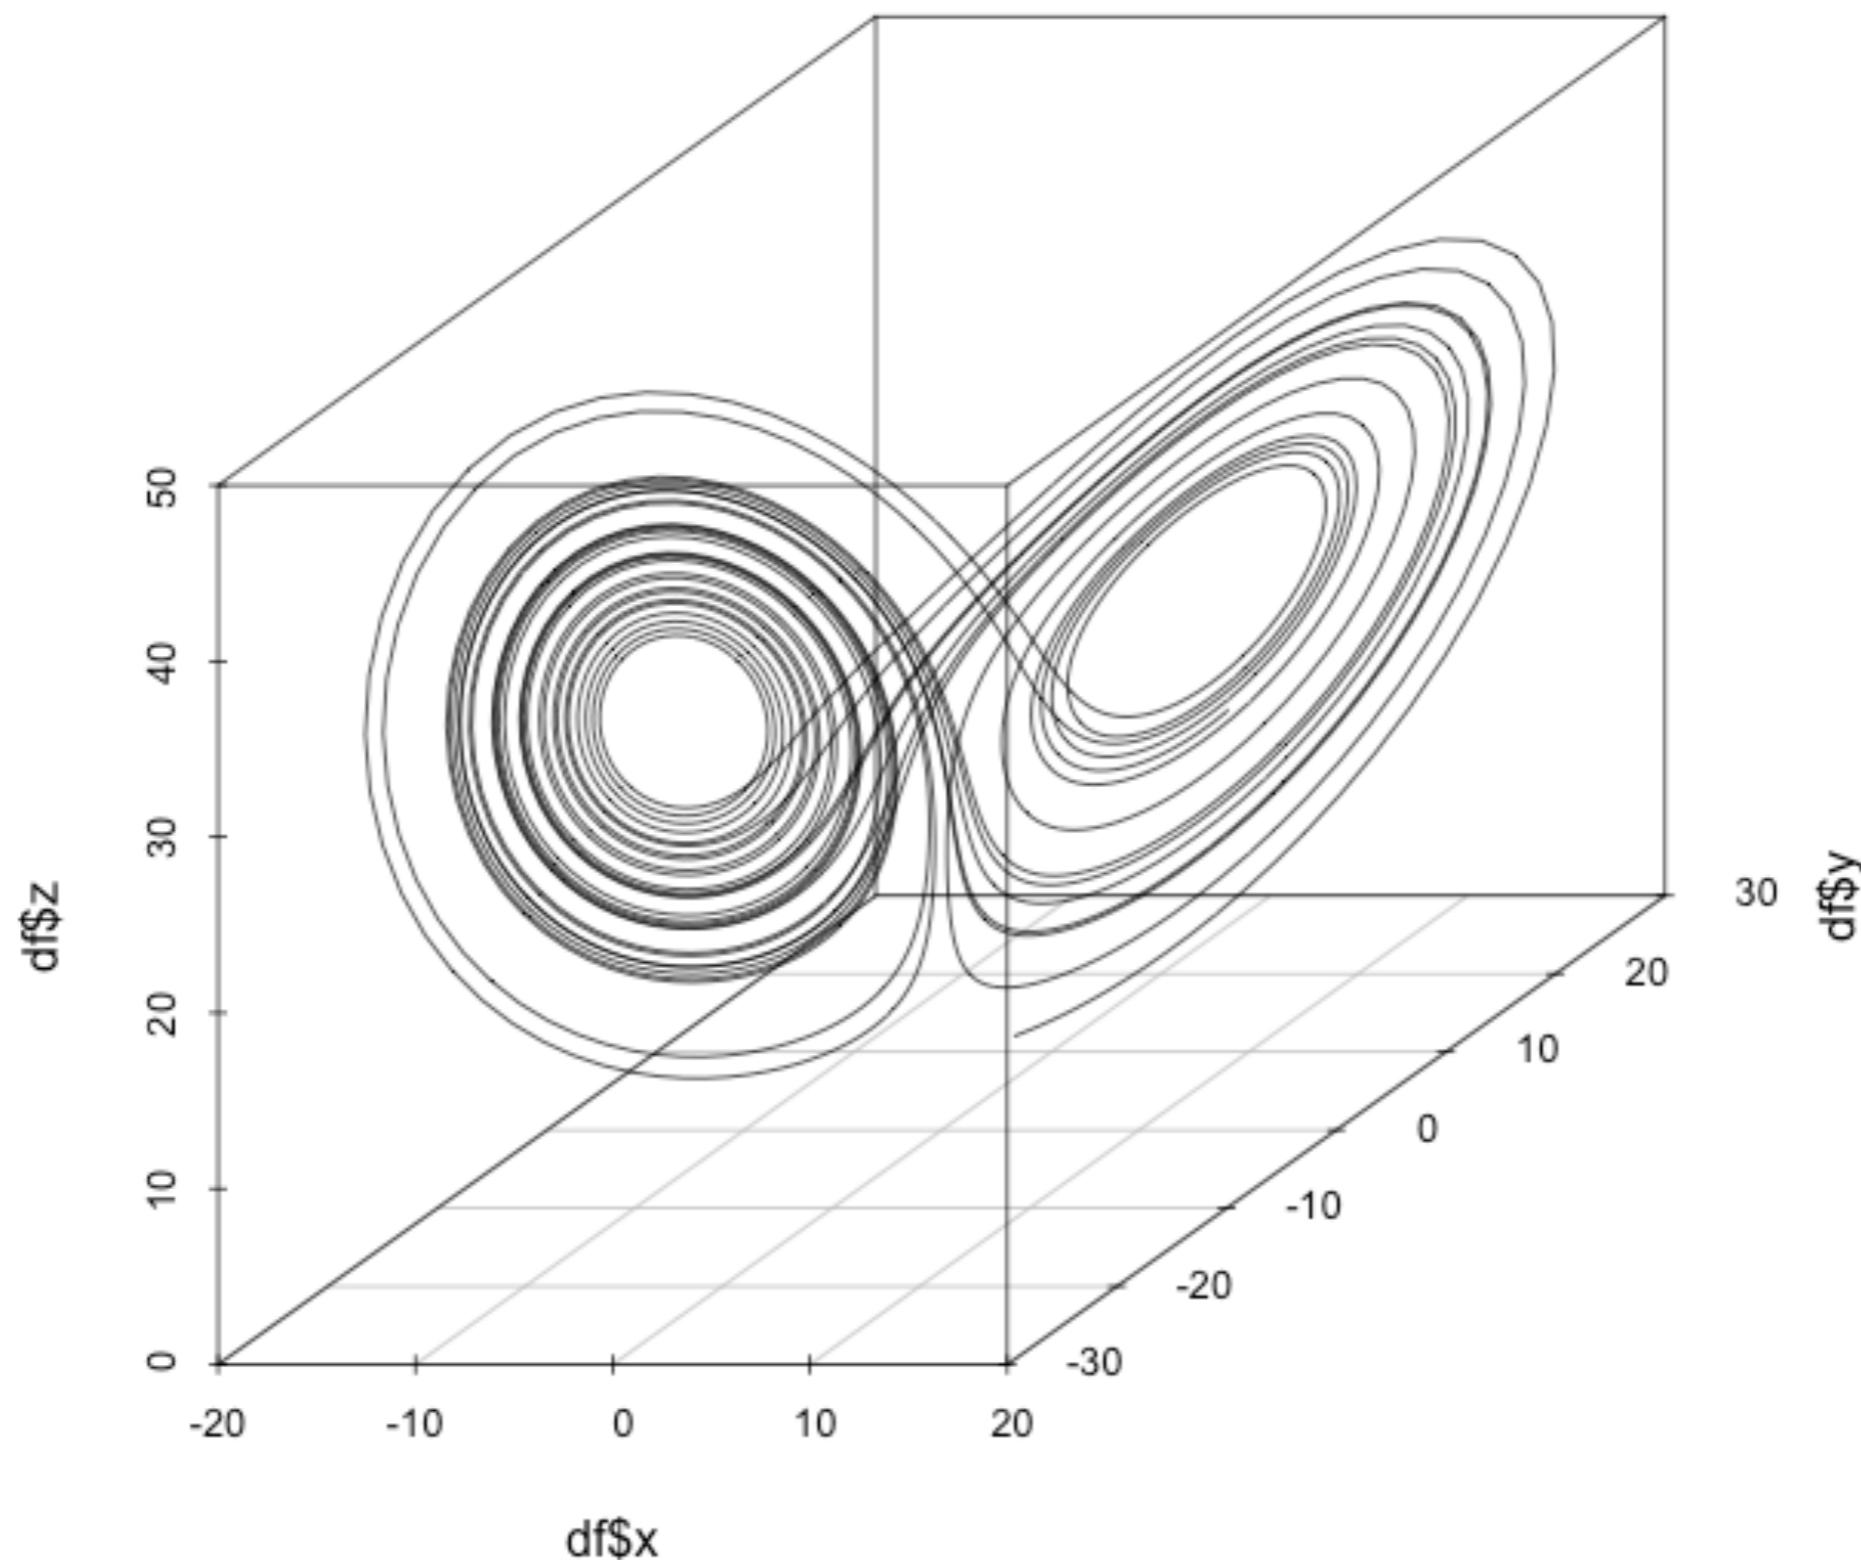


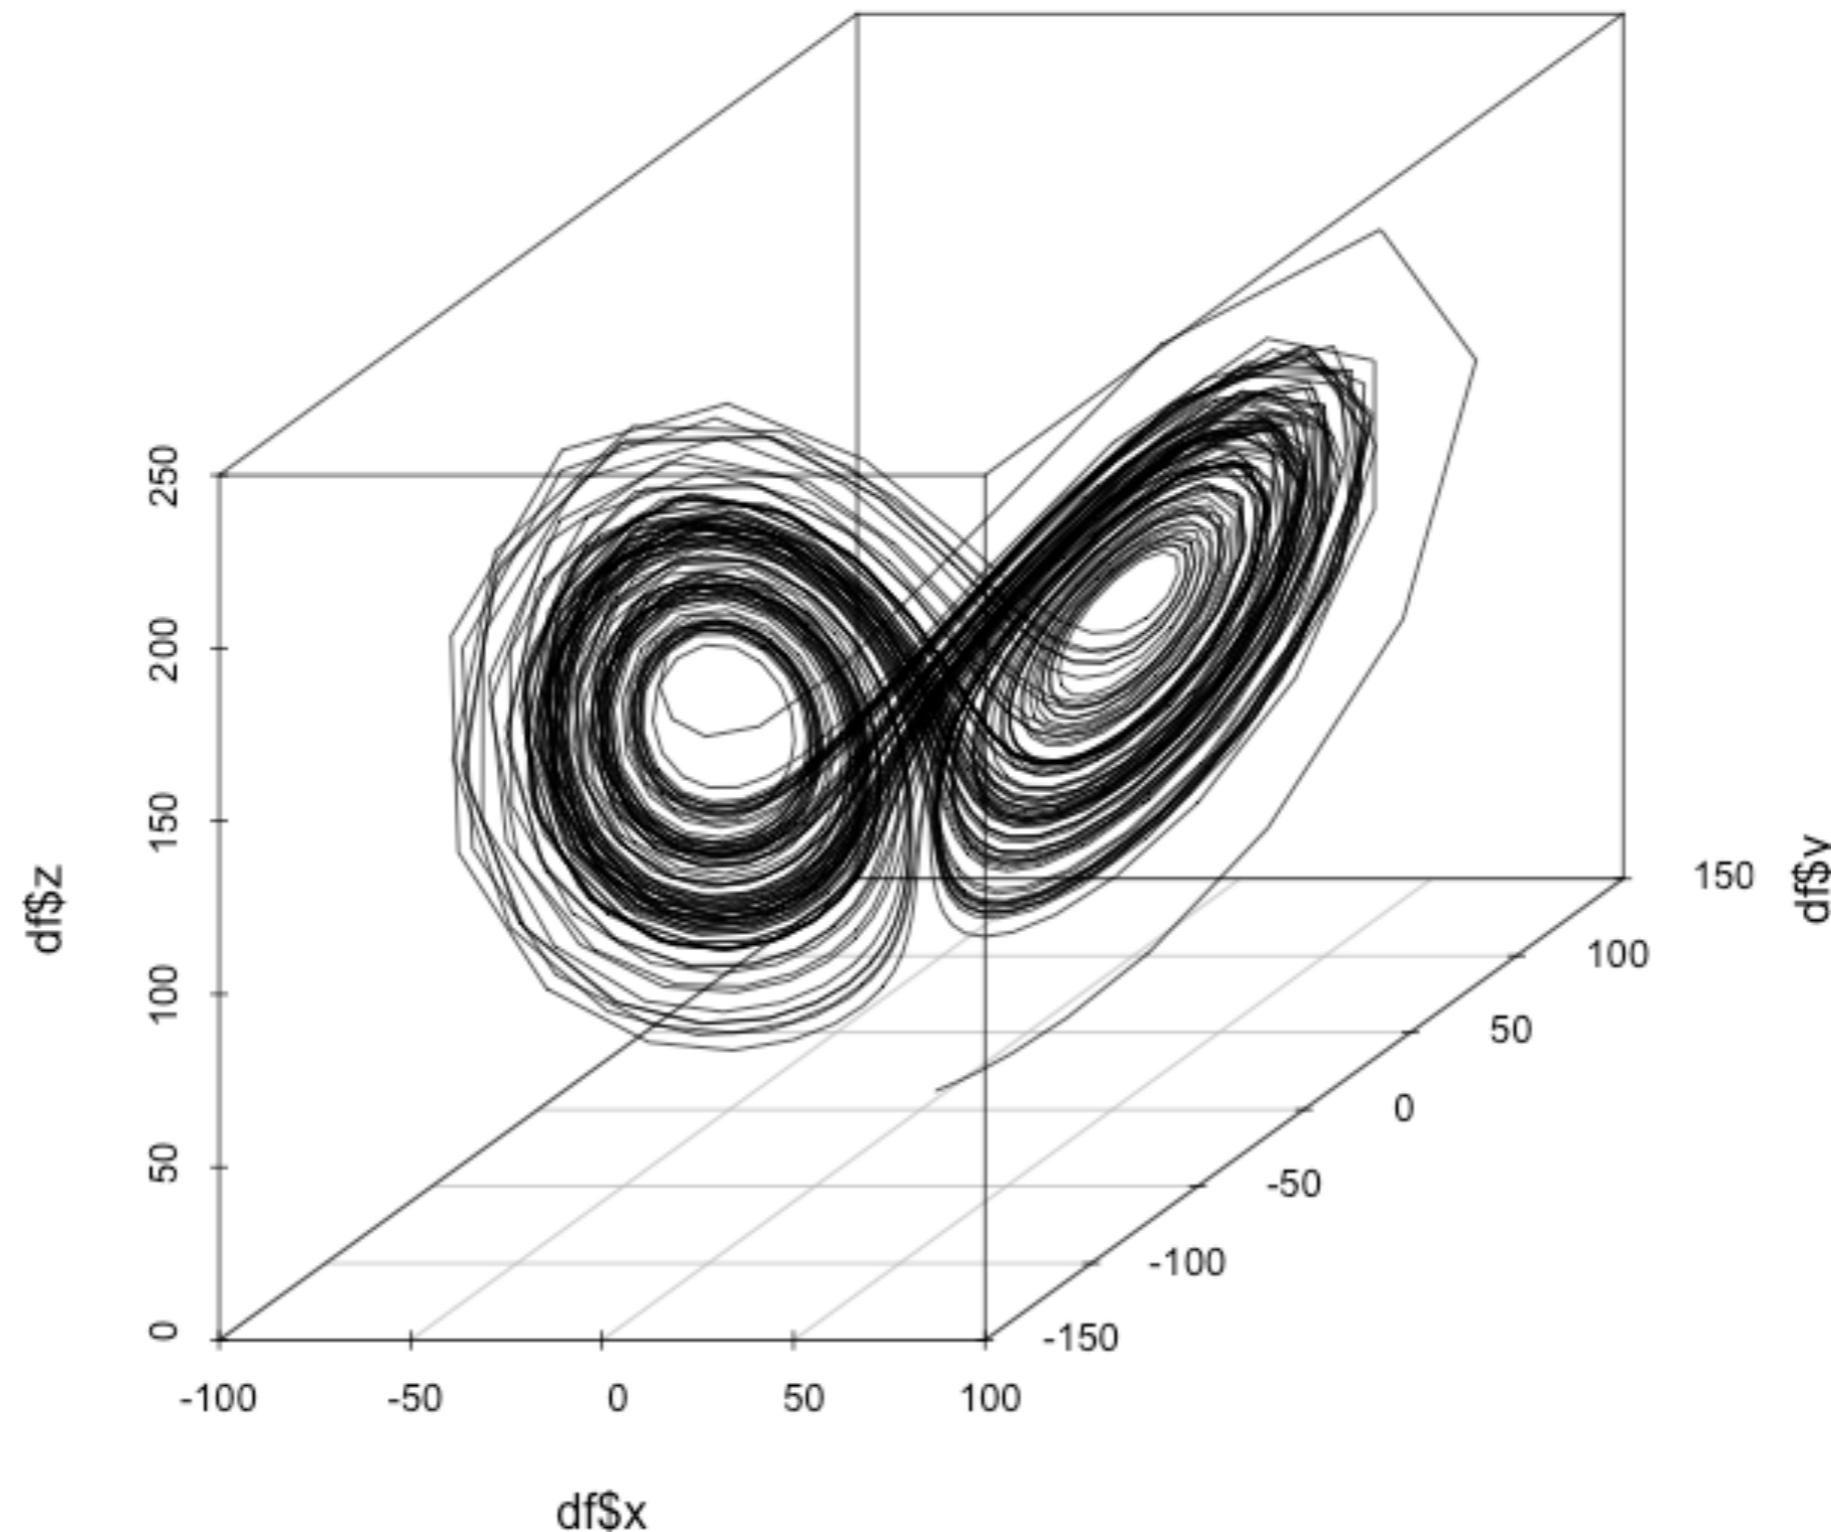
Learning Method	RMSE	% RMSE
Kernelized Regression	1,540	8.3%
Frequency NN	1,251	6.7%
Deep Feedforward NN	1,103	5.9%
Deep Recurrent NN	530	2.8%

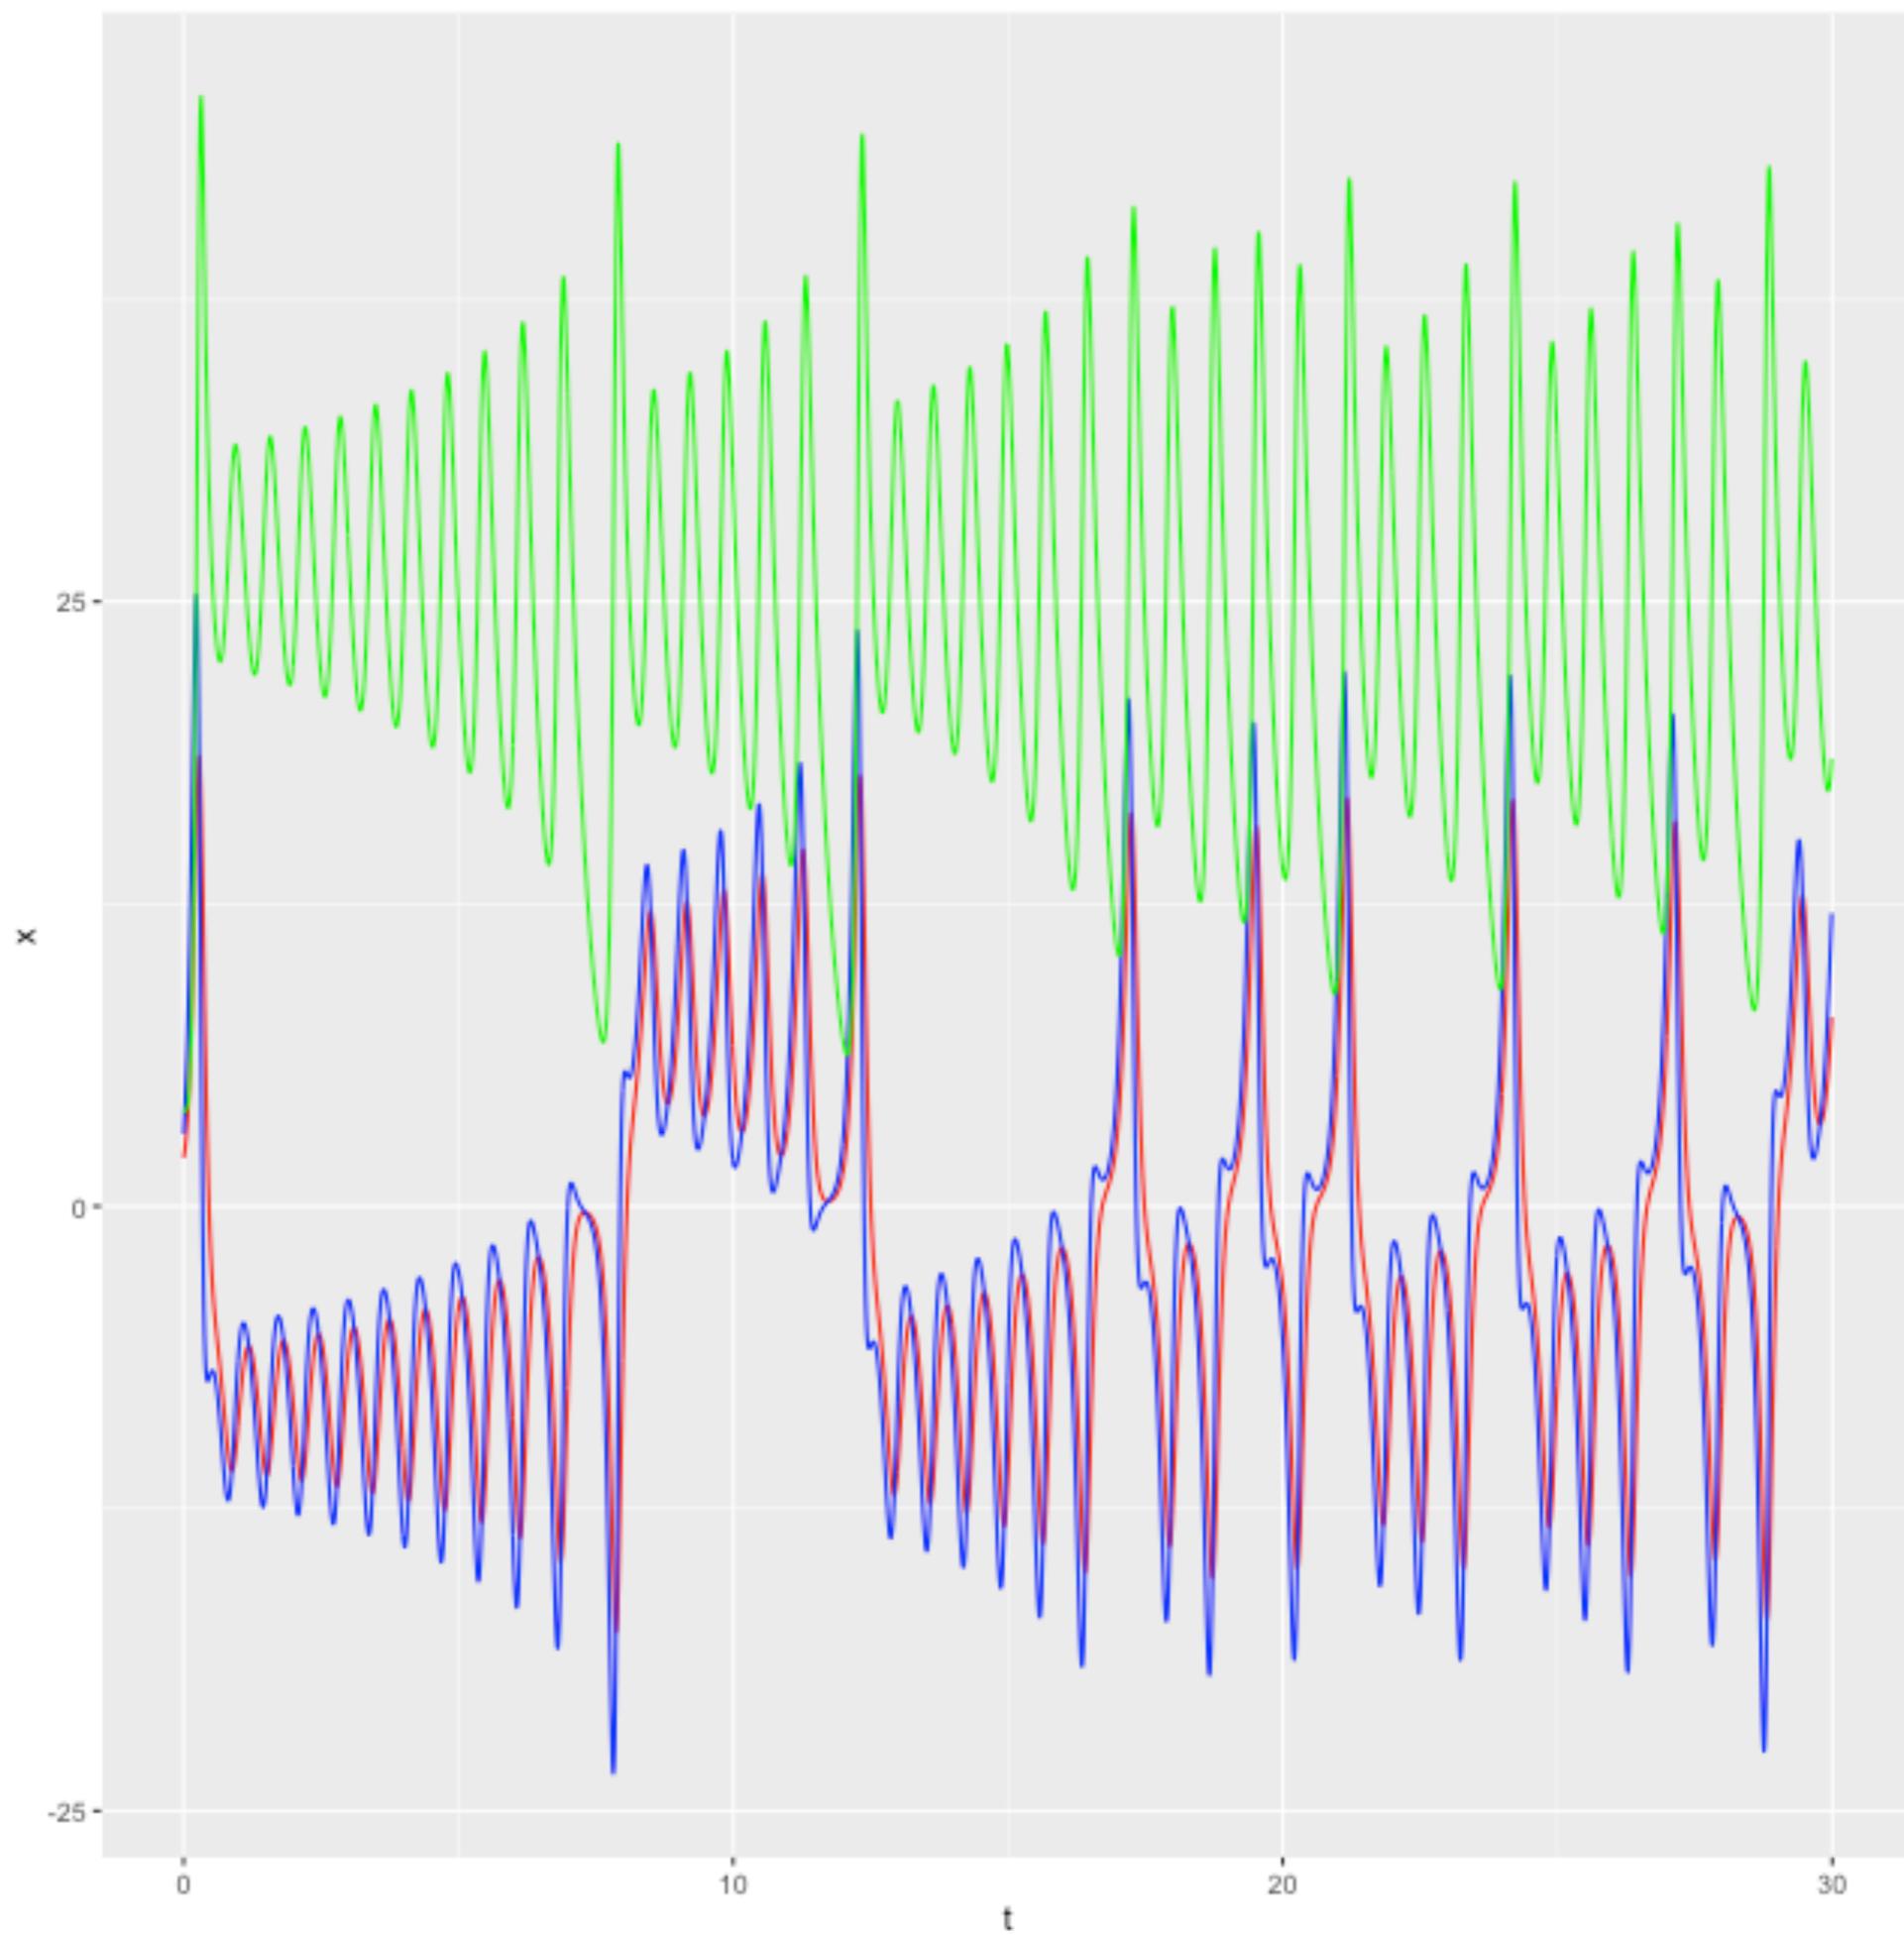


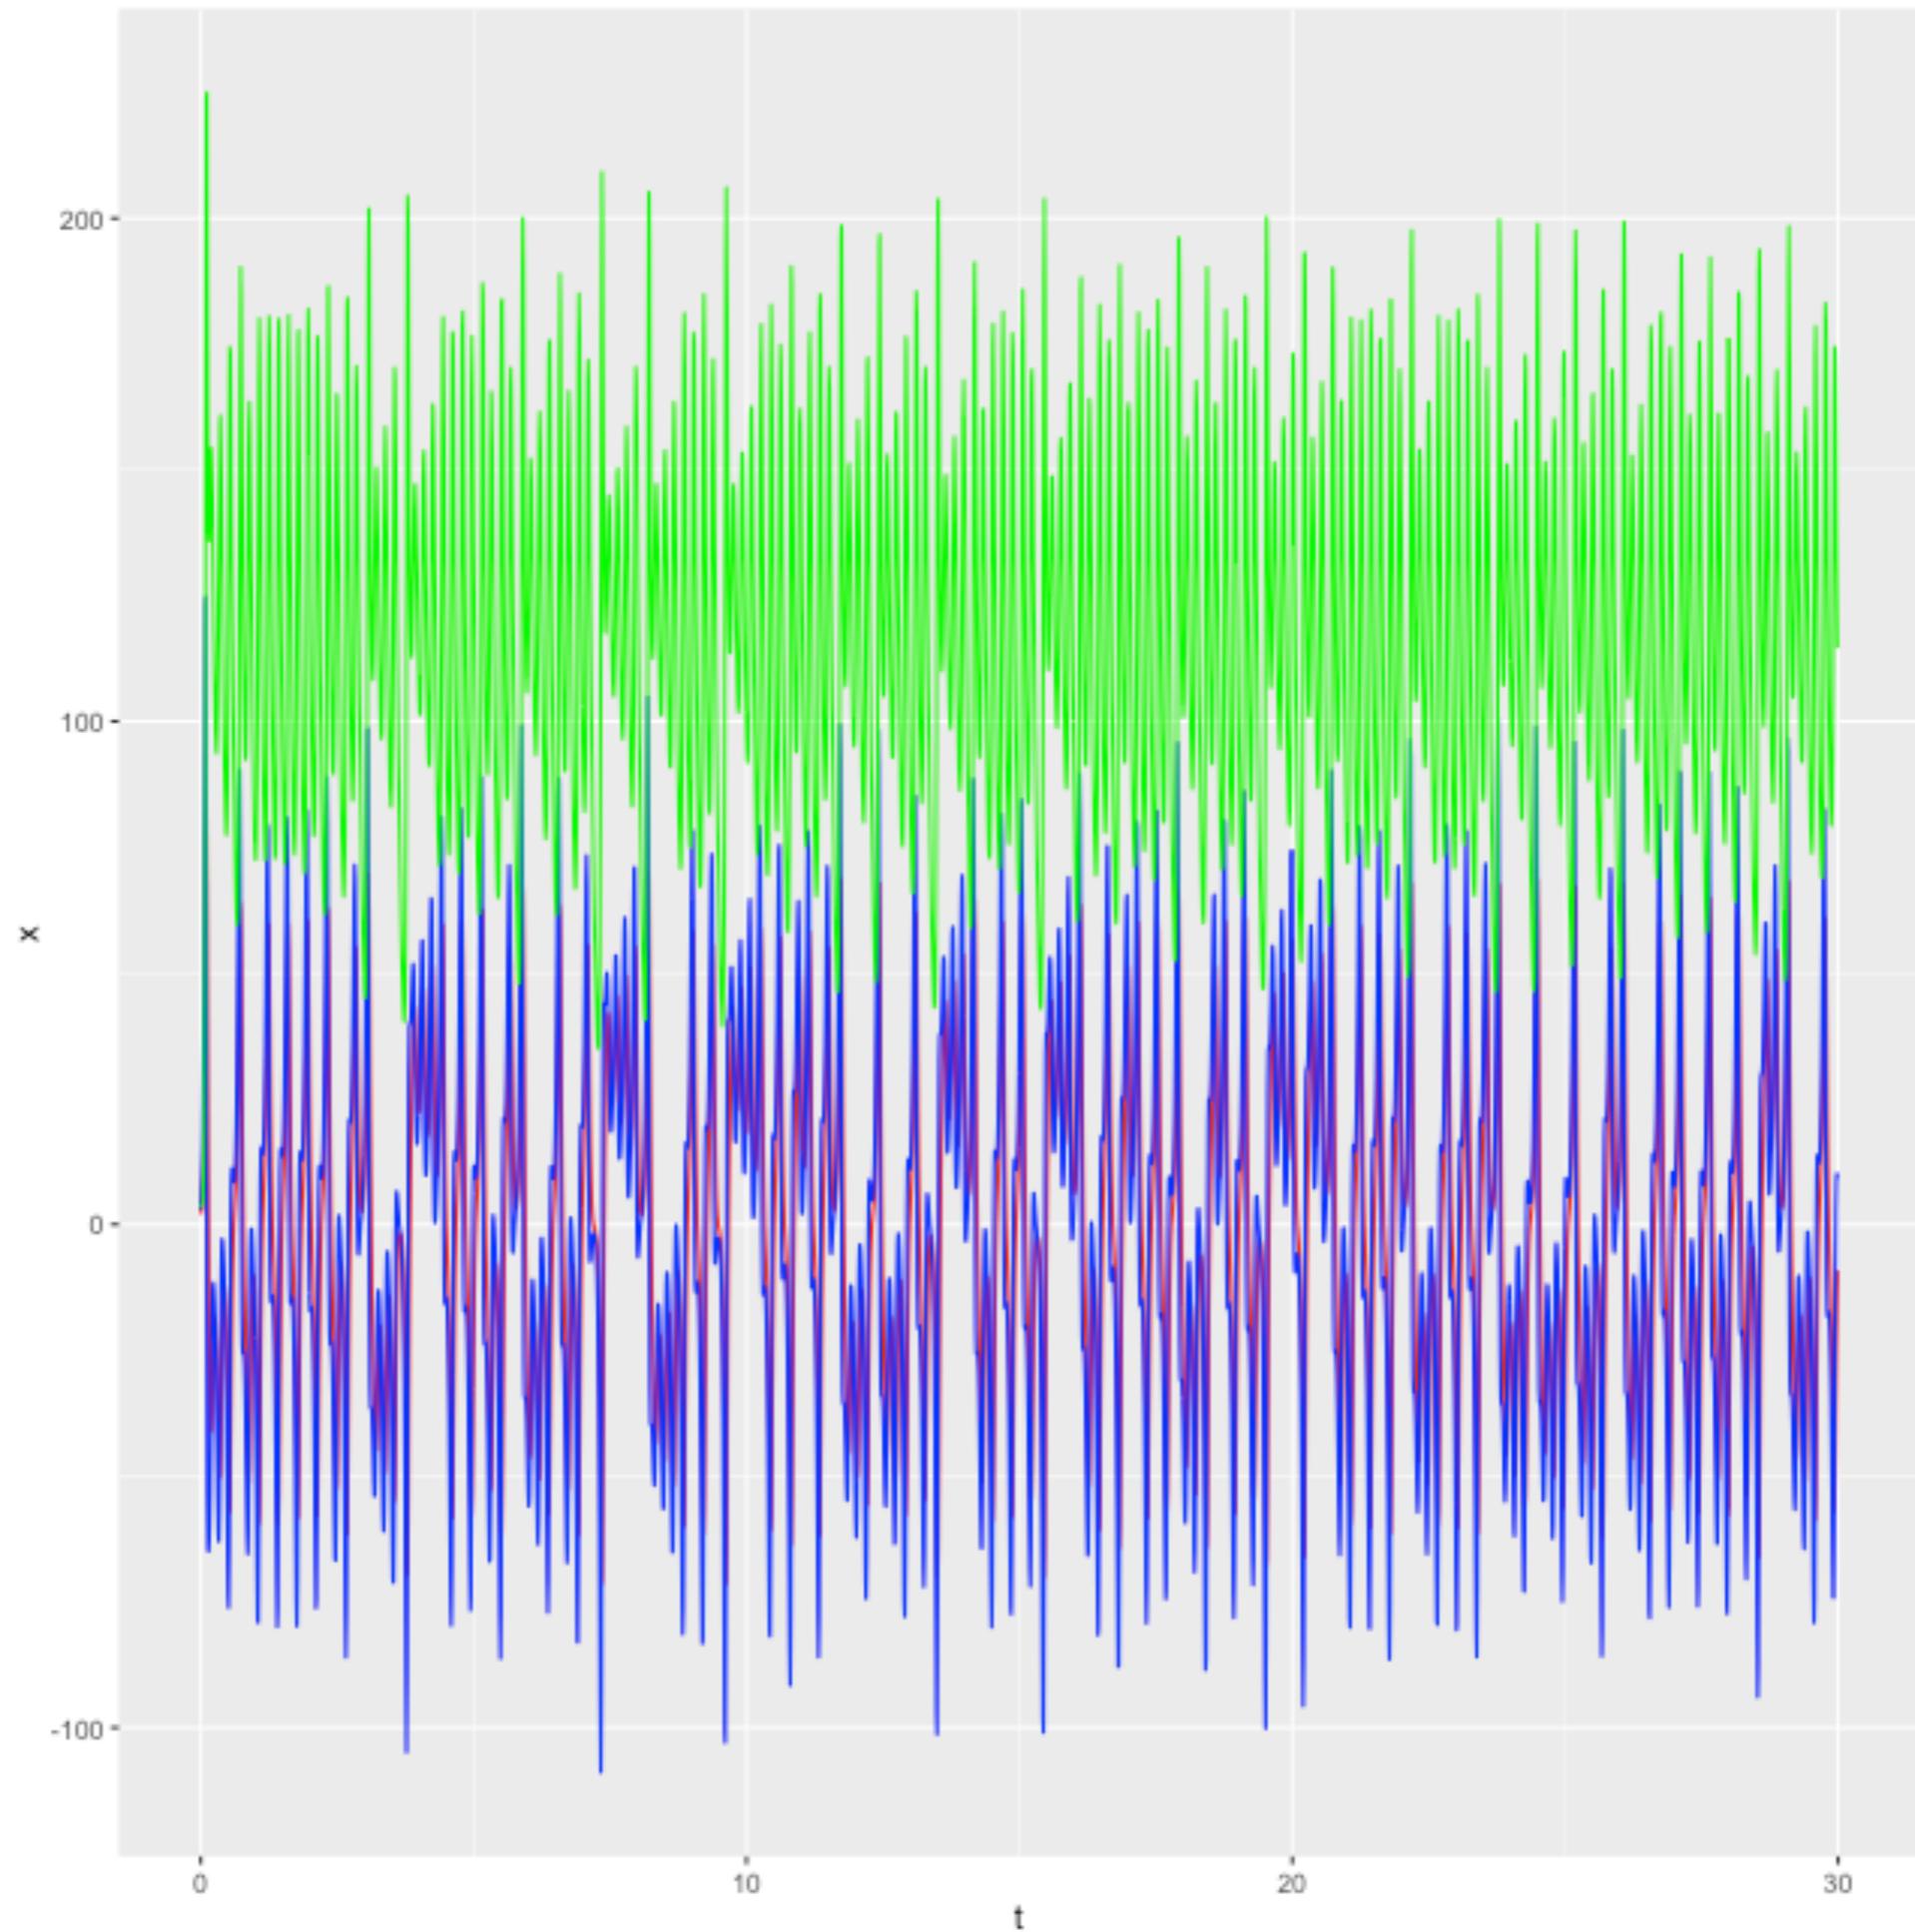


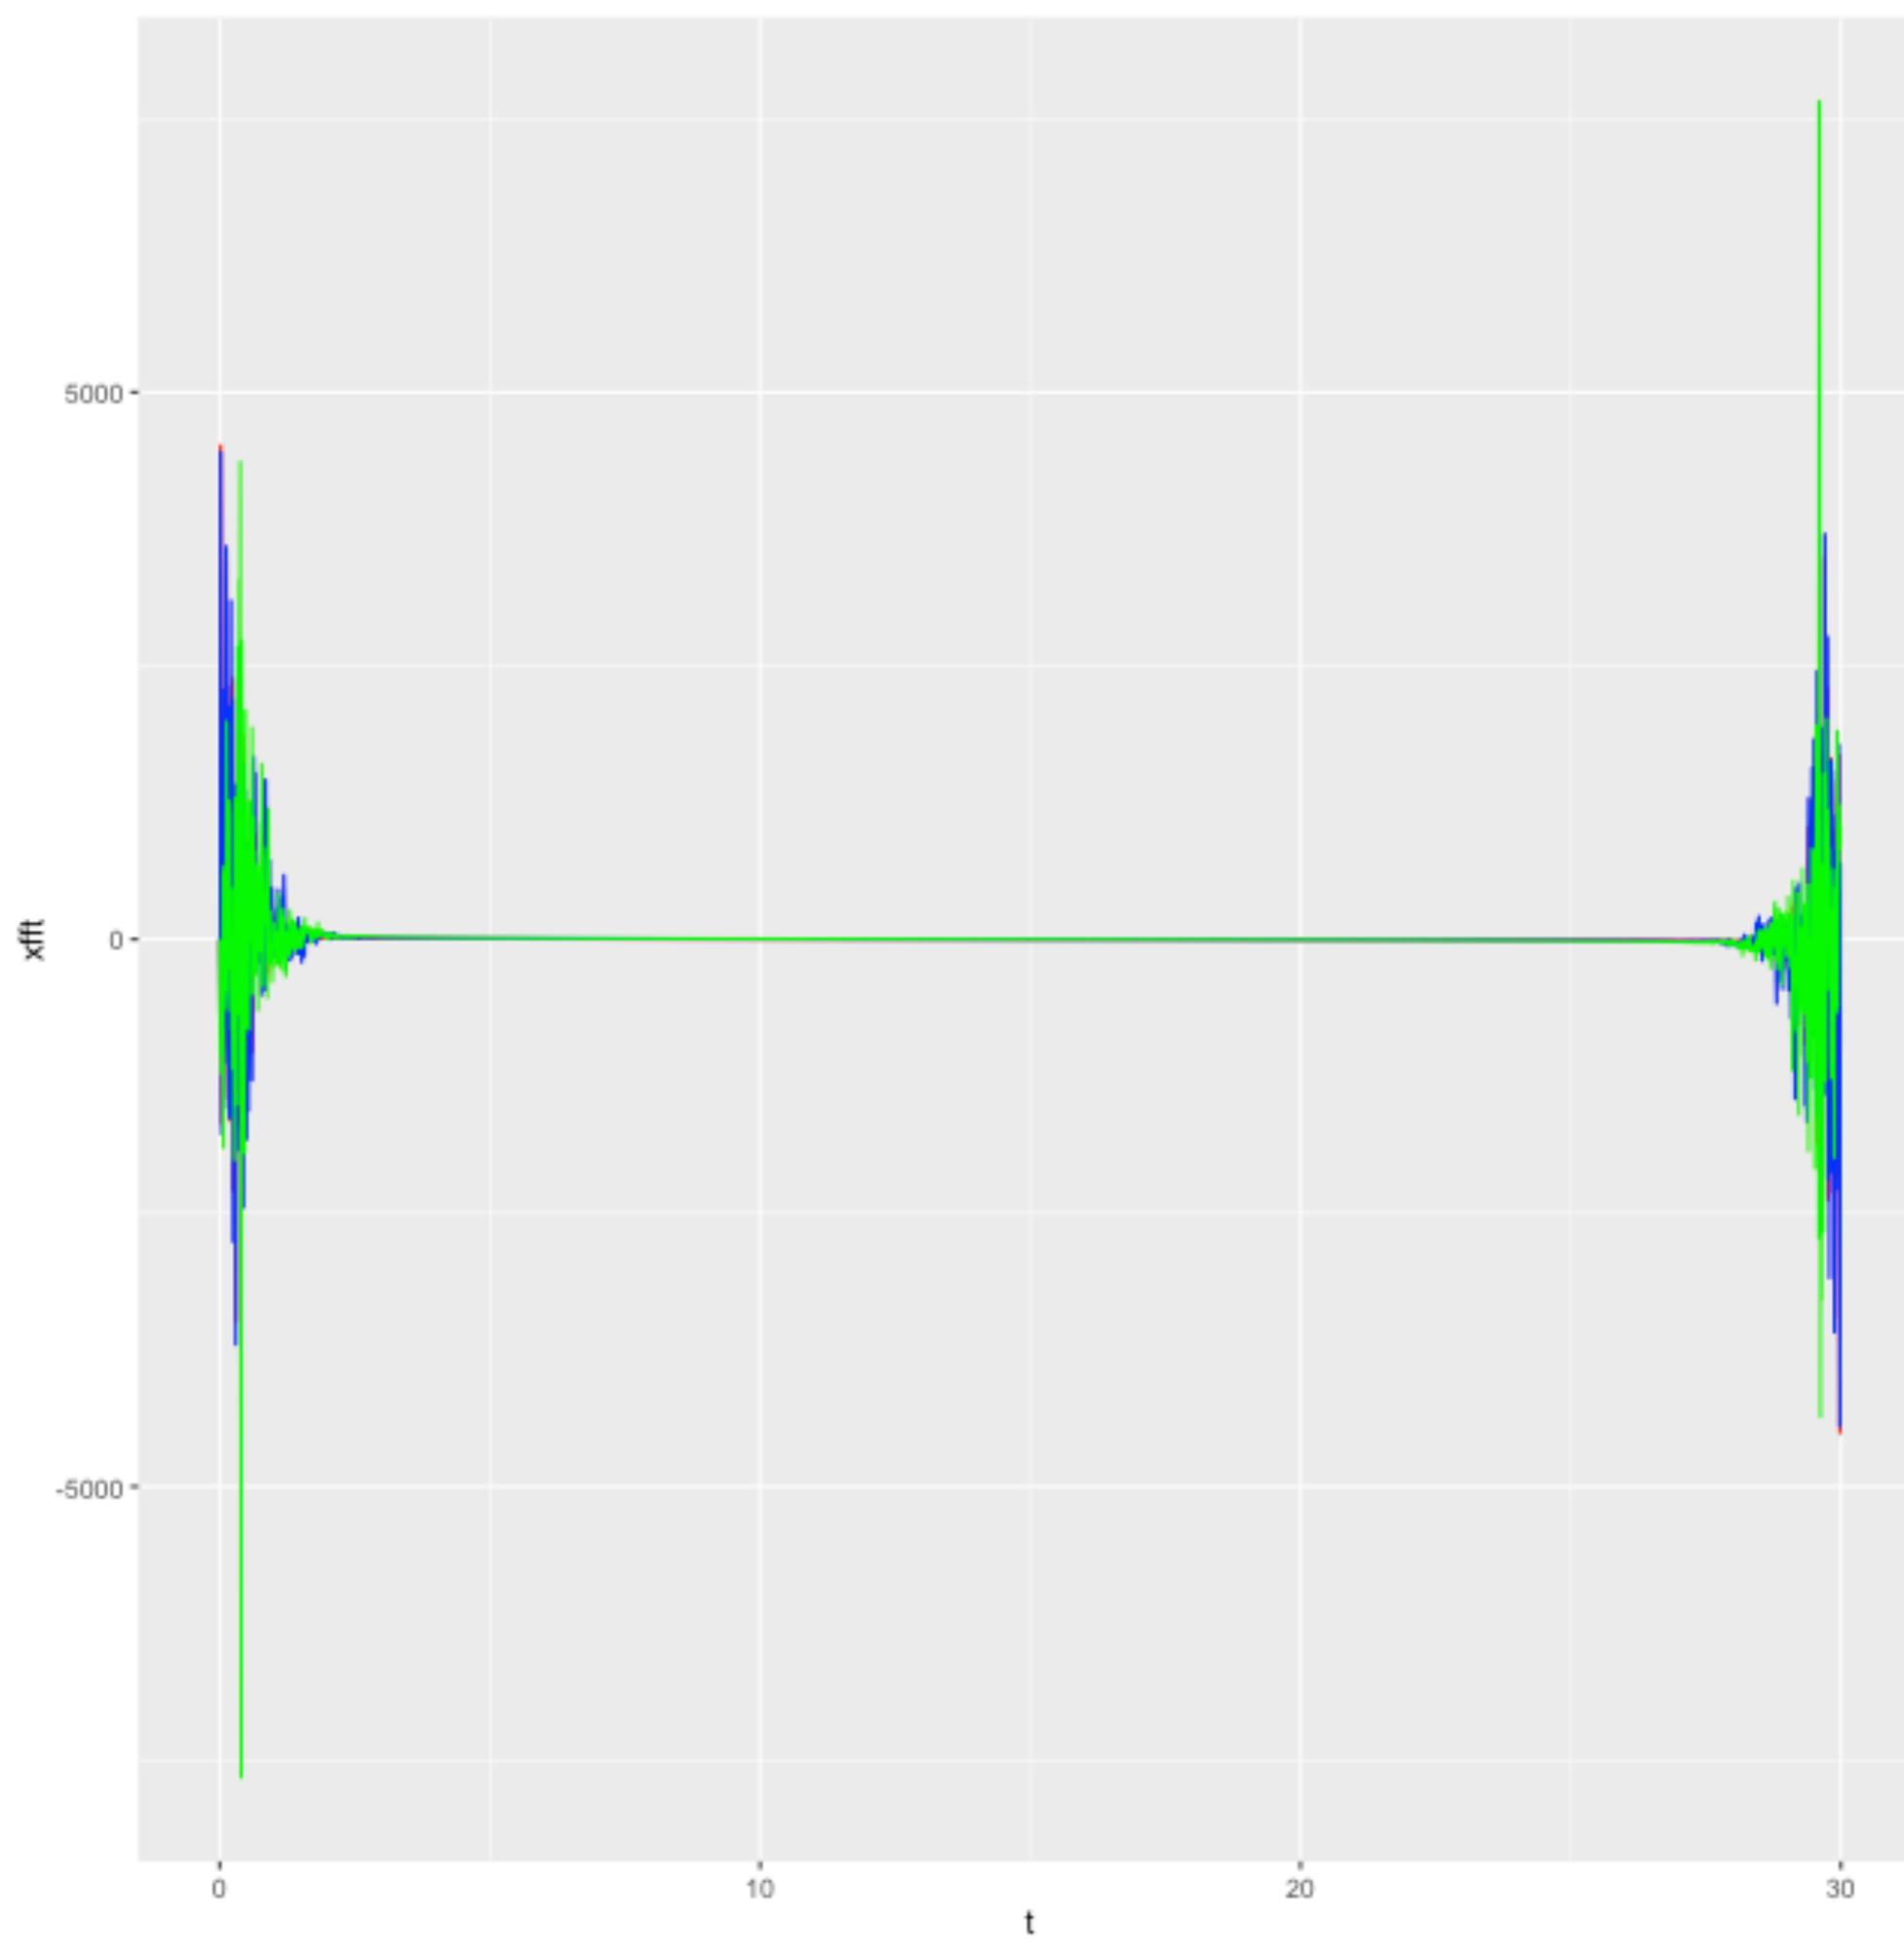


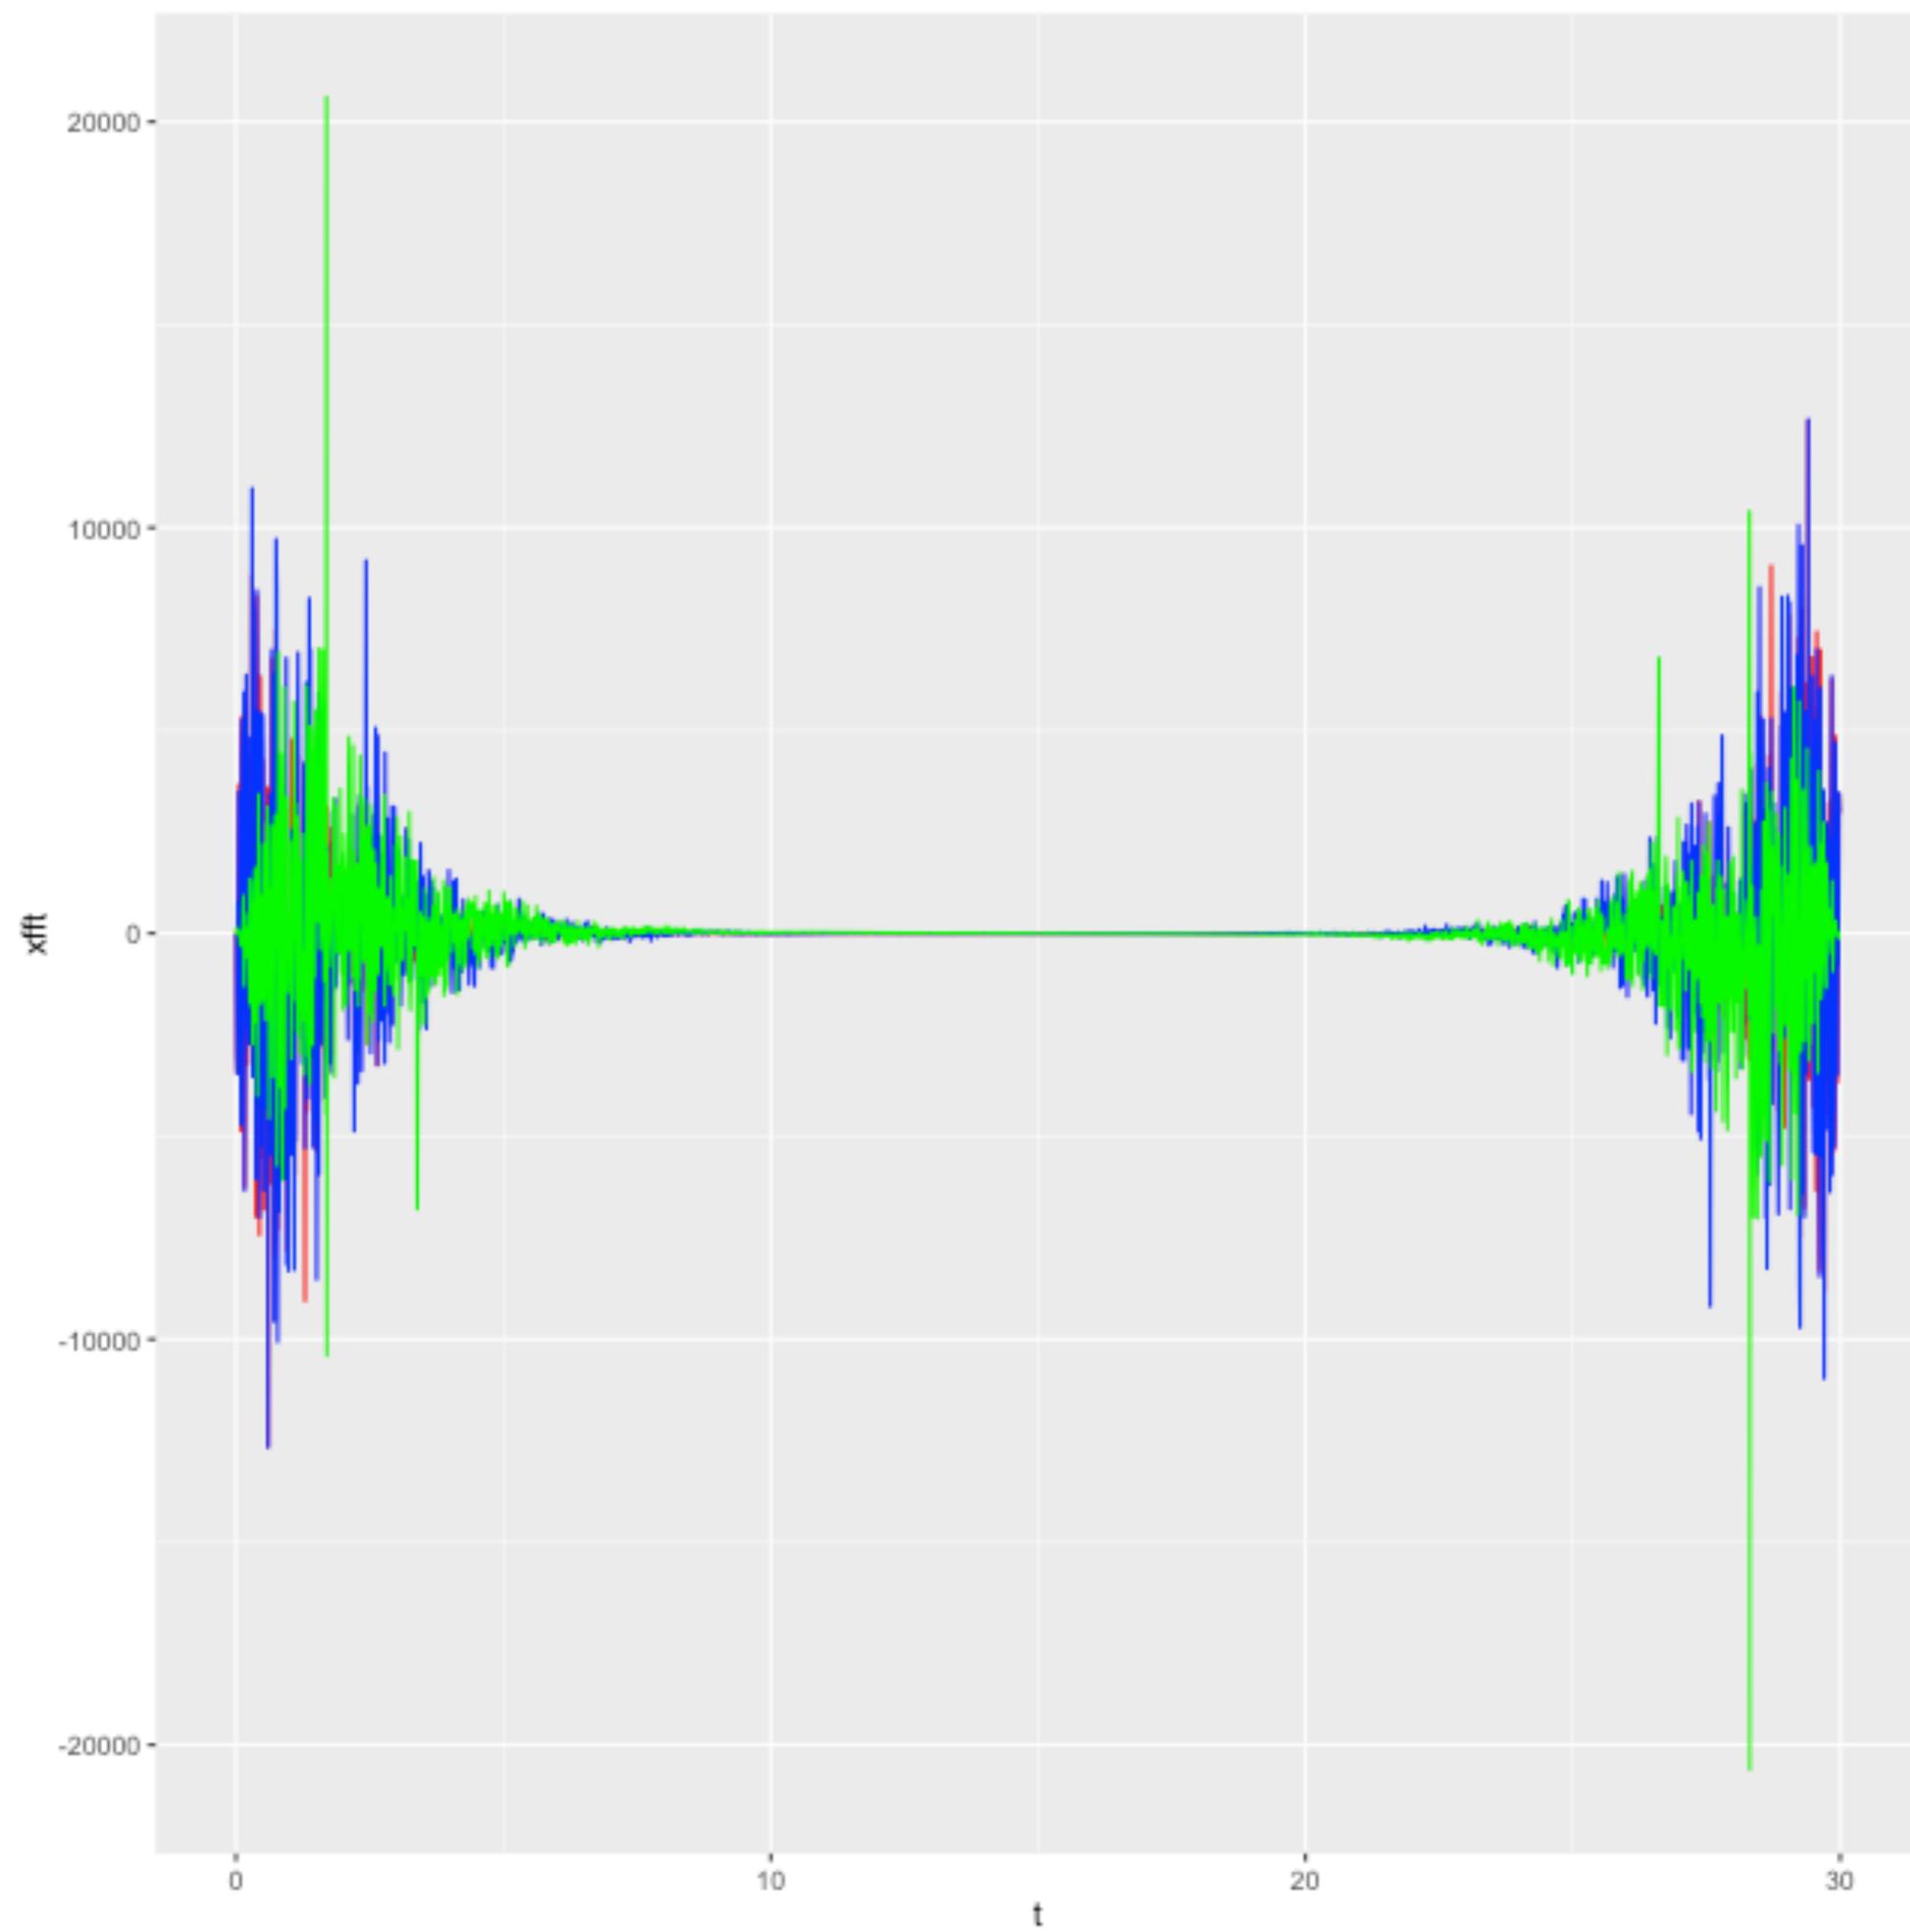












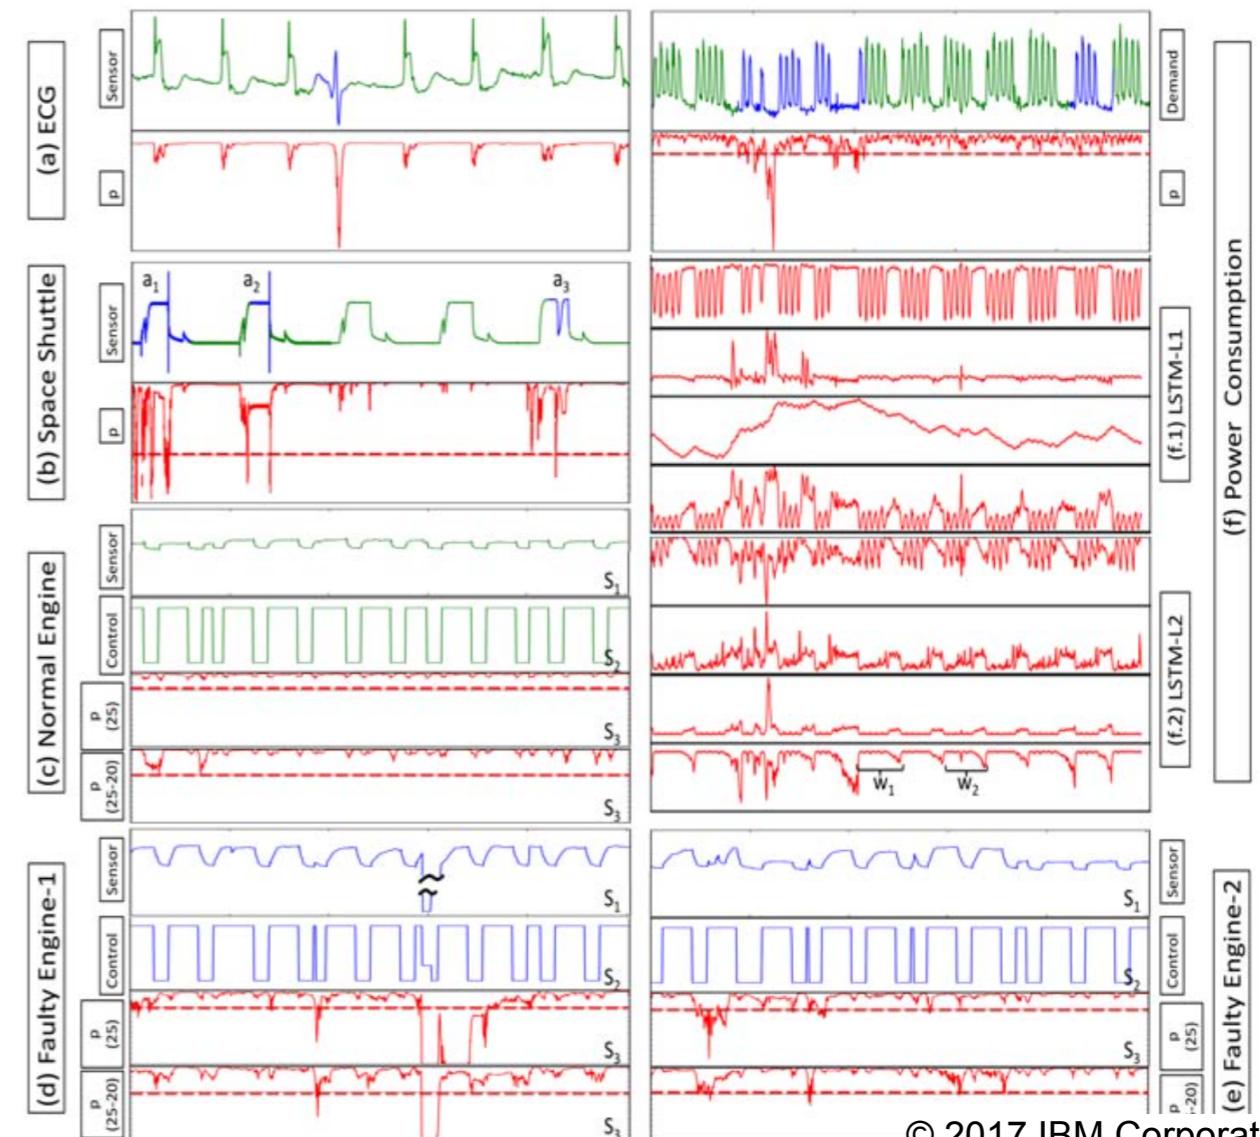
Long Short Term Memory Networks for Anomaly Detection in Time Series

Pankaj Malhotra¹, Lovekesh Vig², Gautam Shroff¹, Puneet Agarwal¹

1- TCS Research, Delhi, India

2- Jawaharlal Nehru University, New Delhi, India

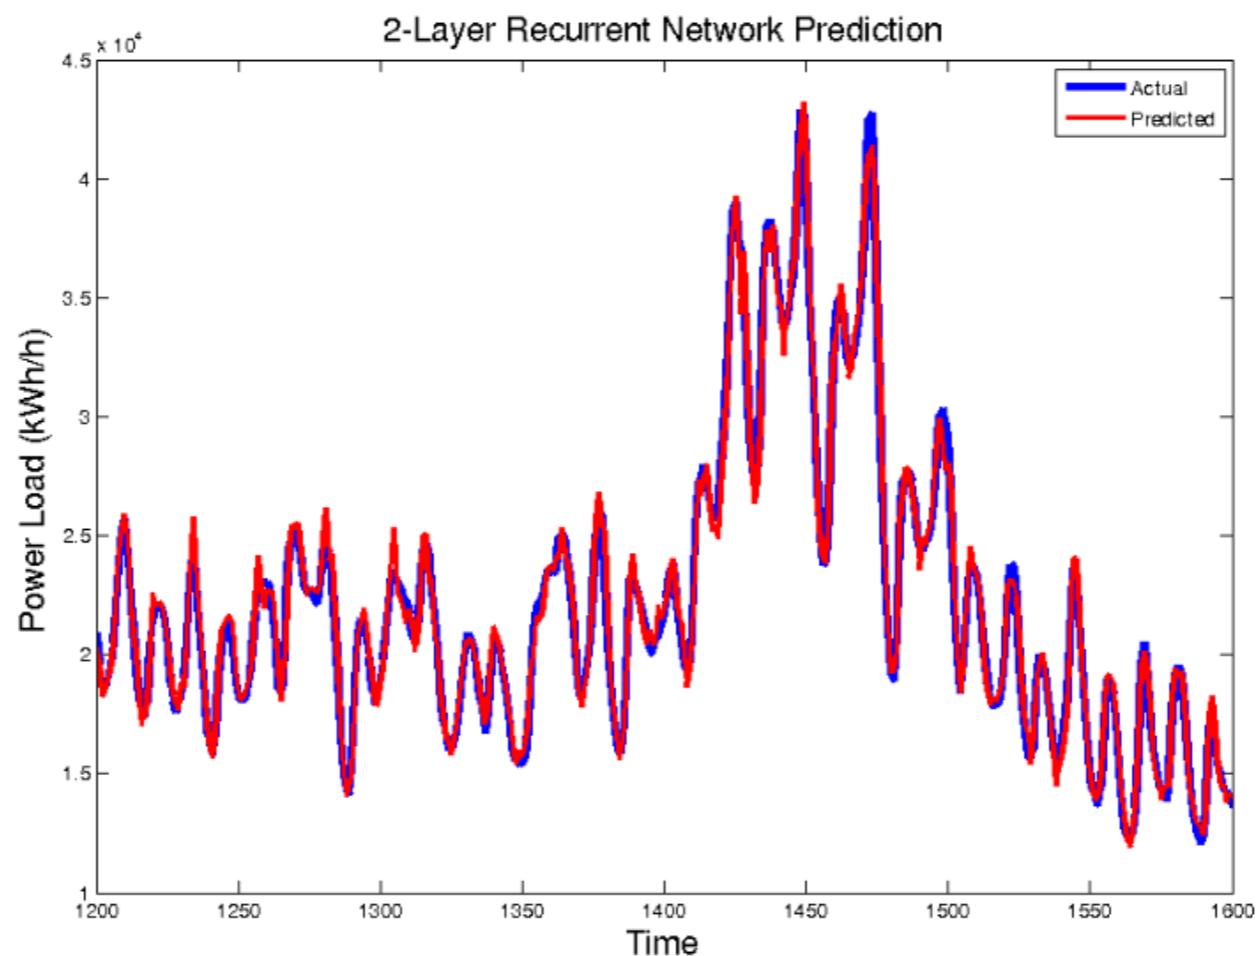
- Outperformed traditional methods, such as
 - cumulative sum (CUSUM)
 - exponentially weighted moving average (EWMA)
- Hidden Markov Models (HMM)
- Learned what “Normal” is
- Raised error if time series pattern haven't been seen before



Deep Learning for Time Series Modeling

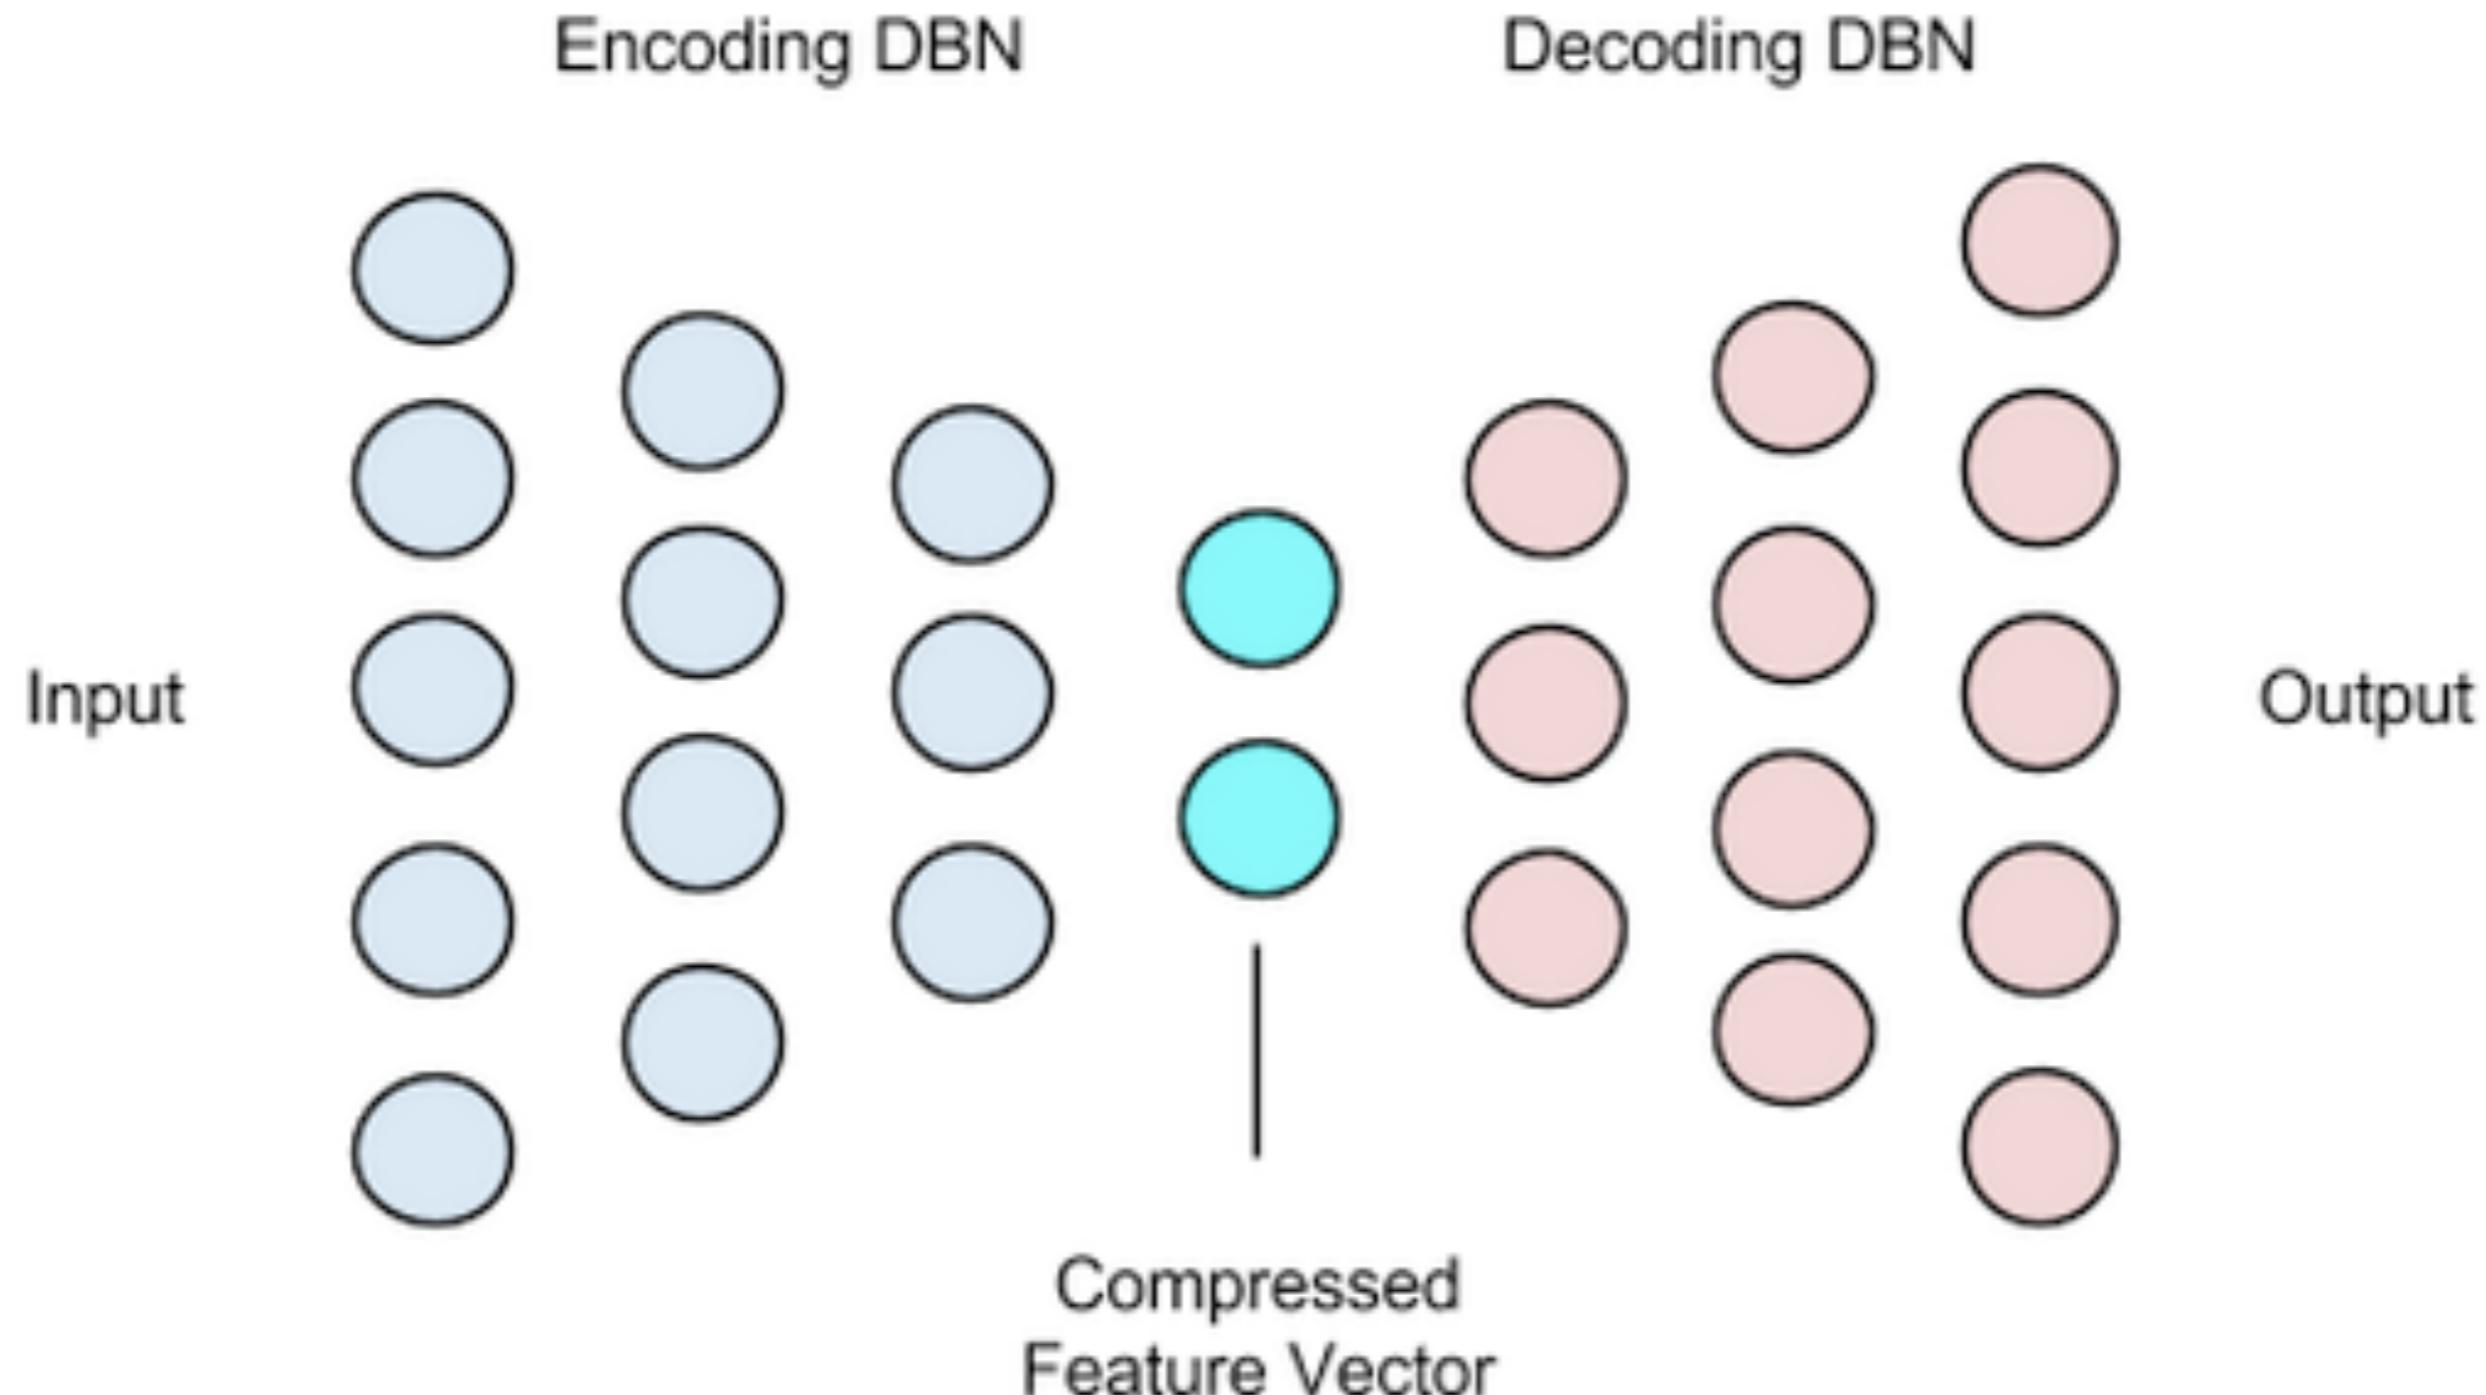
CS 229 Final Project Report

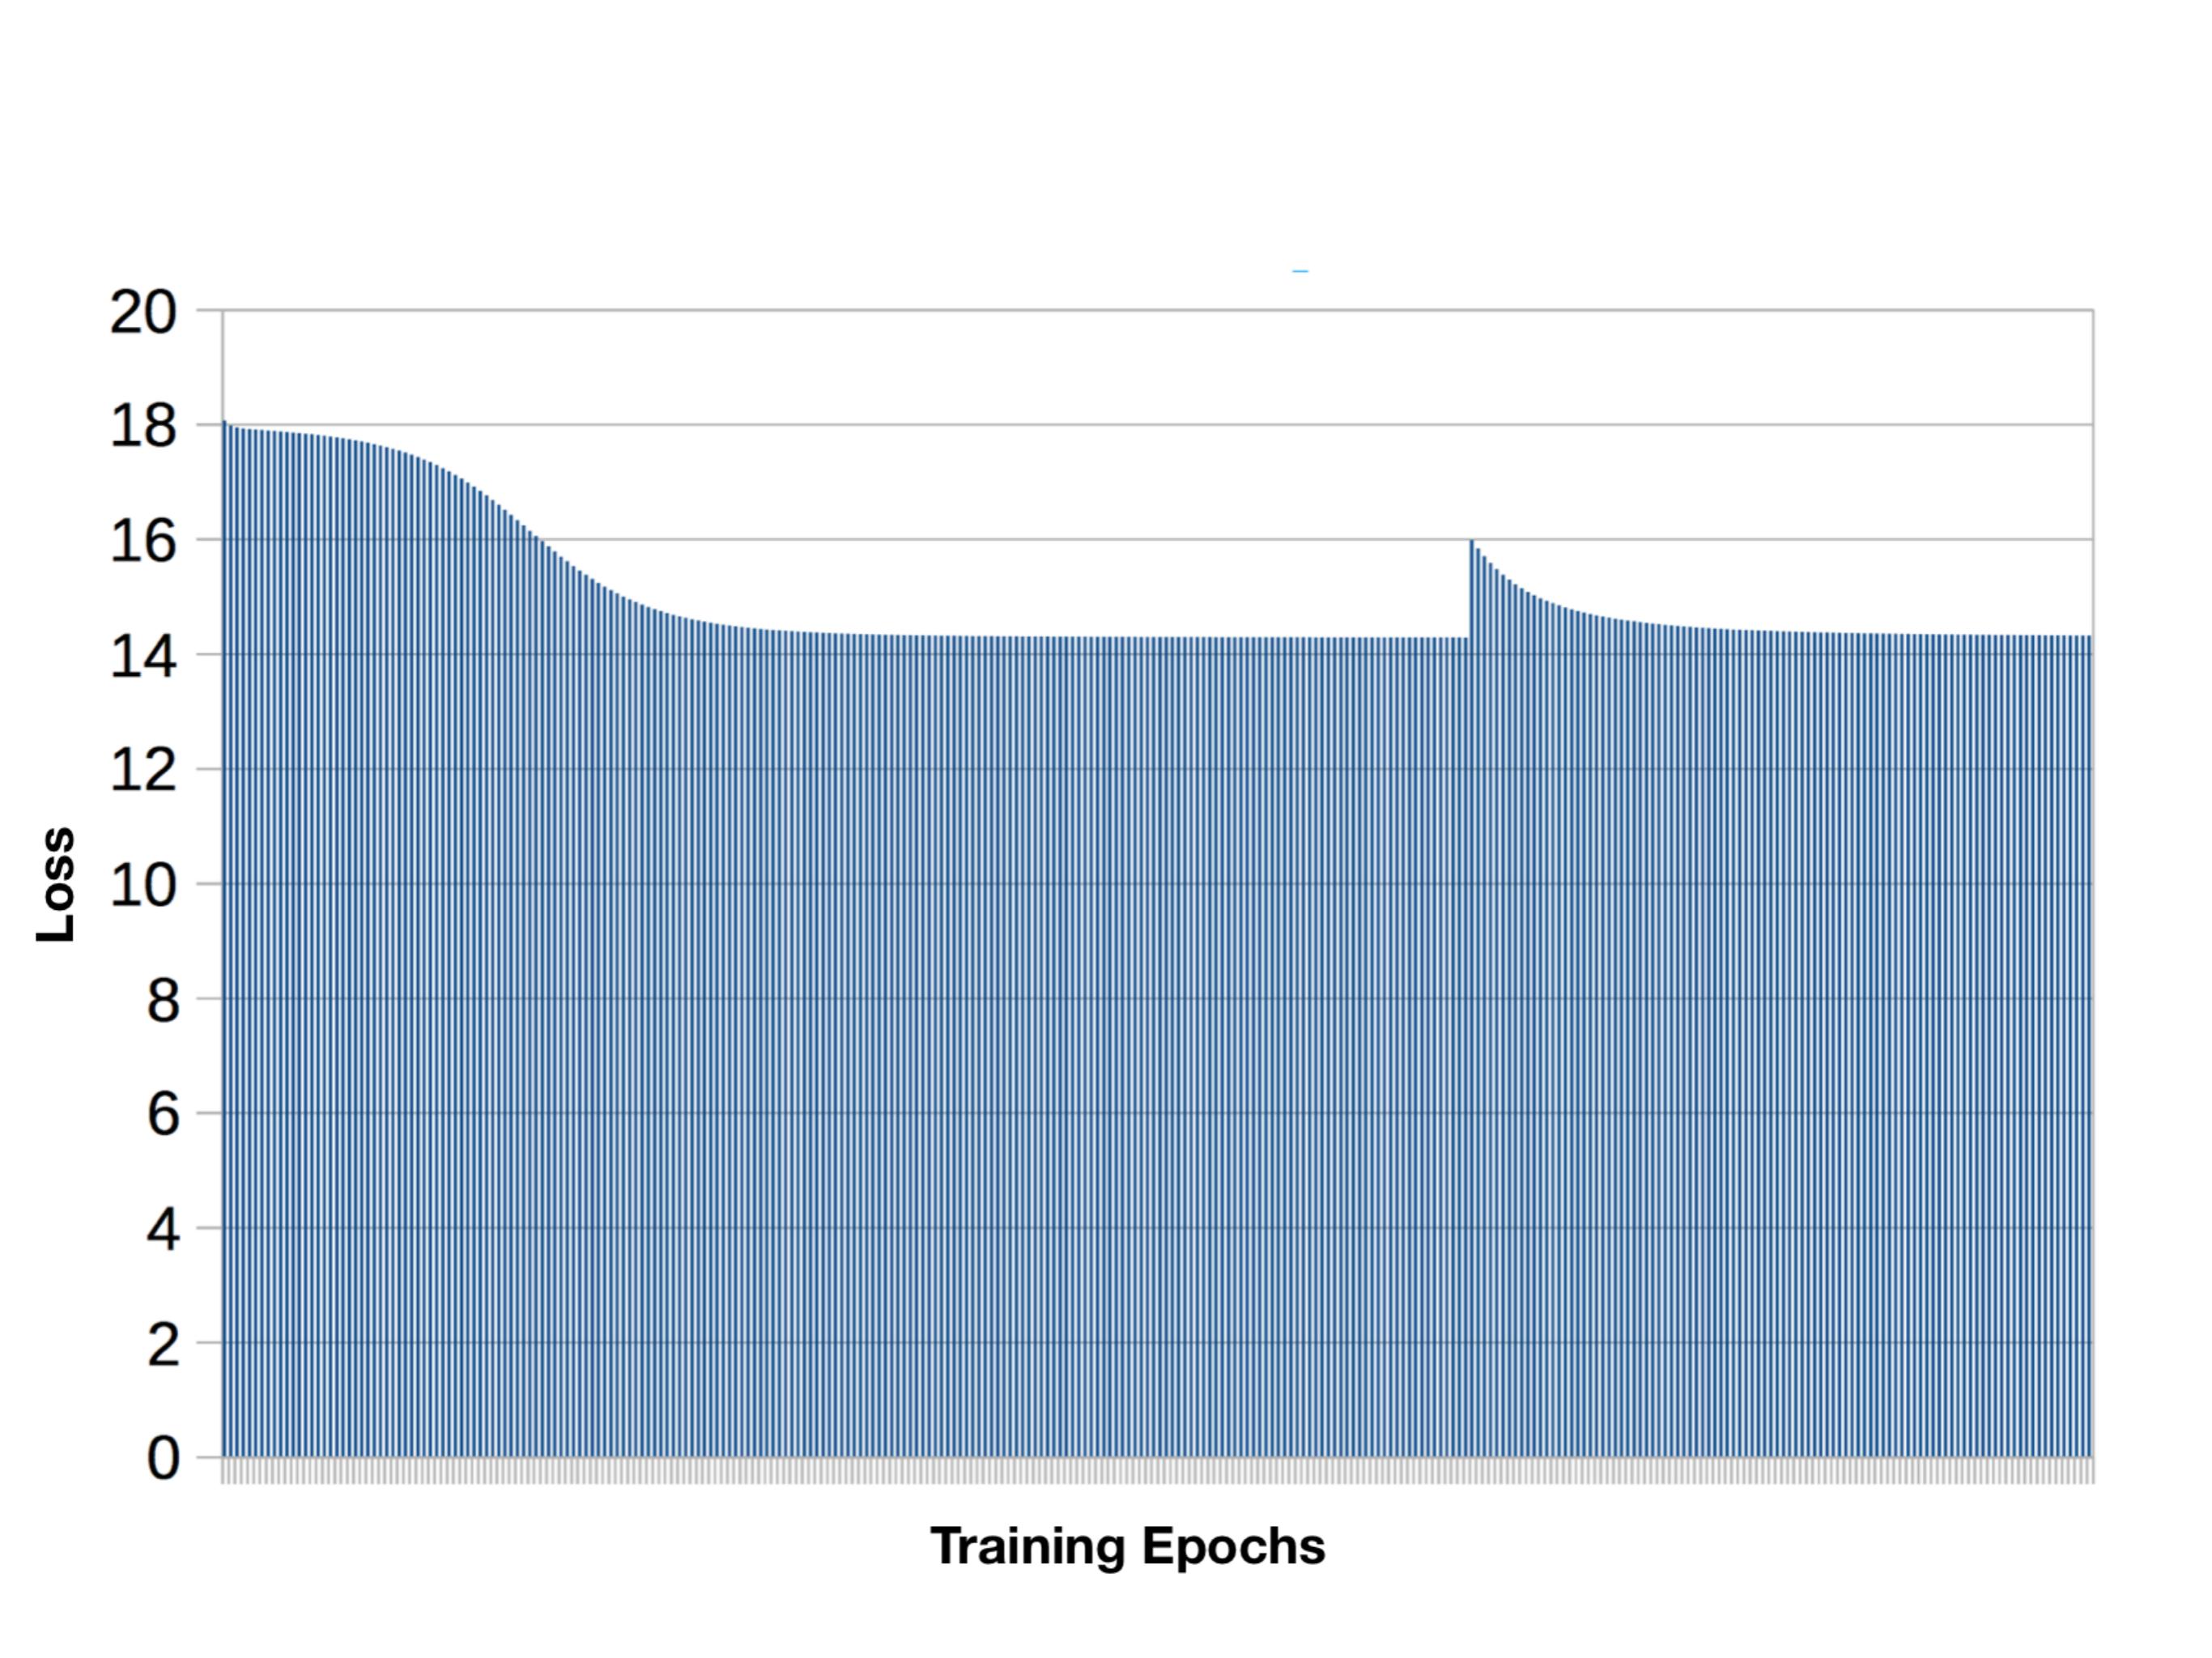
Enzo Busseti, Ian Osband, Scott Wong

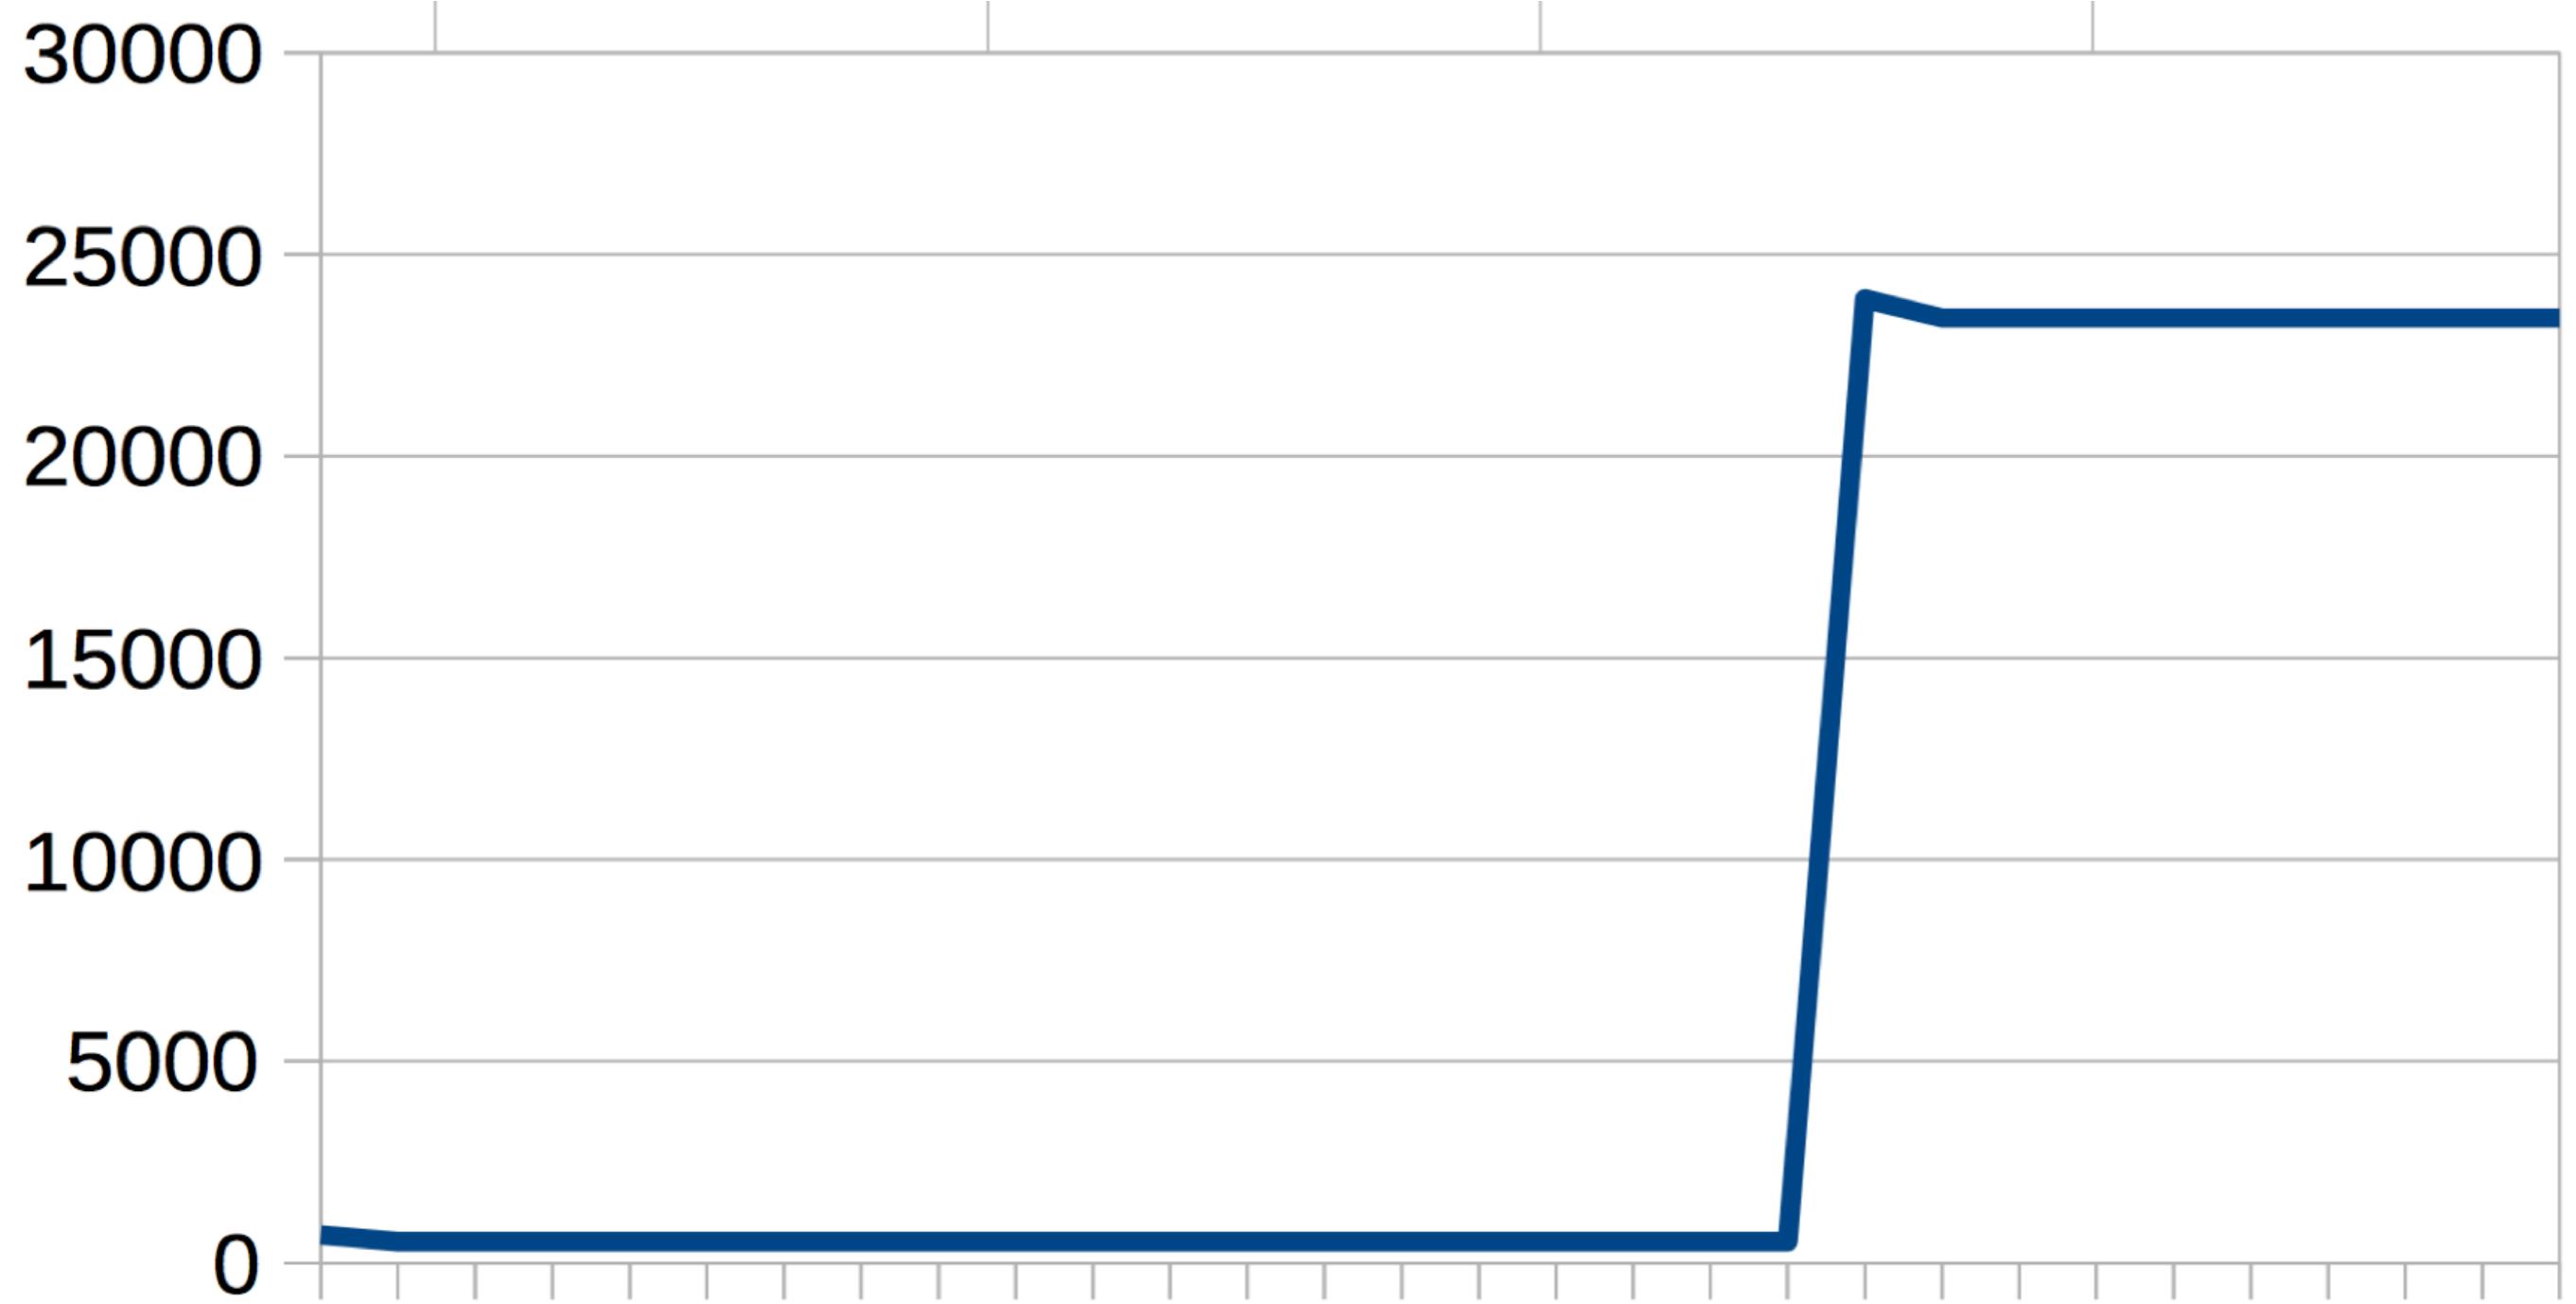


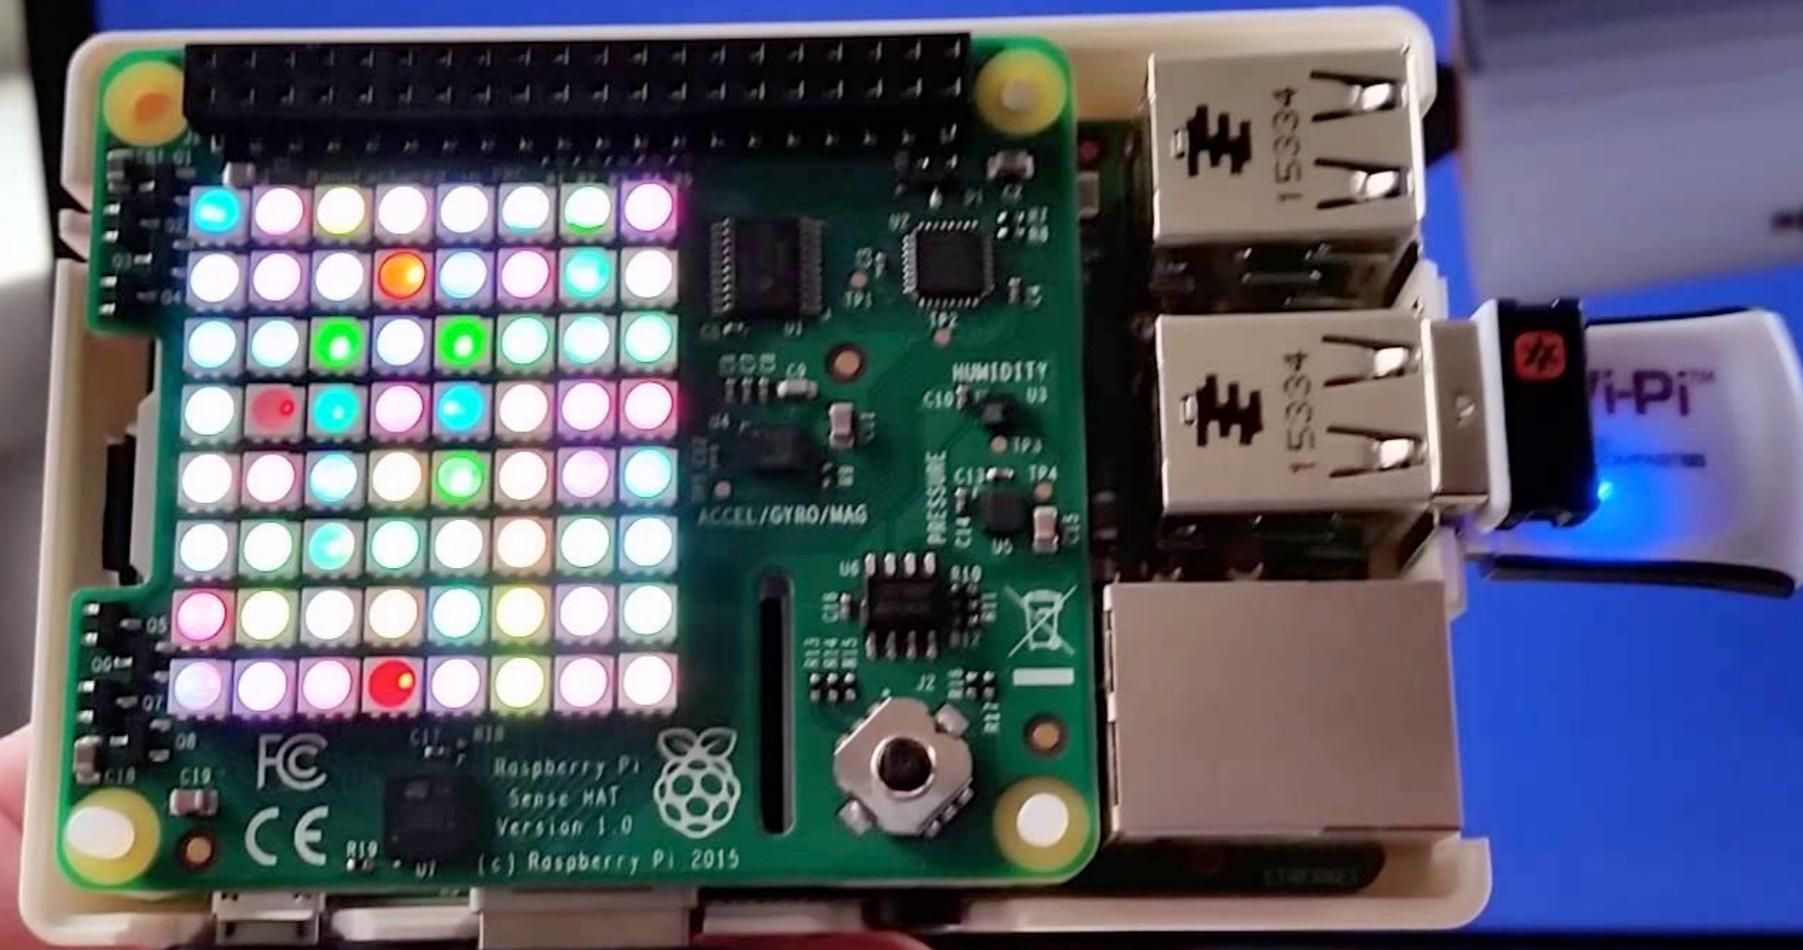
Learning Method	RMSE	% RMSE
Kernelized Regression	1,540	8.3%
Frequency NN	1,251	6.7%
Deep Feedforward NN	1,103	5.9%
Deep Recurrent NN	530	2.8%

Deep Autoencoder



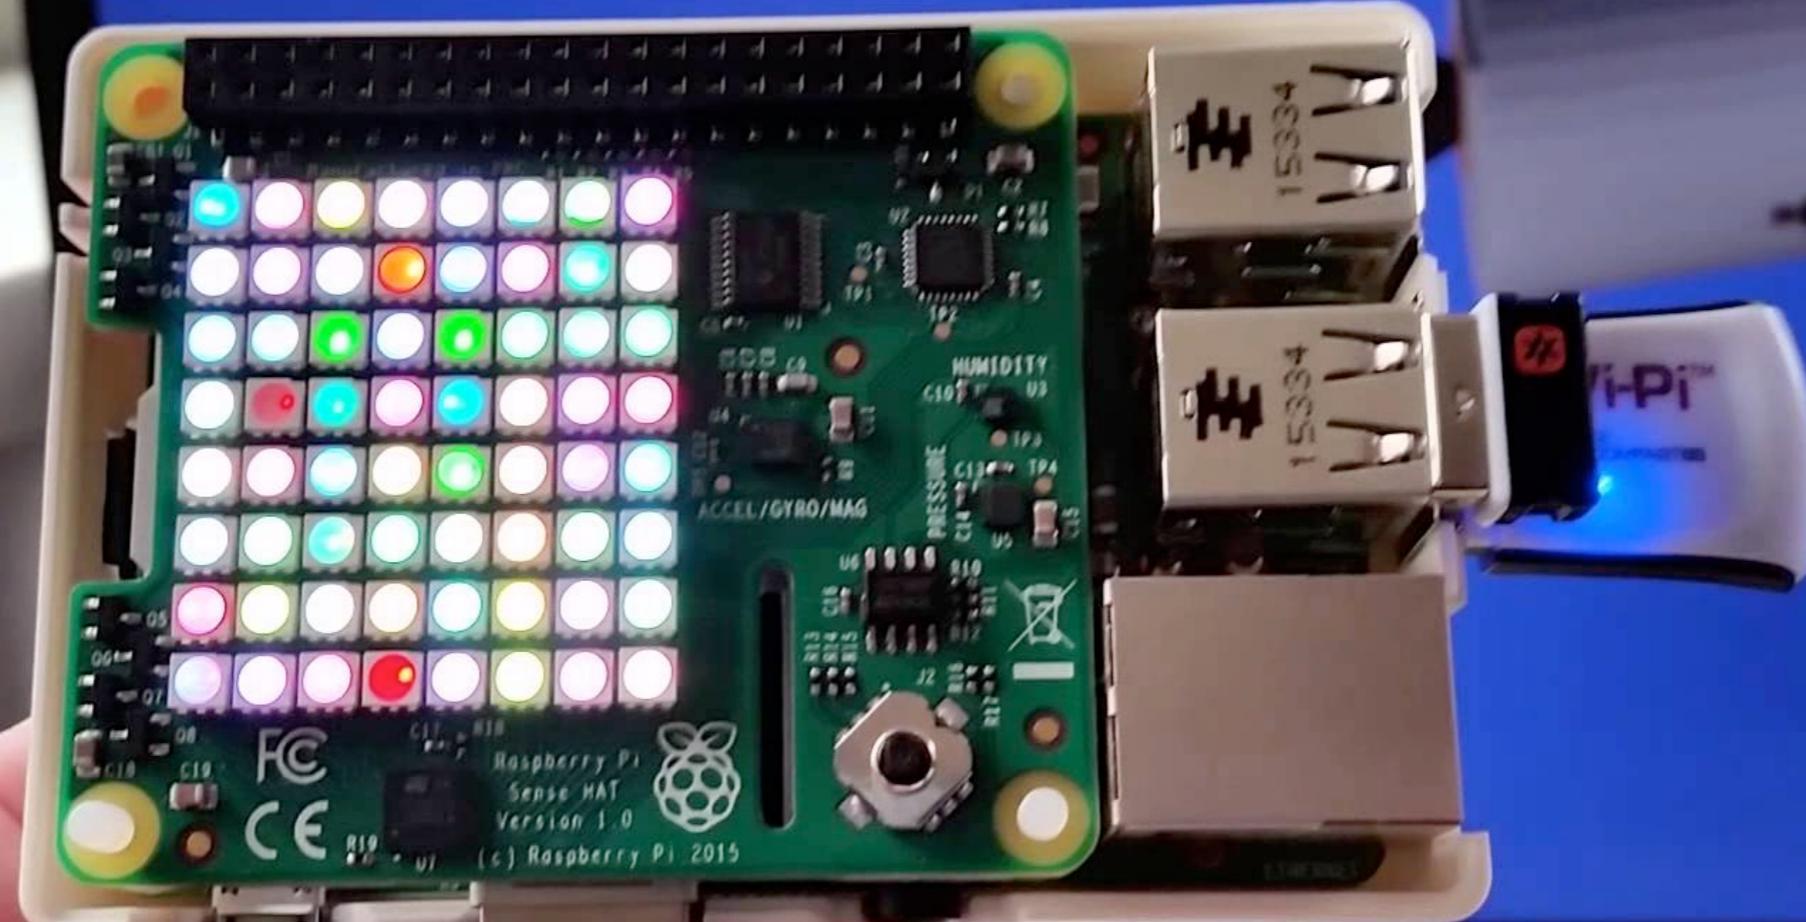






```
import ibmiotf.gateway
from sense_hat import SenseHat
gatewayOptions = {"auth-method": "token", "auth-token": "s*)3uX"}
gatewayCli = ibmiotf.gateway.Client(gatewayOptions)
gatewayCli.connect()
sense = SenseHat()

temp = sense.get_temperature()
myData = { 'temp' : temp}
gatewayCli.publishDeviceEvent(myData, qos=2)
```



```
import ibmiotf.gateway
from sense_hat import SenseHat
gatewayOptions = {"auth-method": "token", "auth-token": "s*)3uX"}
gatewayCli = ibmiotf.gateway.Client(gatewayOptions)
gatewayCli.connect()
sense = SenseHat()
```

```
at Run$MyEventCallback$$.processEvent(Run.scala:43)
at com.ibm.iotf.client.app.ApplicationClient.messageArrived(ApplicationClient.java:973)
at org.eclipse.paho.client.mqttv3.internal.CommsCallback.deliverMessage(CommsCallback.java:477)
```



A developer's guide to the Internet of Things (IoT)

IBM

coursera

A developer's guide
to the Internet of
Things (IoT)

Enroll Now

Starting May 2016

A developer's guide to Exploring and Visualizing IoT Data



Created by: IBM



A developer's guide
to Exploring and
Visualizing IoT Data

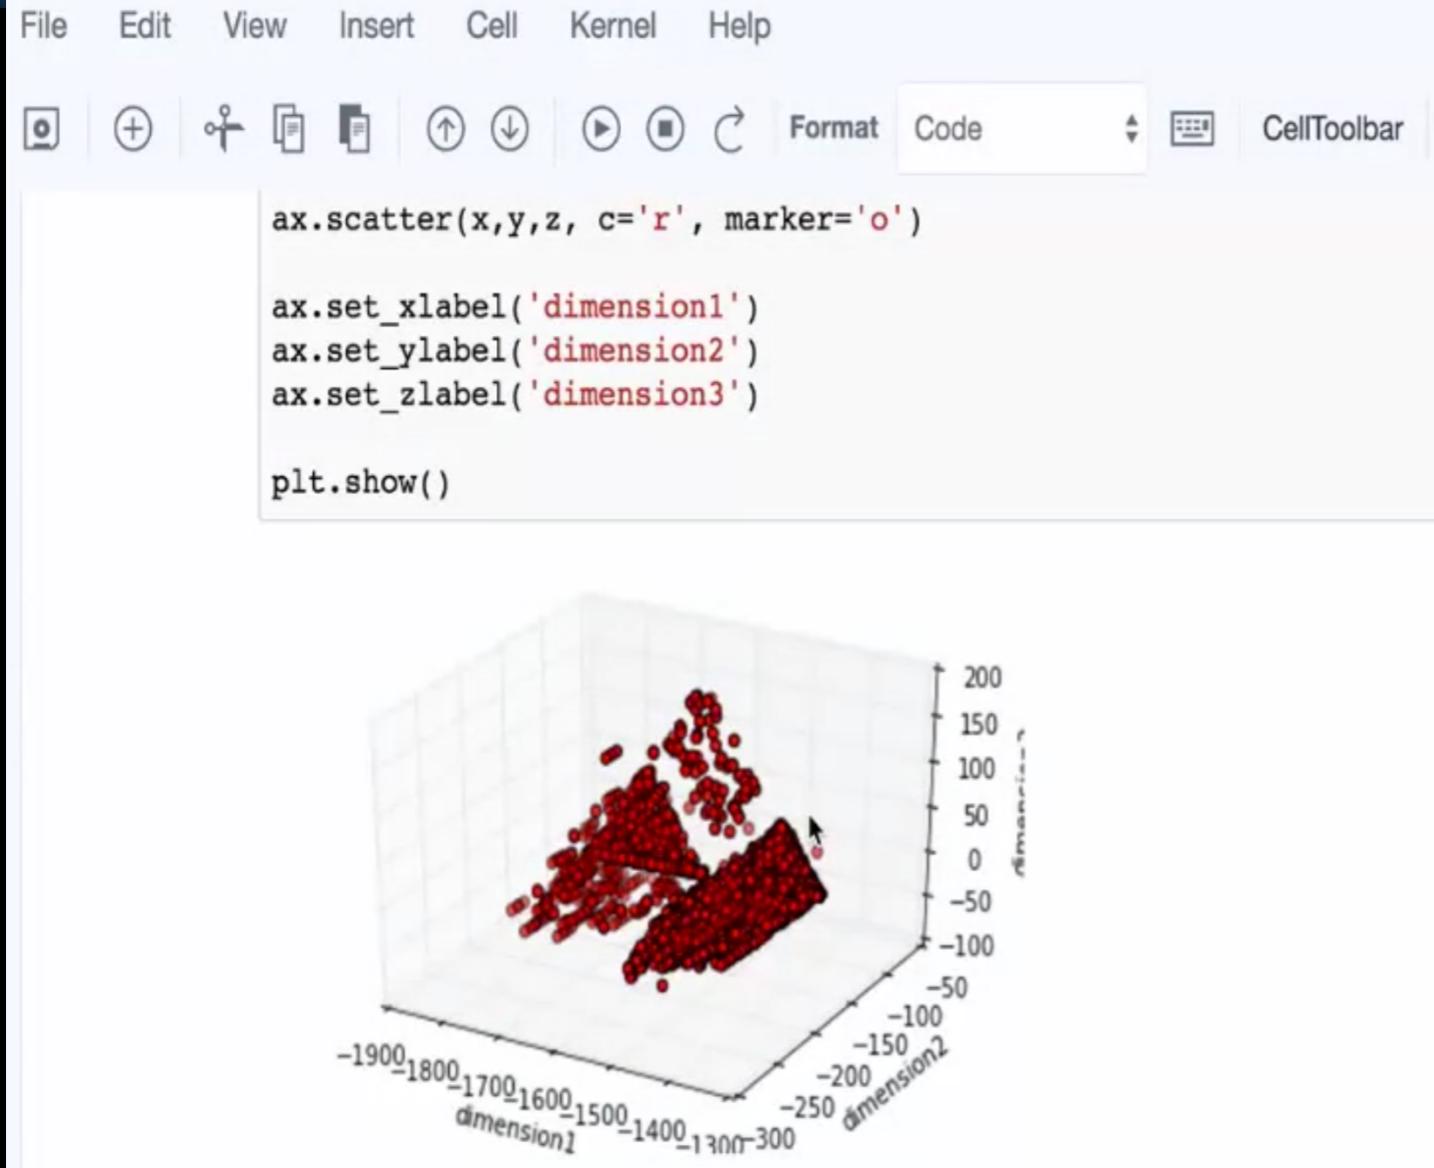
Enroll Now

Starting January 9,
2017



Taught by: Romeo Kienzler, Chief Data Scientist
IBM Watson IoT

A developer's guide to Exploring and Visualizing IoT Data



A developer's guide
to Exploring and
Visualizing IoT Data

Enroll Now
Starting January 9,
2017

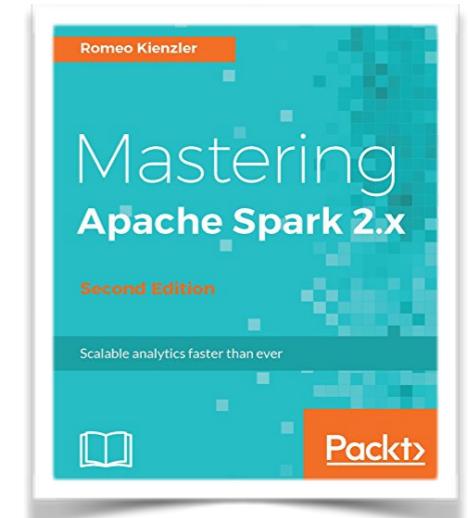


Romeo Kienzler
447 subscribers

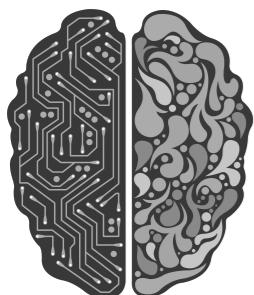
<http://nopenic.ai/tv>



<http://nopenic.ai/slides>



A developer's guide to Exploring and Visualizing IoT Data



<http://nopenic.ai>



<http://bluemix.net>

<http://datascience.ibm.com>

Created by: IBM



Taught by: Romeo Kienzler, Chief Data Scientist
IBM Watson IoT

Disclaimer: This is a beta talk - don't kill me - Disclaimer: This is a beta talk - don't kill me - Disclaimer: This is a beta talk - don't kill me

